

[MUSIC] Hello, and welcome. My name is Dimitri, and I'm happy to see you are interested in competitive data science. Data science is all about machine learning applications. And in data science, like everywhere else, people are looking for the very best solutions to their problems. They're looking for the models that have the best predictive capabilities, the models that make as few mistakes as possible. And the competition for one becomes an essential way to find such solutions. Competing for the prize, participants push through the limits, come up with novel ideas. Companies organize data science competitions to get top quality models for not so high price. And for data scientists, competitions become a truly unique opportunity to learn, well, and of course win a prize. This course is a chance for you to catch up on the trends in competitive data science and learn what we, competition addicts and at the same time, lecturers of this course, have already learned while competing. In this course, we will go through competition solving process step by step and tell you about exploratory data analysis, basic and advanced feature generation and preprocessing, various model validation techniques. Data leakages, competition's metric optimization, model ensembling, and hyperparameter tuning. We've put together all our experience and created this course for you. We've also designed quizzes and programming assignments to let you apply your newly acquired skills. Moreover, as a final project, you will have an opportunity to compete with other students and participate in a special competition, hosted on the world's largest platform for data science challenges called Kaggle. Now, let's meet other lecturers and get started. And now, I want to introduce other lecturers of this course. Alexander, Dmitry, Mikhail, and Marios. Mikhail is aka Cassanova, the person who reached the very top of competitive data science. I will tell you a couple of thoughts about the origins of the course. In year 2014, we started our win in data science by joining competitions. We've been meeting every week and discussing the past competitions, solutions, ideas and tweaks what worked and what did not, this exchange of knowledge and experience helped us to learn quickly from each other and improve our skills. Initially our community was small, but over time more and more people were joining. From the format of groups of discussion. We moved on to the format of well organized meetings. Where a speaker makes an overview of his approach and ideas in front of 50 people. These meetings are called machine learning trainings. Now with the help and support of Yandex and get a hundred of participants. Thus we started from zero and learned everything b

y hard work and collaboration. We had an excellent teacher, Alexander D'yakonov who was top one on Kaggle, he took the course on critical data analysis. In Moscow state university and there we're grateful to him. At some point we started to share our knowledge with other people and some of us even started to read lectures at the university. So now we have decided to summarize everything and make it available for everyone. Together. We've finished and processed in about 20 different competitions only on Kaggle and just as many on other not so famous platforms. All of us have a tremendous amount of skill and experience in competitive data science and now we want to share this experience with you. For all of us, competitive data science opened a number of opportunities as the competitions we took part were dedicated to a large variety of tasks. Mikhail works in e-commerce. Alexander builds predictive model for taxi services, Dmitri works with financial data, Mario develops machinery learning frameworks and I am a deep learning researcher. Competitions, without a doubt, became a stepping stone for our careers and believe me, good comparative record will bring success to you as well. We hope you will find something interesting in this course and wish you good luck. Hello and welcome to our course. In this video, I want to give you a sense for what this course is about and I think the best way to do that is to talk about our course goals, our course assignments and our course schedule. So, at the broadest level, this course is about getting the required knowledge and expertise to successfully participate in data science competitions. That's the goal. Now, we're going to prepare this in a systematic way. We start in week one with a discussion of competitions, what are they, how they work, how they are different from real-life industrial data analysis. Then, we're moving to recap of main machine learning models. Besides this, we're going to review software and hardware requirements and common Python libraries for data analysis. After this is done, we'll go through various feature types, how we preprocess these features and generate new ones. Now, because we sometimes need to extract features from text and images, we will elaborate on most popular methods to do it. Finally, we will start working on the final project, the competition. But then we move on to week two. So, having figured out methods to work with data frames and models, we're starting to cover things you first do in a competition. And this is, by the way, a great opportunity to start working on the final project as we proceed through material. So, first in this week, we'll analyze data set in the exploratory data analysis topic or EDA for short. We'll discuss ways to build intuition about the data, explore anonymized features and clean the data set. Our main instrument here will be logic and visualizations. Okay, now, after making EDA, we switch to validation. And here, we'll spend some time talking about different validation strategies, identifying how data is split into train and test and about what problems we may encounter during validation and ways to address those problems. We finish this week with discussion of data leakage and leaderboard problem. We will define data leakage and understand what are leaks, how to discover various leaks and how to utilize them. So basically, this week, we set up the main pipeline for our final project. And at this point, you should have intuition about the data, reliable validation and data leaks explored. After this pipeline i

s ready, we'll focus on the improvement of our solution and that's already the week three. In that week, we'll analyze various metrics for regression and classification and figure out ways to optimize them both while training the model and afterwards. After we will check that we are correct in measure and improvements of our models, we'll define mean-encodings and work on the encoded features. So here, we start with categorical features, how mean-encoded features lead to overfitting and how we balance overfitting with regularization. Then, we'll discuss several extensions to this approach including applying mean-encodings to numeric features and time series, and this is the point where we move on to other advanced features in the week four. Basically, this include statistics and distance-based features, metrics factorizations, feature interactions and t-SNE. These features often are the key to superior performance in competition, so you should implement and optimize them here for the final project. After this, we'll get to hyperparameters optimization. Here, we will revise your knowledge about model tuning in a systematic way and let you apply to the competition. Then, we move onto the practical guide where all of us have summarized most important moments about competitions which became absolutely clear after few years of participation. These include both some general advice on how to choose and participate in the competition and some technical advice, how to set up your pipeline, what to do first and so on. Finally, we'll conclude this week by working on ensembles with Kaz Anova, the Kaggle top one. We'll start with simple linear ensemble, then we continue with bagging and boosting, and finally we'll cover stacking and stacked net approach. And here by the end of this week, you should already have all required knowledge to succeed in a competition. And then finally, we've got the last week. Here we will work to analyze some of our winning solutions in competitions. But all we are really doing in the last week is wrapping up the course, working on and submitting the final project. So, this basic structure of this course. Now, we move through those sections so that you can practice your skills in the course assignments and there are three basic types of assignments in this class: quizzes, programming assignments and the final project. You don't have to do all of these in order to pass the class, you only need to complete the required assignments and you can see which ones those are by looking on the course website. But let's go ahead and talk about the assignments. We begin with the competition. This is going to be the main assignment for you. In fact, we start working on it on the week two. There we do EDA, exploratory data analysis, set up main pipeline that you'll use for the rest of the course and check the competition for leaks. Then in week three we update our solution by optimizing given metric and adding mean-encoded features. After that, in the week four, we further improve our solution by working on advanced features, tune your hyperparameters and uniting models in ensemble. And in last week, we all are wrapping it up and producing solution by Kaggle winning model standards. We ask you to work on the project at your local machine or your server because Coursera computational resources are limited, and using them for the final project can slow down completing programming assignments for the fellow students. And, in fact, this class is mostly about this program and this competition assignment, but we also have

quizzes and programming assignments for you. We include these to give you an opportunity to refine your knowledge about specific parts of this course: how to check data for leakages, how to implement mean encodings, how to produce an ensemble and so on. You can do them at Coursera site directly but you also can download these notebooks and complete them at your local computer or your server. And this basically is an overview of the course goals, course schedule and course assignments. So, let's go ahead and get started.

Hi everyone. We are starting course about machine learning competitions. In this course, you will learn a lot of tricks and best practices about data science competitions. Before we start to learn advanced techniques, we need to understand the basics. In this video, I will explain the main concept of competitions and you will become familiar with competition mechanics. A variety of machinery competition is very high. In some, participants are asked to process texts. In others, to classify picture or select the best advertising. Despite the variety, all of these competitions are very similar in structure. Usually, they consist of the same elements or concepts which we will discuss in this video. Let's start with a data. Data is what the organizers give us as training material. We will use it in order to produce our solution. Data can be represented in a variety of formats. CSV file with several columns, a text file, an archive with pictures, a database dump, a disabled code or even all together.

With the data, usually there is a description. It's useful to read it in order to understand what we'll work with and which feature can be extracted. Here is an example from Kaggle. From the top, we see several files with data, and below, is their description. Sometimes in addition to data issued by organizers, we can use other data. For example, in order to improve image classification model, one may use a publicly available data set of images. But this depends on a particular competition and you need to check the rules. The next concept is a model. This is exactly what we will build during the competition. It's better to think about model not as one specific algorithm, but something that transforms data into answers. The model should have two main properties. It should produce best possible prediction and be reproducible. In fact, it can be very complicated and contain a lot of algorithms, handcrafted features, use a variety of libraries as this model of the winners of the Homesite competition shown on this slide. It's large and includes many components. But in the course, we will learn how to build such models. To compare our model with the model of other participants, we will send our predictions to the server or in other words, make the submission. Usually, you're asked about predictions only. Sources or models are not required. And also there are some exceptions, cool competitions, where participants submit their code. In this course, we'll focus on traditional challenges where a competitor submit only prediction outputs. Often, I can not just provide a so-called sample submission. An example of how the submission file should look like, look at the sample submission from the Zillow competition. In it is the first column. We must specify the ID of the object and then specify our prediction for it. This is typical format that is used in many competitions. Now, we move to the next concept, evaluation function. When you submit predictions, you need to know how good is your model. The quality of the model is de

defined by evaluation function. In essence and simply the function, the text prediction and correct answers and returns a score characterizes the performance of the solution. The simplest example of such a function is the accurate score. This is just a rate of correct answers. In general, there are a lot of such functions. In our course, we will carefully consider some of them. The description of the competition always indicates which evaluation function is used. I strongly suggest you to pay attention to this function because it is what we will try to optimize. But often, we are not interested in the score itself. We should only care about our relative performance in comparison to other competitors. So we move to the last point we are considering, the leaderboard. The leaderboard is the rate which provides you with information about performance of all participating teams. Most machine learning competition platforms keep your submission history, but the leaderboard usually shows only your best score and position. They cannot as that submission score, reveal some information about data set. And, in extreme cases, one can obtain ground truth targets after sending a lot of submissions. In order to handle this, the set is divided into two parts, public and private. This split is hidden from users and during the competition, we see the score calculated only on public subset of the data. The second part of data set is used for private leaderboard which is revealed after the end of the competition. Only this second part is used for final rating. Therefore, a standard competition routine looks like that. You as the competition, you analyze the data, improve model, prepare submission, send it, see leaderboard score. You repeat this action several times. All this time, only public leaderboard is available. By the end of the competition, you should select submissions which will be used for final scoring. Usually, you are allowed to select two final submissions. Choose wisely. Sometimes public leaderboard scores might be misleading. After the competition deadline, public leaderboard is revealed, and its used for the final rating and defining the winners. That was a brief overview of competition mechanics. Keep in mind that many concepts can be slightly different in a particular competition. All details, for example, where they can join into teams or use external data, you will find in the rules. Strongly suggest you to read the rules carefully before joining the competition. Now, I want to say a few words about competition platforms. Although Kaggle is the biggest and most famous one, there is a number of smaller platforms or even single-competition sites like KDD and VizDoom. Although this list will change over time, I believe you will find the competition which is most relevant and interesting for you. Finally, I want to tell you about the reasons to participate in data science competition. The main reason is that competition is a great opportunity for learning. You communicate with other participants, try new approaches and get a lot of experience. Second reason is that competition often offer you non-trivial problems and state-of-the-art approaches. It allows you to broaden the horizons and look at some everyday task from a different point of view. It's also a great way to become recognizable, get some kind of frame inside data science community and receive a nice job offer. The last reason to participate is that you have a chance for winning some money. It shouldn't be the main goal, just a pleasant bonus. In this video, we ana

lyzed the basic concept of the competition, talked about platforms and reasons for participation. In the next video, we will talk about the difference between real life and competitions.[MUSIC] Hi, everyone. In this video we'll learn how to use Kaggle for participation in data science competitions. Let's open kaggle.com. On the Competitions page, we can see a list of currently running competitions. Every competition has a page which consists of title, short description, price budget, number of participating teams, and time before the end. Information involves all previously running competitions, we can find if we click to All. Let's select some challenge and see how it organized. Here, we see several tabs which we'll explore, and let's start with Overview. In the Description section we see an introduction provided by organizers. In the Description, there is a short story about company and tasks, sometimes with illustration. At the Evaluation page, we see the description of the target metric. In this challenge, target metric is the Mean Absolute Error between the logarithmic transform predictions and ground truth values. This page also contains example of sample submission file, which is typical for such kind of competitions. Now let's move to the Prize page. In the Prize, page we can find information about prizes. Take notice that in the title we have information about the whole money budget, and this page, we see how it will be split among winners. I want to highlight that in order to get money, you need not only be in top three teams, but also beat a Zillow benchmark model. Now let's see, Timeline page, which contains all the information about dates. For example, when competition starts, ends, when will the Team Merger deadline and then what month. All the details about competition, we can find in the Rules. So we need to check really the rules. Here we can find that team limit is three individual, that we have maximum of five submissions per day, that you, for example, should be at least 18 years old to participate. And that, find it, that external data are not allowed. I strongly suggest you to read the rules carefully before joining the competition. And after reading, you should accept it, but I already accepted it. Now, let's check this, Data. Here we have data provided by the organizers, several files which we can download, and sample submission among them, and the description of the data. Here we have description of files, description of data fields, and more importantly a description of train and test split. This is quite useful information in

order to set up right validation scheme. If you have any question about data or other questions to ask, or insights to share, you can go to the forum, which we can find under Discussion tab. Usually it contains a lot of topics or threads, like Welcome, questions about validations, questions about train and test data, and so on and so on. Every topic has title, number of comments, and number of reports. Let's see some of them. Here we have main message, a lot of comments, in this particular we have only one comment. Each we can up vote or down vote and reply to by clicking the reply button. That was a brief overview on forum and now we switch to the Kernels. Usually, I run my code locally, but sometimes it would be handy to check an idea quickly or share code with other participants or teammates. This is what Kernels are for. You can think of Kernel as a small virtual machine in which you write your code, execute it, and share it. Let's take a look at some Kernel, for example for this one. This shows explanatory data analysis on the Zillow competition. It took quite long, contains a lot of pictures, and I believe it's very useful. Here we can see comments for this, different versions. And in order, if you want to make a copy and edit it, we need to Fork this Notebook. It doesn't matter how your predictions were produced, locally or by Kernel, you should submit them through a specialized form. So go back to the competition. Go to submissions. I already submit sample submission, you can do the same. Click submit predictions, and drag and drop file here. Let's look at my submission. After submission, you will see it on the leaderboard. This is my sample submission. Leaderboard contains information about all the teams. So here we have team name or just name in case of single competition team. Score which we produced, number of submissions, time since the last submissions, and position data over seven last days. For example, this means that this guy drops 19 positions during the last week. That was a brief overview of Kaggle interface. Further, I will tell some extra information about the platform. So let's move to Overview page at the bottom. And here, we see information about points and tiers. As mentioned here, the competition will be counting towards ranking points and tiers. If you participate, it will be beneficial for your rating. Sometimes, especially in educational competitions, it's not like that. Information about Kaggle Progression System we can find if we click this link, where we can read info

information about tiers like novice, contributor, master, grandmaster. About medals and ranking points. This ranking points, I use for global User Ranking. Let's check it. So, we have user ranking page, and we see all the users ranked, and with links to their profile. Let's check some profile, for example mine. And here we have photo, name, some information, geo information, information about past competitions, medals, and so on. In addition, I want to say a few words about ability to host competition. Kaggle has this ability. Click Host competition, and there is special Kaggle in class. At in class, everyone can host their own competition for free and invite people to participate. This option is quite often used in various educational competitions. So this was a brief overview of Kaggle platform. Thank for your attention. [MUSIC] In this video, I want to talk about complexity of real world machine learning pipelines and how they differ from data science competitions. Also, we will discuss the philosophy of the competitions. Real world machine learning problems are very complicated. They include several stages, each of them is very important and require attention. Let's imagine that we need to build an anti-spam system and consider the basic steps that arise when building such a system. First of all, before doing any machine learning stuff, you need to understand the problem from a business point of view. What do you want to do? For what? How can it help your users? Next, you need to formalize the task. What is the definition of spam? What exactly is to be predicted? The next step is to collect data. You should ask yourself, what data can we use? How to mine examples of spam and non-spam? Next, you need to take care of how to clean your data and pre-process it. After that, you need to move on to building models. To do this, you need to answer the questions, which class of model is appropriate for this particular task? How to measure performance? How to select the best model? The next steps are to check the effectiveness on the model in real scenario, to make sure that it works as expected and there was no bias introduced by learning process. Does the model actually block spam? How often does it block non-spam emails? If everything is fine, then the next step is to deploy the model. Or in other words, make it available to users. However, the process doesn't end here. You need to monitor the model performance and re-train it on new data. In addition, you need to periodically revise your understanding of the problem and go for the cycle again and again. In contrast, in competitions we have a much simpler situation. All things about formalization and evaluation are already done. All data collected and target metrics fixed. Therefore your main focus on pre-processing the data, picking models and selecting the best ones. But, sometimes you need to understand the business problem in order to get insights or generate a new feature. Also sometimes organizers allow the usage of external data. In such cases, data collection become a crucial part of the solution. I want to show you the difference between real life applications and competitions more thoroughly. This table shows that competitions are much simpler than real world machine learning problems. The hardest part, problem formalization and



d choice of target metric, is already done. Also questions related to deploying out of scope, so participants can focus just on modeling part. One may notice that in this table data collection and model complexity roles have no and yes in competition column. The reason for that, that in some competitions you need to take care of these things. But usually it's not the case. I want to emphasize that as competitors, the only thing we should take care about is target metrics value. Speed, complexity and memory consumption, all this doesn't matter as long as you're able to calculate it and re-produce your own results. Let's highlight key points. Real world machine learning pipelines are very complicated and consist of many stages. Competitions, add weight to a lot of things about modeling and data analysis, but in general they don't address the questions of formalization, deployment and testing. Now, I want to say a few words about philosophy on competitions, in order to form a right impression. We'll cover these ideas in more details later in the course along with examples. The first thing I want to show you is that, machine learning competitions are not only about algorithms. An algorithm is just a tool. Anybody can easily use it. You need something more to win. Insights about data are usually much more useful than a returned ensemble. Some competitions could be solved analytically, without any sophisticated machine learning techniques. In this course, we will show you the importance of understanding your data, tools to use and features you tried to exploit in order to produce the best solution. The next thing I want to say, don't limit yourself. Keep in mind that the only thing you should care about is target metric. It's totally fine to use heuristics or manual data analysis in order to construct golden feature and improve your model. Besides, don't be afraid of using complex solutions, advance feature engineering or doing the huge gritty calculation overnights. Use all the ways you can find in order to improve your model. After passing this course, you will be able to get the maximum gain from your data. And now the important aspect is creativity. You need to know traditional approaches of solid machine learning problems but, you shouldn't be bounded by them. It's okay to modify or hack existing algorithm for your particular task. Don't be afraid to read source codes and change them, especially for deploying stuff. In our course, we'll show you examples of how a little bit of creativity can lead to constructing golden features or entire approaches for solving problems. In the end, I want to say enjoy competitions. Don't be obsessed with getting money. Experience and fun you get are much more valuable than the price. Also, networking is another great advantage of participating in data science competition. I hope you find this course interesting. Hi, everyone. In this video, I want to do a brief overview of basic machine learning approaches and ideas behind them. There are several famous of machine learning algorithms which I want to review. It's a Linear Model, Tree-Based Methods, k-Nearest Neighbors, and Neural Nets. For each of this family, I will give a short intuitive explanation with examples. If you don't remember any of these topics, I strongly encourage you to learn it using links from additional materials. Let's start with Linear Models. Imagine that we have two sets of points, gray points belong to one class and green ones to another. It is very intuitive to separate them with a line. In this case, it's quite sim

ple to do since we have only two dimensional points. But this approach can be generalized for a high dimensional space. This is the main idea behind Linear Models. They try to separate objects with a plane which divides space into two parts. You can remember several examples from this model class like logistic regression or SVM. They all are Linear Models with different loss functions. I want to emphasize that Linear Models are especially good for sparse high dimensional data. But you should keep in mind the limitations of Linear Models. Often, point cannot be separated by such a simple approach. As an example, you can imagine two sets of points that form rings, one inside the other. Although it's pretty obvious how to separate them, Linear Models are not an appropriate choice either and will fail in this case. You can find implementations of Linear Models in almost every machine learning library. Most known implementation in Scikit-Learn library. Another implementation which deserves our attention is Vowpal Wabbit, because it is designed to handle really large data sets.

We're finished with Linear Model here and move on to the next family, Tree-Based Methods. Tree-Based Methods use decision tree as a basic block for building more complicated models. Let's consider an example of how decision tree works. Imagine that we have two sets of points similar to a linear case. Let's separate one class from the other by a line parallel to the one of the axes. We use such restrictions as it significantly reduces the number of possible lines and allows us to describe the line in a simple way. After setting the split as shown at that picture, we will get two sub spaces, upper will have probability of gray=1, and lower will have probability of gray=0.2. Upper sub-space doesn't require any further splitting. Let's continue splitting for the lower sub-space. Now, we have zero probability on gray for the left sub-space and one for the right. This was a brief overview of how decision tree works. It uses divide-and-conquer approach to recur sub-split spaces into sub-spaces. Intuitively, single decision tree can be imagined as dividing space into boxes and approximating data with a constant inside of these boxes. The way of true axis splits and corresponding constants produces several approaches for building decision trees. Moreover, such trees can be combined together in a lot of ways. All this leads to a wide variety of tree-based algorithms, most famous of them being random forest and Gradient Boosted Decision Trees. In case if you don't know what are that, I strongly encourage you to remember these topics using links from additional materials. In general, tree-based models are very powerful and can be a good default method for tabular data. In almost every competitions, winners use this approach. But keep in mind that for Tree-Based Methods, it's hard to capture linear dependencies since it requires a lot of splits. We can imagine two sets of points which can be separated with a line. In this case, we need to grow a tree with a lot of splits in order to separate points. Even in such case, our tree could be inaccurate near decision border, as shown on the picture. Similar to Linear Models, you can find implementations of tree-based models in almost every machine learning library. Scikit-Learn contains quite good implementation of random forest which I personally prefer. All the Scikit-Learn contain implementation of gradient boost decision trees. I prefer to use libraries like XGBoost and LightGBM for their higher speed and accuracy. S

o, here we end the overview of Tree-Based Methods and move on to the k-NN. Before I start the explanation, I want to say that k-NN is abbreviation for k-Nearest Neighbors. One shouldn't mix it up with Neural Networks. So, let's take a look at the familiar binary classification problem. Imagine that we need to predict label for the points shown with question mark at this slide. We assume that points close to each other are likely to have similar labels. So, we need to find the closest point which displayed by arrow and pick its label as an answer. This is how nearest neighbor's method generally works. It can be easily generalized for k-NN, if we will find k-nearest objects and select plus labeled by majority vote. The intuition behind k-NN is very simple. Closer objects will likely to have same labels. In this particular example, we use square distance to find the closest object. In general case, it can be meaningless to use such a distance function. For example, square distance over images is unable to capture semantic meaning. Despite simplicity of the approach, features based on nearest neighbors are often very informative. We will discuss them in more details later in our course. Implementations of k-NN can be found in a lot of machine learning libraries. I suggest you to use implementation from Scikit-Learn since it uses algorithm matrix to speedup calculations and allows you to use several predefined distance functions. Also, it allows you to implement your own distance function. The next big class of model I want to overview is Neural Networks. Neural Nets is a special class of machine learning models, which deserve a separate topic. In general, such methods can be seen in this Black-Box which produce a smooth separating curve in contrast to decision trees. I encourage you to visit TensorFlow playground which is shown on the slide, and play with different parameters of the simple feed-forward network in order to get some intuition about how feed-forward Neural Nets works. Some types of Neural Nets are especially good for images, sounds, text, and sequences. We won't cover details of Neural Nets in this course. Since Neural Nets attracted a lot of attention over the last few years, there are a lot of frameworks to work with them. Packages like TensorFlow, Keras, MXNet, PyTorch, and Lasagne can be used to feed Neural Nets. I personally prefer PyTorch since it provides flexible and user-friendly way to define complex networks. After this brief recap, I want to say a few words about No Free Lunch Theorem. Basically, No Free Lunch Theorem states that there is no methods which outperform all others on all tasks, or in other words, for every method, we can construct a task for which this particular method will not be the best. The reason for that is that every method relies on some assumptions about data or task. If these assumptions fail, Limited will perform poorly. For us, this means that we cannot every competition with just a single algorithm. So we need to have a variety of tools based off different assumptions. Before the end of this video, I want to show you an example from Scikit-Learn library, which plots decision surfaces for different classifiers. We can see the type of algorithm have a significant influence of decision boundaries and consequently on [inaudible]. I strongly suggest you to dive deeper into this example and make sure that you have intuition why these classifiers produce such surfaces. In the end, I want to remind you the main points of this video. First of all, there is no silver bullet algo

rithm which outperforms all the other in all and every task. Next, is that Linear Model can be imagined as splitting space into two sub-spaces separated by a hyper plane. Tree-Based Methods split space into boxes and use constant the predictions in every box. k-NN methods are based on the assumptions that close objects are likely to have same labels. So we need to find closest objects and pick their labels. Also, k-NN approach heavily relies on how to measure point closeness. Feed-forward Neural Nets are harder to interpret but they produce smooth non-linear decision boundary. The most powerful methods are Gradient Boosted Decision Trees and Neural Networks. But we shouldn't underestimate Linear Models and k-NN because sometimes, they may be better. We will show you relevant examples later in our course. Thank you for your attention.

Hi, everyone. In this video, I want to do an overview of hardware and software requirements. You will know what is typical stuff for data science competitions. I want to start from hardware related things. Participating in competitions, you generally don't need a lot of computation resources. A lot of competitions, except imaged based, have under several gigabytes of data. It's not very huge and can be processed on a high level laptop with 16 gigabyte ram and four physical cores. Quite a good setup is a tower PC with 32 gigabyte of ram and six physical cores, this is what I personally use. You have a choice of hardware to use. I suggest you to pay attention to the following things. First is RAM, for this more is better. If you can keep your data in memory, your life will be much, much easier. Personally, I found 64 gigabytes is quite enough, but some programmers prefer to have 128 gigabytes or even more. Next are cores, the more core you have the more or faster experiments you can do. I find it comfortable to work with fixed cores, but sometimes even 32 are not enough. Next thing to pay attention for is storage. If you work with large datasets that don't fit into the memory, it's crucial to have fast disk to read and write chunks of data. SSD is especially important if you train narrowness or large number of images. In case you really need computational resources. For example, if you are part of team or have a computational heavy approach, you can rent it on cloud platforms. They offer machines with a lot of RAMs, cores, and GPUs. There are several cloud providers, most famous are Amazon AWS, Microsoft's Azure, and Google Cloud. Each one has its own pricing, so we can choose which one best fits your needs and budget. I especially want to draw your attention to AWS spot option. Spot instances enable you to be able to use instance, which can lower your cost significantly. The higher your price for

spot instance is set by Amazon and fluctuates depending on supply and demand for spot instances. Your spot instance runs whenever your bid exceeds the current market price. Generally, it's much cheaper than other options. But you always have risk that your bid will get under current market price, and your source will be terminated. Tutorials about how to setup and configure cloud resources you may find in additional materials. Another important thing I want to discuss is software. Usually, rules in competitions prohibit to use commercial software, since it requires to buy a license to reproduce results. Some competitors prefer R as basic language. But we will describe Python's tech as more common and more general. Python is quite a good language for fast prototyping. It has a huge amount of high quality and open source libraries. And I want to reuse several of them. Let's start with NumPy. It's a linear algebra library to work with dimensional arrays, which contains useful linear algebra routines and random number capabilities. Pandas is a library providing fast, flexible, and expressive way to work with a relational or table of data, both easily and intuitively. It allows you to process your data in a way similar to SQL. Scikit-learn is a library of classic machine learning algorithms. It features various classification, regression, and clustering algorithms, including support vector machines, random forest, and a lot more. Matplotlib is a plotting library. It allows you to do a variety of visualization, like line plots, histograms, scatter plots and a lot more. As IDE, I suggest you to use IPython with Jupyter notebook, since they allow you to work interactively and remotely. The last property is especially useful if you use cloud resources. Additional packages contain implementation of more specific tools. Usually, single packages implement single algorithm. XGBoost and LightGBM packages implement gradient-boosted decision trees in a very efficient and optimized way. You definitely should know about such tools. Keras is a user-friendly framework for neural networks. This new package is an efficient implementation of this new projection method which we will discuss in our course. Also, I want to say a few words about external tools which usually don't have any connection despite, but still very used for computations. One such tool is Vowpal Wabbit. It is a tool designed to provide blazing speed and handle really large data sets, which don't fit into memory. Libfm and libffm implement different types of optimization machines, and often used for sparse data like

click-through rate prediction. Rgf is an alternative base method, which I suggest you to use in ensembles. You can install these packages one by one. But as alternative, you can use byte and distribution like Anaconda, which already contains a lot of mentioned packages. And then, through this video, I want to emphasize the proposed setup is the most common but not the only one. Don't overestimate the role of hardware and software, since they are just tools. Thank you for your attention. [MUSIC][NOISE]

Hi. In every competition, we need to pre-process given data set and generate new features from existing ones. This is often required to stay on the same track with other competitors and sometimes careful feature preprocessing and efficient engineering can give you the edge you strive into achieve. Thus, in the next videos, we will cover a very useful topic of basic feature preprocessing and basic feature generation for different types of features. Namely, we will go through numeric features, categorical features, datetime features and coordinate features. And in the last video, we will discuss missing values. Beside that, we also will discuss dependence of preprocessing and generation on a model we're going to use. So the broad goal of the next videos is to help you acquire these highly required skills. To get an idea of following topics, let's start with an example of data similar to what we may encounter in competition. And take a look at well known Titanic dataset. It stores the data about people who were on the Titanic liner during its last trip. Here we have a typical dataframe to work with in competitions. Each row represents a person and each column is a feature. We have different kinds of features here. For example, the values in Survived column are either 0 or 1. The feature is binary. And by the way, it is what we need to predict in this task. It is our target. So, age and fare are numeric features. Sibings p and parch accounts statement and embarked a categorical features. Ticket is just an ID and name is text. So indeed, we have different feature types here, but do we understand why we should care about different features having different types? Well, there are two main reasons for it, namely, strong connection between preprocessing at our model and common feature generation methods for each feature type. First, let's discuss feature preprocessing. Most of times, we can just take our features, fit our favorite model and expect it to get great results. Each type of feature has its own ways to be preprocessed in order to improve quality of the model. In other words,

joys of preprocessing matter, depends on the model we're going to use. For example, let's suppose that target has nonlinear dependency on the pclass feature. Pclass linear of 1 usually leads to target of 1, 2 leads to 0, and 3 leads to 1 again. Clearly, because this is not a linear dependency linear model, one get a good result here. So in order to improve a linear model's quality, we would want to preprocess pclass feature in some way. For example, with the so-called which will replace our feature with three, one for each of pclass values. The linear model will fit much better now than in the previous case. However, random forest does not require this feature to be transformed at all. Random forest can easily put each pclass in separately and predict fine probabilities. So, that was an example of preprocessing. The second reason why we should be aware of different feature text is to ease generation of new features. Feature types different in this and comprehends in common feature generation methods. While gaining an ability to improve your model through them. Also understanding of basics of feature generation will aid you greatly in upcoming advanced feature topics from our course. As in the first point, understanding of a model here can help us to create useful features. Let me show you an example. Say, we have to predict the number of apples a shop will sell each day next week and we already have a couple of months sales history as train in data. Let's consider that we have an obvious linear trend through out the data and we want to inform the model about it. To provide you a visual example, we prepare the second table with last days from train and first days from test. One way to help model neutralize linear train is to add feature indicating the week number past. With this feature, linear model can successfully find an existing linear and dependency. On the other hand, a gradient boosted decision tree will use this feature to calculate something like mean target value for each week. Here, I calculated mean values manually and printed them in the dataframe. We're going to predict number of apples for the sixth week. Note that we indeed have here. So let's plot how a gradient within the decision tree will complete the weak feature. As we do not train Gradient boosting decision tree on the sixth week, it will not put splits between the fifth and the sixth weeks, then, when we will bring the numbers for the 6th week, the model will end up using the wave from the 5th week. As we can see unfortunately,

no users shall land their train here. And vice versa, we can come up with an example of generated feature that will be beneficial for decisions three. And useful spolinari model. So this example shows us, that our approach to feature generation should rely on understanding of employed model. To summarize this feature, first feature preprocessing is necessary instrument you have to use to adapt data to your model. Second, feature generation is a very powerful technique which can aid you significantly in competitions and sometimes provide you the required edge. And at last, both feature preprocessing and feature generation depend on the model you are going to use. So these three topics, in connection to feature types, will be general theme of the next videos. We will thoroughly examine most frequent methods which you can be able to incorporate in your solutions. Good luck. [SOUND] [MUSIC]

Hi. In this video, we will cover basic approach as to feature preprocessing and feature generation for numeric features. We will understand how model choice impacts feature preprocessing. We will identify the preprocessing methods that are used most often, and we will discuss feature generation and go through several examples. Let's start with preprocessing. First thing you need to know about handling numeric features is that there are models which do and don't depend on feature scale. For now, we will broadly divide all models into tree-based models and non-tree-based models. For example, decision trees classifier tries to find the most useful split for each feature, and it won't change its behavior and its predictions. It can multiply the feature by a constant and to retrain the model. On the other side, there are models which depend on these kind of transformations. The model based on your nearest neighbors, linear models, and neural network. Let's consider the following example. We have a binary classification test with two features. The object in the picture belong to different classes. The red circle to class zero, and the blue cross to class one, and finally, the class of the green object is unknown. Here, we will use a one nearest neighbor's model to predict the class of the green object. We will measure distance using square distance, which is also called altometric. Now, if we calculate distances to the red circle and to the blue cross, we will see that our model will predict class one for the green object because the blue cross of class one is much closer than the red circle. But if we multiply the first feature by 10, the red circle will become the closest object, and we will get an opposite prediction. Let's now consider two extreme cases. What will happen if we multiply the first feature by zero and by one million? If the feature is multiplied by zero, then every object will have feature value of zero, which results in KNN ignoring that feature. On the opposite, if the feature is multiplied by one million, slightest differences in that feature's values will impact prediction, and this will result in KNN favoring that feature over all others. Great, but what about other models? Linear models are also experiencing difficulties with differently scaled features. First, we want regularization to be applied to linear models coefficients for features in equal amount. But in fact, regulariza



tion impact turns out to be proportional to feature scale. And second, gradient descent methods can go crazy without a proper scaling. Due to the same reasons, neural networks are similar to linear models in the requirements for feature preprocessing. It is important to understand that different features scalings result in different models quality. In this sense, it is just another hyper parameter you need to optimize. The easiest way to do this is to rescale all features to the same scale. For example, to make the minimum of a feature equal to zero and the maximum equal to one, you can achieve this in two steps. First, we subtract at minimum value. And second, we divide the difference base maximum. It can be done with `MinMaxScaler` from `sklearn`. Let's illustrate this with an example. We apply the so-called `MinMaxScaler` to two features from the detaining dataset, `Age` and `SibSp`. Looking at histograms, we see that the features have different scale, `ages` between zero and 80, while `SibSp` is between zero and 8. Let's apply `MinMaxScaling` and see what it will do. Indeed, we see that after this transformation, both `age` and `SibSp` features were successfully converted to the same value range of 0,1. Note that distributions of values which we observe from the histograms didn't change. To give you another example, we can apply a scalar named `StandardScaler` in `sklearn`, which basically first subtract mean value from the feature, and then divides the result by feature standard deviation. In this way, we'll get standardized distribution, with a mean of zero and standard deviation of one. After either of `MinMaxScaling` or `StandardScaling` transformations, features impacts on non-tree-based models will be roughly similar. Even more, if you want to use KNN, we can go one step ahead and recall that the bigger feature is, the more important it will be for KNN. So, we can optimize scaling parameter to boost features which seems to be more important for us and see if this helps.

When we work with linear models, there is another important moment that influences model training results. I'm talking about outliers. For example, in this plot, we have one feature, `X`, and a target variable, `Y`. If you fit a simple linear model, its predictions can look just like the red line. But if you do have one outlier with `X` feature equal to some huge value, predictions of the linear model will look more like the purple line. The same holds, not only for features values, but also for target values. For example, let's imagine we have a model trained on the data with target values between zero and one. Let's think what happens if we add a new sample in the training data with a target value of 1,000. When we retrain the model, the model will predict abnormally high values. Obviously, we have to fix this somehow. To protect linear models from outliers, we can clip features values between two chosen values of lower bound and upper bound. We can choose them as some percentiles of that feature. For example, first and 99s percentiles. This procedure of clipping is well-known in financial data and it is called winsorization. Let's take a look at this histogram for an example. We see that the majority of feature values are between zero and 400. But there is a number of outliers with values around -1,000. They can make life a lot harder for our nice and simple linear model. Let's clip this feature's value range and to do so, first, we will calculate lower bound and upper bound values as features values at first and 99s percentiles. After we clip the features values, we can see t

that features distribution looks fine, and we hope now this feature will be more useful for our model. Another effective preprocessing for numeric features is the rank transformation. Basically, it sets spaces between proper assorted values to be equal. This transformation, for example, can be a better option than `MinMaxScaler` if we have outliers, because rank transformation will move the outliers closer to other objects. Let's understand rank using this example. If we apply a rank to the source of array, it will just change values to their indices. Now, if we apply a rank to the not-sorted array, it will sort this array, define mapping between values and indices in this source of array, and apply this mapping to the initial array. Linear models, KNN, and neural networks can benefit from this kind of transformation if we have no time to handle outliers manually. Rank can be imported as a random data function from `scipy`. One more important note about the rank transformation is that to apply to the test data, you need to store the creative mapping from features values to their rank values. Or alternatively, you can concatenate, train, and test data before applying the rank transformation. There is one more example of numeric features preprocessing which often helps non-tree-based models and especially neural networks. You can apply log transformation through your data, or there's another possibility. You can extract a square root of the data. Both these transformations can be useful because they drive too big values closer to the features' average value. Along with this, the values near zero are becoming a bit more distinguishable. Despite the simplicity, one of these transformations can improve your neural network's results significantly. Another important moment which holds true for all preprocessings is that sometimes, it is beneficial to train a model on concatenated data frames produced by different preprocessings, or to mix models training differently-preprocessed data. Again, linear models, KNN, and neural networks can benefit hugely from this. To this end, we have discussed numeric feature preprocessing, how model choice impacts feature preprocessing, and what are the most commonly used preprocessing methods. Let's now move on to feature generation. Feature generation is a process of creating new features using knowledge about the features and the task. It helps us by making model training more simple and effective. Sometimes, we can engineer these features using prior knowledge and logic. Sometimes we have to dig into the data, create and check hypothesis, and use this derived knowledge and our intuition to derive new features. Here, we will discuss feature generation with prior knowledge, but as it turns out, an ability to dig into the data and derive insights is what makes a good competitor a great one. We will thoroughly analyze and illustrate this skill in the next lessons on exploratory data analysis. For now, let's discuss examples of feature generation for numeric features. First, let's start with a simple one. If you have columns, Real Estate price and Real Estate squared area in the dataset, we can quickly add one more feature, price per meter square. Easy, and this seems quite reasonable. Or, let me give you another quick example from the Forest Cover Type Prediction dataset. If we have a horizontal distance to a water source and the vertical difference in heights within the point and the water source, we as well may add combined feature indicating the direct distance to the water from this point. Among

other things, it is useful to know that adding, multiplications, divisions, and other features interactions can be of help not only for linear models. For example, although gradient within decision tree is a very powerful model, it still experiences difficulties with approximation of multiplications and divisions. And adding size features explicitly can lead to a more robust model with less amount of trees. The third example of feature generation for numeric features is also very interesting. Sometimes, if we have prices of products as a feature, we can add new feature indicating fractional part of these prices. For example, if some product costs 2.49, the fractional part of its price is 0.49. This feature can help the model utilize the differences in people's perception of these prices. Also, we can find similar patterns in tasks which require distinguishing between a human and a robot. For example, if we will have some kind of financial data like auctions, we could observe that people tend to set round numbers as prices, and there are something like 0.935, blah, blah,, blah, very long number here. Or, if we are trying to find spambots on social networks, we can be sure that no human ever read messages with an exact interval of one second. Great, these three examples should have provided you an idea that creativity and data understanding are the keys to productive feature generation.

All right, let's summarize this up. In this video, we have discussed numeric features. First, the impact of feature preprocessing is different for different models. Tree-based models don't depend on scaling, while non-tree-based models usually depend on them. Second, we can treat scaling as an important hyperparameter in cases when the choice of scaling impacts predictions quality. And at last, we should remember that feature generation is powered by an understanding of the data. Remember this lesson and this knowledge will surely help you in your next competition. Hi.

In this video, we will cover categorical and ordinal features. We will overview methods to work with them. In particular, what kind of pre-processing will be used for each model type of them?

What is the difference between categorical and ordinal features and how we can generate new features from them? First, let's look at several rows from the Titanic dataset and find categorical features here. Their names are: Sex, Cabin and Embarked. These are usual categorical features but there is one more special, the Pclass feature. Pclass stands for ticket class, and has three unique values: one, two, and three. It is ordinal or, in other words, order categorical feature. This basically means that it is ordered in some meaningful way. For example, if the first class was more expensive than the second, or the more the first should be more expensive than the third. We should make an important note here about differences between ordinal and numeric features. If Pclass would have been a numeric feature, we could say that the difference between first, and the second class is equal to the difference between second and the third class, but because Pclass is ordinal, we don't know which difference is bigger. As these numeric features, we can't sort and integrate an ordinal feature the other way, and expect to get similar performance. Another example for ordinal feature is a driver's license type. It's either A, B, C, or D. Or another example, level of education, kindergarten, school, undergraduate, bachelor, master, and doctoral. These categories are sorted in increasingly complex order.

r, which can prove to be useful. The simplest way to encode a categorical feature is to map its unique values to different numbers. Usually, people referred to this procedure as label encoding. This method works fine with two ways because tree-methods can split feature, and extract most of the useful values in categories on its own. Non-tree-based-models, on the other side, usually can't use this feature effectively. And if you want to train linear model kNN on neural network, you need to treat a categorical feature differently. To illustrate this, let's remember example we had in the beginning of this topic. What if Pclass of one usually leads to the target of one, Pclass of two leads to zero, and Pclass of three leads to one. This dependence is not linear, and linear model will be confused. And indeed, here, we can put linear models predictions, and see they all are around 0.5. This looks kind of set but three on the other side, we'll just make two splits select in each unique value and reaching it independently. Thus, this entries could achieve much better score here using these feature. Let's take now the categorical feature and again, apply label encoding. Let this be the feature Embarked. Although, we didn't have to encode the previous feature Pclass before using it in the model. Here, we definitely need to do this with embarked. It can be achieved in several ways. First, we can apply encoding in the alphabetical or sorted order. Unique way to solve of this feature namely S, C, Q. Thus, can be encoded as two, one, three. This is called label encoder from sklearn works by default. The second way is also labeling coding but slightly different. Here, we encode a categorical feature by order of appearance. For example, s will change to one because it was meant first in the data. Second then c, and we will change c to two. And the last is q, which will be changed to three. This can make sense if all were sorted in some meaningful way. This is the default behavior of pandas.factorize function. The third method that I will tell you about is called frequency encoding. We can encode this feature via mapping values to their frequencies. Even 30 percent for us embarked is equal to c and 50 to s and the rest 20 is equal to q. We can change this values accordingly: c to 0.3, s to 0.5, and q to 0.2. This will preserve some information about values distribution, and can help both linear and three models. First ones, can find this feature useful if value frequency is correlated to its target value. While the second ones can help with less number of split because of the same reason. There is another important moment about frequency encoding. If you have multiple categories with the same frequency, they won't be distinguishable in this new feature. We might apply or run categorization here in order to deal with such ties. It is possible to do like this. There are other ways to do label encoding, and I definitely encourage you to be creative in constructing them. Okay. We just discussed label encoding, frequency encoding, and why this works fine for tree-based-methods. But we also have seen that linear models can struggle with label encoded feature. The way to identify categorical features to non-tree-based-models is also quite straightforward. We need to make new code for each unique value in the future, and put one in the appropriate place. Everything else will be zeroes. This method is called, one-hot encoding. Let's see how it works on this quick example. So here, for each unique value of Pclass feature, we just created a n

ew column. As I said, this works well for linear methods, kNN, or neural networks. Furthermore, one-hot encoding feature is already scaled because minimum this feature is zero, and maximum is one. Note that if you care for a few important numeric features, and hundreds of binary features are used by one-hot encoding, it could become difficult for tree-methods they use first ones efficiently. More precisely, tree-methods will slow down, not always improving their results. Also, it's easy to imply that if categorical feature has too many unique values, we will add too many new columns with a few non-zero values. To store these new array efficiently, we must know about sparse matrices. In a nutshell, instead of allocating space in RAM for every element of an array, we can store only non-zero elements and thus, save a lot of memory. Going with sparse matrices makes sense if number of non-zero values is far less than half of all the values. Sparse matrices are often useful when they work with categorical features or text data. Most of the popular libraries can work with these sparse matrices directly namely, XGBoost, LightGBM, sklearn, and others. After figuring out how to pre-processed categorical features for tree based and non-tree based models, we can take a quick look at feature generation. One of most useful examples of feature generation is feature interaction between several categorical features. This is usually useful for non tree based models namely, linear model, kNN. For example, let's hypothesize that target depends on both Pclass feature, and sex feature. If this is true, linear model could adjust its predictions for every possible combination of these two features, and get a better result. How can we make this happen? Let's add this interaction by simply concatenating strings from both columns and one-hot encoding get. Now linear model can find optimal coefficient for every interaction and improve. Simple and effective. More on features interactions will come in the following weeks especially, in advanced features topic. Now, let's summarize this features. First, ordinal is a special case of categorical feature but with values sorted in some meaningful order. Second, label encoding, basically replace this unique values of categorical features with numbers. Third, frequency encoding in this term, maps unique values to their frequencies. Fourth, label encoding and frequency encoding are often used for tree-based methods. Fifth, One-hot encoding is often used for non-tree-based-methods. And finally, applying One-hot encoding combination one heart and chords into combinations of categorical features allows non-tree-based-models to take into consideration interactions between features, and improve. Fine. We just sorted out it feature pre-process for categorical features, and took a quick look on feature generation. Now, you will be able to apply these concepts in your next competition and get better results. Hi. In this video, we will discuss basic visual generation approaches for datetime and coordinate features. They both differ significantly from numeric and categorical features. Because we can interpret the meaning of datetime and coordinates, we can come up with specific ideas about future generation which we'll discuss here. Now, let's start with datetime. Datetime is quite a distinct feature

because it isn't relying on your nature, it also has several different tiers like year, day or week. Most new features generated from datetime can be divided into two categories. The first one, time moments in a period, and the second one, time passed since particular event. First one is very simple. We can add features like second, minute, hour, day in a week, in a month, on the year and so on and so forth. This is useful to capture repetitive patterns in the data. If we know about some non-common materials which influence the data, we can add them as well. For example, if we are to predict efficiency of medication, but patients receive pills one time every three days, we can consider this as a special time period. Okay now, time seems particular event. This event can be either row-independent or row-dependent. In the first case, we just calculate time passed from one general moment for all data. For example, from here to thousand. Here, all samples will become pairable between each other on one time scale. As the second variant of time since particular event, that date will depend on the sample we are calculating this for. For example, if we are to predict sales in a shop, like in the ROSSMANN's store sales competition. We can add the number of days passed since the last holiday, weekend or since the last sales campaign, or maybe the number of days left to these events. So, after adding these features, our dataframe can look like this. Date is obviously a date, and sales are the target of this task. While other columns are generated features. Week day feature indicates which day in the week is this, daynumber since year 2014 indicates how many days have passed since January 1st, 2014. is\_holiday is a binary feature indicating whether this day is a holiday and days\_till\_holidays indicate how many days are left before the closest holiday. Sometimes we have several datetime columns in our data. The most for data here is to subtract one feature from another. Or perhaps subtract generated features, like once we have, we just have discussed. Time moment inside the period or time passed in zero dependent events. One simple example of third generation can be found in churn prediction task. Basically churn prediction is about estimating the likelihood that customers will churn. We may receive a valuable feature here by subtracting user registration date from the date of some action of his, like purchasing a product, or calling to the customer service. W

e can see how this works  
 on this data dataframe. For every user, we know  
 last\_purchase\_date and last\_call\_date. Here we add the difference between  
 them as new feature named date\_diff. For clarity,  
 let's take a look at this figure. For every user, we have his  
 last\_purchase\_date and his last\_call\_date. Thus, we can add date\_diff  
 feature which indicates number of days between these events. Note that after generation feature is  
 from date time, you usually will get either numeric features like  
 time passed since the year 2000, or categorical features like day of week. And these features now are needed  
 to be treated accordingly with necessary pre-processings we have discussed earlier. Now having discussed feature  
 generation for datetime, let's move onto feature generation for coordinates. Let's imagine that we're trying to  
 estimate the real estate price. Like in the Deloitte competition named  
 Western Australia Rental Prices, or in the Sberbank Russian Housing Market  
 competition. Generally, you can calculate distances to important points on the map. Keep this wonderful map. If you have additional data with  
 infrastructural buildings, you can add as a feature distance to the nearest  
 shop to the second by distance hospital, to the best school in the neighborhood and  
 so on. If you do not have such data, you can extract interesting points on  
 the map from your trained test data. For example, you can do a new  
 map to squares, with a grid, and within each square, find the most expensive flat, and for every other object in this  
 square,  
 add the distance to that flat. Or you can organize your data points into clusters, and then use centers of clusters as such important points. Or again, another way. You can find some special areas,  
 like the area with very old buildings and add distance to this one. Another major approach to use coordinates  
 is to calculate aggregated statistics for objects surrounding area. This can include number of lets  
 around this particular point, which can then be interpreted as a polarity. Or we can add mean realty price, which will indicate how expensive  
 area around selected point is. Both distances and aggregate statistics are often  
 useful in tasks with coordinates. One more trick you need to know about  
 coordinates, that if you train decision trees from them, you can add slightly  
 rotated coordinates as new features. And this will help a model make

more precise selections on the map. It can be hard to know what exact rotation we should make, so we may want to add all rotations to 45 or 22.5 degrees. Let's look at the next example of a relative price prediction. Here the street is dividing an area in two parts. The high priced district above the street, and the low priced district below it. If the street is slightly rotated, trees will try to make a lot of space here. But if we will add new coordinates in which these two districts can be divided by a single split, this will hugely facilitate the rebuilding process. Great, we just summarize the most frequent methods used for future generation from datetime and coordinates. For datetime, these are applying periodicity, calculates in time passed since particular event, and engine differences between two datetime features. For coordinates, we should recall extracting interesting samples from trained test data, using places from additional data, calculating distances to centers of clusters, and adding aggregated statistics for surrounding area. Knowing how to effectively handle datetime and coordinates, as well as numeric and categorical features, will provide you reliable way to improve your score. And to help you devise that specific part of solution which is often required to beat very top scores. [SOUND] Often we have to deal with missing values in our data. They could look like not numbers, empty strings, or outliers like minus 999. Sometimes they can contain useful information by themselves, like what was the reason of missing value occurring here? How to use them effectively? How to engineer new features from them? We'll do the topic for this video. So what kind of information missing values might contain? How can they look like? Let's take a look at missing values in the Springfield competition. This is metrics of samples and features. People mainly reviewed each feature, and found missing values for each column. This latest could be not a number, empty string, minus 1, 99, and so on. For example, how can we find out that -1 can be the missing value? We could draw a histogram and see this variable has uniform distribution between 0 and 1. And that it has small peak of -1 values. So if there are no not numbers there, we can assume that they were replaced by -1. Or the feature distribution plot can look like the second figure. Note that x axis has lock scale. In this case, not a numbers probably were few by features mean value. You can easily generalize this



logic to apply to other cases. Okay on this example we just learned this, missing values can be hidden from us. And by hidden I mean replaced by some other value beside not a number. Great, let's talk about missing value importation. The most often examples are first, replacing not a number with some value outside fixed value range. Second, replacing not a number with mean or median. And third, trying to reconstruct value somehow. First method is useful in a way that it gives three possibility to take missing value into separate category. The downside of this is that performance of linear networks can suffer. Second method usually beneficial for simple linear models and neural networks. But again for trees it can be harder to select object which had missing values in the first place. Let's keep the feature value reconstruction for now, and turn to feature generation for a moment. The concern we just have discussed can be addressed by adding new feature isnull indicating which rows have missing values for this feature. This can solve problems with trees and neural networks while computing mean or median. But the downside of this is that we will double number of columns in the data set. Now back to missing values importation methods. The third one, and the last one we will discuss here, is to reconstruct each value if possible. One example of such possibility is having missing values in time series. For example, we could have everyday temperature for a month but several values in the middle of months are missing. Well of course, we can approximate them using nearby observations. But obviously, this kind of opportunity is rarely the case. In most typical scenario rows of our data set are independent. And we usually will not find any proper logic to reconstruct them. Great, to this moment we already learned that we can construct new feature, isnull indicating which rows contains not numbers. What are other important moments about feature generation we should know? Well there's one general concern about generating new features from one with missing values. That is, if we do this, we should be very careful with replacing missing values before our feature generation. To illustrate this, let's imagine we have a year long data set with two features. Daytime feature and temperature which had missing values. We can see all of this on the figure. Now we fill missing values with some value, for example with median. If you have data over the whole

year  
median probably will be near zero so it should look like that. Now we want to add feature like difference between temperature today and yesterday, let's do this. As we can see, near the missing values this difference usually will be abnormally huge. And this can be misleading our model. But hey, we already know that we can approximate missing values sometimes here by interpolation the error by points, great. But unfortunately, we usually don't have enough time to be so careful here. And more importantly, these problems can occur in cases when we can't come up with such specific solution. Let's review another example of missing value importation. Which will be substantially discussed later in advanced feature [INAUDIBLE] topic. Here we have a data set with independent rows. And we want to encode the categorical feature with the numeric feature. To achieve that we calculate mean value of numeric feature for every category, and replace categories with these mean values. What happens if we fill not the numbers in the numeric feature, with some value outside of feature range like -999. As we can see, all values we will be doing them closer to -999. And the more the row's corresponding to particular category will have missing values. The closer mean value will be to -999. The same is true if we fill missing values with mean or median of the feature. This kind of missing value importation definitely can screw up the feature we are constructing. The way to handle this particular case is to simply ignore missing values while calculating means for each category. Again let me repeat the idea of these two examples. You should be very careful with early non importation if you want to generate new features. There's one more interesting thing about missing values. [INAUDIBLE] boost can handle a lot of numbers and sometimes using this approach can change score drastically. Besides common approaches we have discussed, sometimes we can treat outliers as missing values. For example, if we have some easy classification task with songs which are thought to be composed even before ancient Rome, or maybe the year 2025. We can try to treat these outliers as missing values. If you have categorical features, sometimes it can be beneficial to change the missing values or categories which present in the test data but do not present in the train data. The intention for doing so appeals to the fact that the model which didn't have that category in the train data

will eventually treat it randomly. Here and of categorical features can be of help. As we already discussed in our course, we can change categories to its frequencies and thus to its categories was in before based on their frequency. Let's walk through the example on the slide. There you see from the categorical feature, they not appear in the train. Let's generate new feature indicating number of where the occurrence is in the data. We will name this feature categorical\_encoded. Value A has six occurrences in both train and test, and that's value of new feature related to A will be equal to 6. The same works for values B, D, or C. But now new features various related to D and C are equal to each other. And if there is some dependence in between target and number of occurrences for each category, our model will be able to successfully visualize that. To conclude this video, let's overview main points we have discussed. The choice of method to fill not a numbers depends on the situation. Sometimes, you can reconstruct missing values. But usually, it is easier to replace them with value outside of feature range, like -999 or to replace them with mean or median. Also missing values already can be replaced with something by organizers. In this case if you want know exact rows which have missing values you can investigate this by browsing histograms. More, the model can improve its results using binary feature isnull which indicates what roles have missing values. In general, avoid replacing missing values before feature generation, because it can decrease usefulness of the features. And in the end, Xgboost can handle not a numbers directly, which sometimes can change the score for the better. Using knowledge you have derived from our discussion, now you should be able to identify missing values. Describe main methods to handle them, and apply this knowledge to gain an edge in your next computation. Try these methods in different scenarios and for sure, you will succeed.[MUSIC] Hi. Often in computations, we have data like text and images. If you have only them, we can apply approach specific for this type of data. For example, we can use search engines in order to find similar text. That was the case in the Allen AI Challenge for example. For images, on the other hand, we can use conditional neural networks, like in the Data Science Bowl, and a whole bunch of other competitions. But if we have text or images as additional data, we usually must grasp different features, which can be edited as complementary to our

main data frame of samples and features. Very simple example of such case we can see in the Titanic dataset we have called name, which is more or less like text, and to use it, we first need to derive the useful features from it. Another most surest example, we can predict whether a pair of online advertisements are duplicates, like slightly different copies of each other, and we could have images from these advertisements as complimentary data, like the Avito Duplicates Ads Detection competition. Or you may be given the task of classifying documents, like in the Tradeshift Text Classification Challenge. When feature extraction is done, we can treat extracted features differently. Sometimes we just want to add new features to existing dataframe. Sometimes we even might want to use the right features independently, and in end, make stake in with the base solution. We will go through stake in and we will learn how to apply it later in the topic about ensembles, but for now, you should know that both ways first to acquire, to of course extract features from text and images somehow. And this is exactly what we will discuss in this video. Let's start with featured extraction from text. There are two main ways to do this. First is to apply bag of words, and second, use embeddings like word to vector. Now, we'll talk about a bit about each of these methods, and in addition, we will go through text pre-processings related to them. Let's start with the first approach, the simplest one, bag of words. Here we create new column for each unique word from the data, then we simply count number of occurrences for each word, and place this value in the appropriate column. After applying the separation to each row, we will have usual dataframe of samples and features. In a scalar, this can be done with CountVectorizer. We also can post process calculated metrics using some pre-defined methods. To make out why we need post-processing let's remember that some models like kNN, like neural regression, and neural networks, depend on scaling of features. So the main goal of post-processing here is to make samples more comparable on one side, and on the other, boost more important features while decreasing the scale of useless ones. One way to achieve the first goal of making a sample small comparable is to normalize sum of values in a row. In this way, we will count not

occurrences but frequencies of words. Thus, texts of different sizes will be more comparable. This is the exact purpose of term frequency transformation. To achieve the second goal, that is to boost more important features, we'll make post process our matrix by normalizing data column wise. A good idea is to normalize each feature by the inverse fraction of documents, which contain the exact word corresponding to this feature. In this case, features corresponding to frequent words will be scaled down compared to features corresponding to rarer words. We can further improve this idea by taking a logarithm of these numberization coefficients. As a result, this will decrease the significance of widespread words in the dataset and do require feature scaling. This is the purpose of inverse document frequency transformation. General frequency, and inverse document frequency transformations, are often used together, like an sklearn, in Tfidf Vectorizer. Let's apply Tfidf transformation to the previous example. First, TF. Nice. Occurrences which are switched to frequencies, that means some of variance for each row is now equal to one. Now, IDF, great. Now data is normalized column wise, and you can see, for those of you who are too excited, IDF transformation scaled down the appropriate feature. It's worth mentioning that there are plenty of other variants of Tfidf which may work better depending on the specific data. Another very useful technique is Ngrams. The concept of Ngram is simple, you add not only column corresponding to the word, but also columns corresponding to inconsequent words. This concept can also be applied to sequence of chars, and in cases with low N, we'll have a column for each possible combination of N chars. As we can see, for  $N = 1$ , number of these columns will be equal to 28. Let's calculate number of these columns for  $N = 2$ . Well, it will be 28 squared. Note that sometimes it can be cheaper to have every possible char Ngram as a feature, instead of having a feature for each unique word from the dataset. Using char Ngrams also helps our model to handle unseen words. For example, rare forms of already used words. In a scaled count vectorizer has appropriate parameter for using Ngrams, it is called `Ngram_range`. To change from word Ngrams to char Ngrams, you may use parameter named `analyzer`. Usually, you may want to preprocess text, even before applying bag of words, and sometimes, careful text p

reprocessing can help bag of words drastically. Here, we will discuss such methods as converting text to lowercase, lemmatization, stemming, and the usage of stopwords. Let's consider simple example which shows utility of lowercase. What if we applied bag of words to the sentence very, very sunny? We will get three columns for each word. So because Very, with capital letter, is not the same string as very without it, we will get multiple columns for the same word, and again, Sunny with capital letter doesn't match sunny without it. So, first preprocessing what we want to do is to apply lowercase to our text. Fortunately, configurer from sklearn does this by default. Now, let's move on to lemmatization and stemming. These methods refer to more advanced preprocessing. Let's look at this example. We have two sentences: I had a car, and We have cars. We may want to unify the words car and cars, which are basically the same word. The same goes for had and have, and so on. Both stemming and lemmatization may be used to fulfill this purpose, but they achieve this in different ways. Stemming usually refers to a heuristic process that chops off ending of words and thus unite duration of related words like democracy, democratic, and democratization, producing something like, democr, for each of these words. Lemmatization, on the hand, usually means that you have want to do this carefully using knowledge or vocabulary, and morphological analogies of force, returning democracy for each of the words below. Let's look at another example that shows the difference between stemming and lemmatization by applying them to word saw. While stemming will return on the letter s, lemmatization will try to return either see or saw, dependent on the word's meaning. The last technique for text preprocessing, which we will discuss here, is usage of stopwords. Basically, stopwords are words which do not contain important information for our model. They are either insignificant like articles or prepositions, or so common they do not help to solve our task. Most languages have predefined list of stopwords which can be found on the Internet or logged from NLTK, which stands for Natural Language Toolkit Library for Python. CountVectorizer from sklearn also has parameter related to stopwords, which is called max\_df. max\_df is the threshold of words we can see, after we see in which, the word will be removed from text corpus. Good, we just have discussed classical

feature extraction pipeline for text. At the beginning, we may want to pre-process our text. To do so, we can apply lowercase, stemming, lemmatization, or remove stopwords. After preprocessing, we can use bag of words approach to get the matrix where each row represents a text, and each column represents a unique word. Also, we can use bag of words approach for Ngrams, and in new columns for groups of several consecutive words or chars. And in the end, when we post process these metrics using TFIDF, which often prove to be useful. Well, then now we can add extracted features to our basic data frame, or putting the dependent model on them to create some tricky features. That's all for now. In the next video, we will continue to discuss feature extraction. We'll go through two big points. First, we'll talk about approach for texts, and second, we will discuss feature extraction for images. [MUSIC]Hi and welcome back. In this video, we'll talk about Word2vec approach for texts and then we'll discuss feature extraction for images. After we've summarized pipeline for feature extraction with Bag of Words approach in the previous video, let's overview another approach, which is widely known as Word2vec. Just as the Bag of Words approach, we want to get vector representations of words and texts, but now more concise than before. Word2vec is doing exactly that. It converts each word to some vector in some sophisticated space, which usually have several hundred dimensions. To learn the word embedding, Word2vec uses nearby words. Basically, different words, which often are used in the same context, will be very close in these vectoring representation, which, of course, will benefit our models. Furthermore, there are some prominent examples showing that we can apply basic operations like addition and subtraction on these vectors and expect results of such operations to be interpretable. You should already have seen this example by now somewhere. Basically, if we calculate differences between the vectors of words queen and king, and differences between the vectors of words woman and man, we will find that these differences are very similar to each other. And, if we try to see this from another perspective, and subtract the vector of woman from the vector of king and then add the vector of man, will pretty much again the vector of the word queen. Think about it for a moment. This is fascinating fact and indeed creation of Word2vec approach led to many extensive and far reaching results in the field. There are several implementations of this embedding approach besides Word2vec namely Glove, which stands for Global Vector for word representation. FastText and few others. Complications may occur, if we need to derive vectors not for words but for sentences. Here, we may take different approaches. For example, we can calculate mean or sum of words vectors or we can choose another way and go with special models like Doc2vec. Choice all the way to proceed here depends on and particular situation. Usually, it is better to check both approaches and select the best. Training of

Word2vec can take quite a long time, and if you work with text or some common origin, you may find useful pre-trained models on the internet. For example, ones which are trained on the Wikipedia. Otherwise, remember, the training of Word2vec doesn't require target values from your text. It only requires text to extract context for each word. Note, that all pre-processing we had discussed earlier, namely lowercase stemming, lemmatization, and the usage of stopwords can be applied to text before training Word2vec models. Now, we're ready to summarize difference between Bag of Words and the Word2vec approaches in the context of competition. With Bag of Words, vectors are quite large but is a nice benefit. Meaning of each value in the vector is known. With Word2vec, vectors have relatively small length but values in a vector can be interpreted only in some cases, which sometimes can be seen as a downside. The other advantage of Word2vec is crucial in competitions, is that words with similar meaning will have similar vector representations. Usually, both Bag of Words and Word2vec approaches give quite different results and can be used together in your solution. Let's proceed to images now. Similar to Word2vec for words, convolutional neural networks can give us compressed representation for an image. Let me provide you a quick explanation. When we calculate network output for the image, beside getting output on the last layer, we also have outputs from inner layers. Here, we will call these outputs descriptors. Descriptors from later layers are better way to solve texts similar to one network was trained on. In contrary, descriptors from early layers have more text independent information. For example, if your network was trained on images and data set, you may successfully use its last layer representation in some Kar model classification text. But if you want to use your network in some medical specific text, you probably will do better if you will use an earlier for connected layer or even retrain network from scratch. Here, you may look for a pre-trained model which was trained on data similar to what you have in the exact competition. Sometimes, we can slightly tune network to receive more suitable representations using targets values associated with our images. In general, process of pre-trained model tuning is called fine-tuning. As in the previous example, when we are solving some medical specific task, we can find tune VGG ResNet or any other pre-trained network and specify it to solve these particular texts. Fine-tuning, especially for small data sets, is usually better than training standalone model on descriptors or a training network from scratch. The intuition here is pretty straightforward. On the one hand, fine-tuning is better than training standalone model on descriptors because it allows to tune all networks parameters and thus extract more effective image representations. On the other hand, fine-tuning is better than training network from scratch if we have too little data, or if the text we are solving is similar to the text model was trained on. In this case, model can you use the my knowledge already encoded in networks parameters, which can lead to better results and the faster retraining procedure. Lets discuss the most often scenario of using the fine-tuning on the online stage or the Data Science Game 2016. The task was to classify these laid photos of roofs into one of four categories. As usual, logo was first chosen to the other metric. Competitors had 8,000 different images. In this setting,



it was a good choice to modify some pre-trained network to predict probabilities for these four classes and fine tune it. Let's take a look at VGG-16 architecture because it was trained on the 1000 classes from VGG RestNet, it has output of size 1000. We have only four classes in our text, so we can remove the last layer with size of 1000 and put in its place a new one with size of four. Then, we just retrain our model with very smaller rate is usually about 1000 times lesser than our initial low rate. That is fine-tuning is done, but as we already discussed earlier in this video, we can benefit from using model pre-trained on the similar data set. Image in by itself consist of very different classes from animals to cars from furniture to food could define most suitable pre-trained model. We just could take model trained on places data set with pictures of buildings and houses, fine-tuning this model and further improve their result. If you are interested in details of fine-tuning, you can find information about it in almost every neural networks library namely Keras, PyTorch, Caffe or other. Sometimes, you also want to increase number of training images to train a better network. In that case, image augmentation may be of help. Let's illustrate this concept of image augmentation. On the previous example, we discussed classification of roof images. For simplicity, let's imagine that we now have only four images one for each class. To increase the number of training samples. let's start with rotating images by 180 degrees. Note, that after such rotation, image of class one again belongs to this class because the roof on the new image also has North-South orientation. Easy to see that the same is true for other classes. Great. After doing just one rotation, we already increase the amount of our trained data twice. Now, what will happen if we rotate image from the first class by 90 degrees? What class will it belong to? Yeah, it will belong to the second class and eventually, if you rotate images from the third and the fourth classes by 90 degrees, they will stay in the same class. Look, we just increase the size of our trained set four times although adding such augmentations isn't so effective as adding brand new images to the trained set. This is still very useful and can boost your score significantly. In general case, augmentation of images can include groups, rotations, and the noise and so on. Overall, this reduces over fitting and allows you to train more robust models with better results. One last note about the extracting vectors from images and this note is important one. If you want to fine-tuning convolutional neural network or train it from scratch, you usually will need to use labels from images in the trained set. So be careful with validation here and do not over fit. Well then, let's recall main points we have discussed here. Sometimes, you have a competition with texts or images as additional data. In this case, usually you want to extract the useful features from them to improve your model. When you work with text, pre-processing can prove to be useful. These pre-processing can include all lowercase, stemming, lemmatization, and removing the stopwords. After that pre-processing is done, you can go either Bag of Words or with the Word2vec approach. Bag of Words guarantees you clear interpretation. Each feature are tuned by means of having a huge amount of features one for each unique word. On other side, Word2vec produces relatively small vectors by meaning of each feature value can be hazy. The other

Another advantage of Word2vec that is crucial in competitions is that words with similar meaning will have similar vector representation. Also, Ngrams can be applied to include words interactions for text and TFIDF can be applied to post-process metrics produced by Bag of Words. Now images. For images, we can use pre-trained convolutional neural networks to extract the features. Depending on the similarity between the competition data and the data neural network was trained on, we may want to calculate descriptors from different layers. Often, fine-tuning of neural network can help improve quality of the descriptors. For the purpose of effective fine-tuning, we may want to augment our data. Also, fine-tuning and data augmentation are often used in competitions where we have no other data except images. Besides, there are a number of pre-trained models for convolutional neural networks and Word2vec on the internet. Great. Now, you know how to handle competitions with additional data like text and images. By applying and adapting ideas we have discussed, you will be able to gain an edge in this kind of setting.