

[MUSIC] Hello everyone. In this video we will analyze the Crowdflower competition. We, I mean me, Stanislav Semenov and Dmitry Altukhov participated as a team and took second place. I will explain most important parts of our solution along with some details. The outline of the video is as follows. First, I will tell you about the contest itself, what kind of data and metrics were provided. After that we will discuss our approach, features and tricks. And then I will summarize what is worth to look for on the competition. Some of the competition were organized by CrowdFlower Company. The goal of this competition is to measure relevance of the search results given the queries and resulting product descriptions from living e-commerce sites. The task is to evaluate the accuracy of their search algorithms. The challenge of the competition is to predict the relevance score of a given query and problem description. On the picture we see assessor's user interface, which has a query, a search query, and some information about an item. Assessors were asked to give each query-product pairing a score of 1, 2, 3, or 4, with 4 indicating the item completely satisfied the search query, and 1 indicating the item doesn't match. As the training data, we have only median and variance of these scores. Data set consists of three text fields, request query, result and products title, and product description, and two columns related to the target, median and variance of scores. To ensure that the algorithm is robust enough to handle any HTML snippets in the real world, the data provided in the program description field is raw and sometimes contain permissions that is relevant to the product. For example, a string of text or object IDs. To discourage hand-labeling, the actual set was augmented with extra data. This additional data was ignored to when the public and private leaderboard scores were calculated. And of course, campaign scores have no idea which objects we had already used to calculate the scores. So we have 10,000 samples in train, and 20,000 in test, but it's good data. I mean, validation works well and local scores are close enough to leaderboard scores. Effective solutions, non-standard metric was used, quadratic weighted kappa. Let's take a closer look at it. You can find a detailed description of the metric on the competition evaluation page. We have already discussed the metric in our course, but let me recall it. Submissions are scored based on the quadratic weighted kappa which measures the agreement between two

o ratings. This metric typically will rise from 0, random agreement between raters, to 1, complete agreement between raters. In case there is less agreement between the raters than expected by random, the metric may go below 0. In order to understand the metric, let's consider an example of how to calculate it. First, we need  $N$  by  $n$  confusion matrix  $C$ , which is constructed and normalized. Our vertical axis by its failures, or horizontal predicted failures. In our case,  $N$  is equal to 4 as the number of possible ratings. Second we need  $N$  by  $n$  histogram matrix of expected matrix  $E$ , which is calculated assuming that there is no correlation between ratings cost. This is calculated as within histogram vectors of ratings. Also we need  $N$  by  $N$  matrix of weights,  $W$ , which is calculated based on its elements positions. This particular formula for weights uses square distance between indexes  $i$  and  $j$ , so this is why the method is called quadratic weighted kappa. Finally, kappa can be calculated as one minus a fraction of this weighted sum of confusion matrix in the numerator, and weighted sum of expectation matrix in the denominator. I want to notice that kappa has properties similar both to classification loss and regression loss. The more distant predicted and true ratings are, the more penalties there will be. Remember that thing, we will use it later in our video. Okay, let's now talk about my team solution. It's turned out to be quite complex, consisting of several parts. Like extending of queries, per-query models, bumper features, and so on. I will tell you about main points. So let's start with text features. We have three text fields, a query, a title, and a description. You can apply all techniques that we discuss in our course and calculate various measures of similarity. That's exactly what we did. That is for query title and query description pair, we calculated the number of magic words, cosine distance between TF-IDF representations, distance between the average word2vec vectors, and Levensthein distance. In fact, this is a standard set of features that should be used for similar task. And there is nothing outstanding. The support should be considered as a baseline. And now we turn to the interesting things. In addition, we found it was useful to use symbolic n-grams. You can work with them in the same way as with word-based, if each letter is interpreted as a separate word. We use symbolic n-grams from one letter, to five letters. After getting a large parse metrics of

n-grams, we apply it as to them, and took only close to 300 combinations as features. You remember we discussed this portion of our course, there is an example. Looking at the data we notice three interesting properties. Queries are very short, numbers of unique queries is 261, and queries are the same in train and test sets. We decided to use these observations to build extended versions of query as follows. For each query, we get the most relevant corresponding items, those with relevance equal to four. We join all words from the title of the relevant items and take ten most popular words. This is what we called extended query and it's used to build all these text features that I mentioned earlier. Notice that this trick is applicable only within the context framework. Due to the fact that queries in test are exactly the same as in train. In real life we couldn't do so because it's unrealistic to ignore relevant results for every search query. The fact that sets of queries in train and test are the same, give us an opportunity to split our problem into many small subtasks. Specifically, for each query, you can build a separate model that will only predict relevance of corresponding items. Again, in real life, such tricks can't be applied, but within the context framework, it's totally fine. So, for each unique query, we get corresponding samples, calculate word to word similarities, back and forth presentation, and fit a random fourth classifier. Finally, we have 261 model which predictions were used as a feature in final example. Do remember that every product item is by several people. Median in a variance of the ratings. The variance show in ratings are inconsistent or not. Intuitively, if the variance is large, then such an object need to be taken into account with a smaller weight. One way to do it is to assign items weight depending on query answers, which was a heuristics  $\frac{1}{1 + \text{variance}}$  as the weight. Another method is to restore individual observation using median and variance statistics. We found that supposing there are three assessors, we can almost always certainly restore our original labels. For example, if we have a median equal to three, and variance equal to 0.66, there of course are two, three, and four, which by this approach and for each source sample, generated three once. However, using them as training data took longer to train, and did not improve the score. And simple heuristic works quite well, and they use it in final solution. In the competition, you need

to choose a metric, we need to predict class, but penalty for the error is not symmetric. We decided to take it into account, adding several artificially created binary delimiters between classes as features. In other words, we're trying to classify to answer the question, is it true that target class number is greater than 1, greater than 2, and so on. We call these features bumpers, since they are kind of separators between classes. We build them in fashion, similar to how we construct predictions instead. It was very useful for the final solution. All mentioned features will be used in an ensemble. We build several models on different subsets of features. Some feel relatively simple like SVM and some look quite complex. You can see the part of complex model created by me. It uses bumper features, all sorts of similarities and query features in different combinations. Also there is a lot of frostage models, which are mixed up with second stage random forest. In fact, each participant of the team made his own model, and then for competition we simply mixed our model linearly for final submission. Let's remember that the metric on the one hand has some proper classification It's necessary to predict the class. But for regression answer, we can analyze more. We have built models for regression, but we have had to somehow turn real-valued predictions into classes. A simple approach with would work poorly, so we decided to pick up thresholds. For the purpose, we were right. Thresholds and choose the best one weighted on cross-validation score. The buff procedure gave us a very significant increase in quality. in fact it often happens in competitions with non-standard metrics that [INAUDIBLE] grades symmetric optimization gives a significant improvement. So let's sum up. In the competition it was really important to use the [INAUDIBLE] ideas. First symbolic and grams, once since they give a significant increase in the score and it was not solved that you should use that. Second, expansion of queries led to significant increase in this course. Also optimization of thresholds was a crucial part of our solution. I hope you will re-use some of these tricks in your competitions, though. Thank you for the attention. [MUSIC][MUSIC] Hi, throughout the course, we use the Springleaf competition as a useful example of EDA, mean encodings and

features based on nearest neighbors. Back then, we took the third place in this competition together with and. And now in this video, I will describe the last part of our solution, which is the usage of stacking and ensembles. On this picture, you can see the final stacking scheme we produced on the level 0 features, on the first level, predictions by basic models. On the level one plus combination. So these predictions and some accurately chosen features on the second level models on this new set of features. And finally, on the third level, their linear combination. In this video, we will go through each level as it builds up to this non-trivial ensembled scheme. But first, let's quickly remind ourselves about the problem. This was a binary classification task with area under curve metric. We had 145,000 samples in training data and about 2,000 anonymized features. These were useful insights derived by us while doing EDA. And you can check out EDA done by earlier in our course to refresh your memory. So now let's start with features. Here we have four packs of features. First two are the basic dataset and the processed dataset. To keep it simple, we just used insights derived from EDA to clean data [INAUDIBLE] and to generate new features. For example, we remove duplicated features and edit some feature interaction based on scatter plots and correlations. Then, we mean-encoded all categorical features using growth relation loop and sign data and smoothing. We further used the mean-encoded dataset to calculate features based on nearest neighbors. Like, what is the least in closest object of the class zero? And how many objects out of ten nearest neighbors belong to class one? You can review how this could be done in related topics introduced by Dmitri Altihof. So finally, these four packs of feature were level 0 of our solution. And the second level was represented by several different gradient within decision tree models, and one neural network. The main idea here is that meta features should be diverse. Each meta feature should bring new information about the target. So we use both distinct parameters and different sets of features for our models. For the neural network, we additionally pre-processed features with common scalars, ranks and power transformation. The problem there was in huge outliers which skew network training results. So ranks and power transformation

helped to handle this problem. After producing meta features who is gradual in boosting decision to it and neural networks, we calculated pay rise differences on them to help next level models. Note that this is also an interesting trick to force the model to utilize the differences in the first level models predictions. Here we edit two datasets of features based on nearest neighbors. One was taken directly from level 0 and they contain the same features. But it was calculated on the mean-encoded dataset to the power of one-half. The point here was that these features were not completely utilized by the first level models. And indeed, they brought new pieces of information to this level. Now we already have autofold predictions from the first level and we will train with the models on them. Because we could have target leakage here because of other folk, and also because features not very good and there are almost no patterns left in the data for models to discover. We chose simple classifiers, keeping in mind that predictions should be diverse. We used four different models. Gradient boosted decision tree, neural networks, random forest and logistic regression. So this is all with the second level models. And finally, we took a linear in your combination of the second level models. Because a linear model is not inclined to that we estimated coefficients directly using these four predictions and our target for throwing in data. So, this is it. We just went through each level of this stacking scheme and then the student. Why we need this kind of complexity? Well, usually it's because different models utilize different patterns in the data and we want to unite all of this patterns in one mighty model. And stacking can do exactly that for us. This may seem too complicated. Of course, it takes time to move up to this kind of scheme in a competition. But be sure that after completion our course, you already have enough knowledge about how to do this. These schemes never appear in the final shape at the beginning of the competition. Most work here usually is done on the first level. So you try to create diverse meta features and unite them in one simple model. Usually, you start to create the high grade second level of stacking, when you have only a few days left. And after that, you mostly work on the improvement of this scheme. That said, you already have the required knowledge and now you just need to get some practice out there. Be diligent, and without a doubt,

you will succeed. [SOUND] [MUSIC][MUSIC] Hi, in this video, I will tell you about Microsoft Malware Classification Challenge. We were participating in a team with other lecturers of this course, Mikhail Trofimov and Stanislav Semenov. We got the third place in this competition. The plan for the presentation is the following. We will start with the data description and a little bit of EGA. We'll then talk about feature extraction, then we will discuss how we did feature processing and selection. We will see how we did the modeling, and finally we will explain several tricks that allowed us to get higher on the league of board. So let's start with problem description. In this competition, we were provided about half of terabyte of malicious software executables. Each executable was represented in two forms. The first is, HEX dump. HEX dump is just a representation of file as a sequence of bytes, it is row format. The file on the disk is stored as a sequence of bytes, that is exactly what we see in HEX dump. The second form is a listing generated by interactive disassembler, or IDA in short. IDA tries to convert low level sequence of bytes from the HEX dump to a sequence of assembler columns. Take a look at the bottom screenshot. On the left, we see sequences of bytes from HEX dump. And on the right, we see what IDA has converted them to. The task was to classify malware files into nine families. We will provide with train sets of about 10,000 examples with known labels, and a testing set of singular size. And the evaluation metric was multi-class logarithmic loss. Note that the final loss in this competition was very, very low. The corresponding accuracy was more than 99.7%, it's enormously high accuracy. Remember, it is sometimes beneficial to examine metadata. In this competition, we were given seven zip archives with files. And the filenames look like hash values. Actually, we did not find any helpful metadata for our models, but it was interesting how train test split was done by organizers. The organizers were using the first letter of the file names for the split. On this plot, on the x axis, we see the first characters of the file names. And on the y axis, we have the number of files with their names, starting with the given character. The plots actually look cool, but the only information we get from them is the train test split is in fact random, and we cannot use them to improve our models. So in this competi

tion we were given a raw data, so we needed to extract the features ourselves. Let's see what features we extracted. Well, in fact, no one from our team had domain knowledge or previous experience in malware classification. We did not even know what IDA disassembly is. So we started really simple. Remember issue executable in the dataset was represented as two files. So our first two features were the sizes of those files, or equivalently their length. And surprisingly we got 88% accuracy just by using these two features. On the plot you see x axis correspond to an index of an object in the train set. We actually sorted the objects by their class here. And the y axis shows the file sizes of a HEX dump file for every object. And we see this feature is quite demonstrative. The most simple features we could derive out of sequence are account of their elements, right? So this is basically what function value comes from pandas does. So it is what we did. We counted bytes in HEX dump files, and that is how we get 257 features. And 257 is because we have 256 byte values, and there was one special symbol. We achieved almost 99% accuracy using those features. At the same time that we started to read about this assembly format, and papers on malware classification. And so, we got an idea what feature is to generate. We looked into this assembly file. We use regular expressions to extract various things. Many of the features we extracted were thrown away immediately, some were really good. And what we did actually first, we counted system calls. You see that on the slide, they are also called API calls in Windows as we read in the paper. And here is the error rate we got with this feature, it's pretty good. We counted assembler commons like move, push, goal, and assembler, they work rather well. We also try to extract common sequences like common end grams and extract features out of them, but we didn't manage to get a good result. The best features we had were section count. Just count the number of lines in this assembly file, which start with .text or .data, or something like that. The classification accuracy with that feature was more than 99% using these features. By incorporating all of these, we were able to get an error rate less than 0.5%. From our text mining experience, we knew that n-grams could be a good discriminative feature. So we computed big grams. We found a paper which proposed to use 4-grams. We computed them too, and we even went further and extracted 10-grams. Of course, we couldn't work



with 10 grams directly, so we performed a nice feature selection, which was one of the most creative ideas for this competition. We will see later in this video how exactly we did it. Interestingly, just by applying the feature selection to 10-grams and fitting in XGBoost, we could get a place in the top 30. Another feature we found interesting is an entropy. We computed entropy of small segments of wide sequence by moving the sliding window over the sequence. And computing entropy inside each window. So we've got another sequence that could contain an explicit information about the encrypted parts of the executable. See, we expect the entropy to be high if the data is encrypted and compressed, and low if the data is structured. So the motivation is that some malwares are injected into a normal executables, and they are stored in encrypted format. When the program starts, the malware extracts itself in background first and then executes. So entropy features could kind of help us to detect and encrypt the trojans in the executables, and thus, detect some classes of malware. But we got an entropy sequence of variable length. We couldn't use those sequences as features, right? So we generated some statistics of entropy distribution like mean and median. And also we computed a 20 percentiles and inverse percentiles, and use them as features. The same features were extracted out of entropy changes, that is we first apply diff function to entropy sequence and then compute the statistics. It was possible to extract three things from hex dump by looking for printable sequences that ends with 0 element. We didn't use strings themselves, but we computed strings lengths. Distribution for each file and extract the similar statistics, the statistics that we extracted from entropy sequences. Okay, we finished with feature extraction. So let's see how we pre-process them and perform feature selection. Those moment when we generated a lot of features. We wanted to incorporate all of them in the classifier, but we could not fit the classifier efficiently when we got say 20,000 features. Most features will probably useless, so we tried different feature selection method and transformation techniques. Let's consider the built-in features. There are 257 of them, not much, but probably there is still a redundancy. We tried non-negative matrix factorization, NMF. And the principal component analysis, PCA, in order to remove this redundancy. I will remind you that both NMF and PCA are trying to factorize object feature matrix  $x$  into

the product of two low-rank matrices. You can think of that of a finding a small number of basis vectors in the feature space, so that every object can be approximated by a linear combination of those basis vectors. And this coefficients of approximation can be treated as new features for each object. So the only difference between NMF and PCA is that NMF requires all the components of floor rank matrices to be non-negative. We set the number of basis vectors to 15, and here we see plots between the first two extracted features. One for NMF and one for PCA. So these are the coefficient for most important basis vectors actually. We used 3D based model and it is obvious that NMF features were a lot better for trees. So NMF works good when the non-negativity of the data is essential. And in our case we worked with counts, which are non-negative by nature. Simple trick to get another pack of features by doing almost nothing is to apply a log transform to the data and calculate NMF on the transformed data again. Let's think what happened here. NMF protest originally uses MSE laws to measure the quality of approximation it build. This trick implicitly changes the laws NMF optimizes from MSE to RMSLE. Just recall that RMSLE is MSE in the log space. So now the decomposition process pays attention to different things due to different loss, and thus produces different features. We used the small pipeline to select 4-grams features. We removed rare features, applied linear SVM to the L1 regularization, as such model tends to select features. And after that, we thresholded random forest feature importances to get final feature set of only 131 feature. The pipeline was a little bit more complicated with 10-grams. First, we used the hashing to reduce the dimensionality of original features. We actually did it online while computing 10-grams. We then selected about 800 features based on L1 regularize SVM and the RandomForest importances. This is how we got about 800 features instead of 2 to the power of 28. But we went even further. This is actually the most interesting part for me. The main problem with feature selection, that I've just described, is that we've done it for 10-grams independently of other features that we already had to that moment. After selection, we could end up with really good features. But those features could contain the same information that we already had in other features. And we actually wanted to improve our features with 10-grams. So here is what we did instead. We generated out of fold prediction for the train set using all the features that we had. We sorted the object by their true cla

ss  
predicted probability and try to find the features that would separate the most error prone object from the others. So actually, we use another model for it. We've created a new data set with error prone objects having label 1, and others having label 0. We trained random forest and selected 14 10-grams, well, actually hashes to be precise. We had a nice performance increase on the leaderboard when incorporating those 14 features. This method actually could lead us to severe overfitting, but it fortunately worked out nicely. All right, let's get to modeling. We didn't use stacking in this competition as we usually do now. It became popular slightly after this competition. We first used Random Forest everywhere, but it turned out it needs to be calibrated for log loss. So we switched to XGBoost and used it for all our modeling. Every person in the team extracted his own features and trained his own models with his own parameters. So our final solution was a combination of diverse models. We found bagging worked quite well in this data set. We even did the bagging in a very peculiar way. We concatenated our twin set with a seven times larger set of objects sampled from train set with replacement. And we used the resulting data set that is eight times larger than original train set, to train XGBoost. Of course, we averaged about 20 models to account for randomness. And finally, let's talk about several more tricks we used in this competition. The accuracy of the models was quite high, and we all know that the more data we have the better. So we've decided to try to use testing data for training our models. We just need some labels for the testers, right? We either use predicted class or we sample the label from the predicted class distribution. Generated test set labels are usually called pseudo labels. So how do we perform cross validation with pseudo labels? We split our train set into folds as we usually do when performing cross validation. In this example, we split data in two folds. But what's different in cross validation is that now, before training the model in a particular fold, we can concatenate this fold with the test data. Then we switch the faults and do the same thing. See we trained the objects to be denoted with green color and predict the once shown with red. Okay, and to get predictions for the test set, we again do kind of cross validation thing, like on the previous slide. But now we divide the test set in faults and concatenate train set to each fold of the test set. In the end, we get out of all predictions

for the test set, that is our submission. One of the crucial things to understand about this method, that sometimes we need two different models for it. And this is the case where one model is very confident in its predictions. That is, the predictive probabilities are very close to 0 and 1. In this case, if we train a model, predict task set, sample labels or take it the most probable label and retrain same model with the same features or get no improvement at all. We just did not introduce any information with pseudo labeling in this way, but if I ask say, Stanislav, to predict that with his model. And then I use his labels for a test and create my model, that is where I will get a nice improvement. This actually becomes just another way to incorporate knowledge from two diverse models. And the last thing that helped us a lot is per-class mixing. At the time of the competition, people usually mixed models linearly. We went further and mixed models joining coefficients for each class separately. In fact, if you think about it, it's very similar to stacking with a linear model of a special kind at a second level. But the second became popular in a month after this competition, and what we did here is a very simplest manual form of stacking. We published our source code online, so if you are interested you can check it out. All right, this was an interesting competition. It was challenging from technical view point as we needed to manipulate more than half of terabyte of data, but it was very interesting. The data was given in a raw format and the actual challenge was to come up with nice features. And that is where we could be creative, thank you. [MUSIC] Hi. In this video, I'm going to talk about Walmart trip type classification challenge which was held in Kaggle couple of years ago. I won the first place in that competition. And now, I will tell you about most interesting parts of the problem and about my solution. That said, this presentation consists of four parts. First, we will state the problem. Second, we will understand what data format and data reprocessing. And third, we will talk about models, their relative quality and their relation to the general stacking scheme. And finally, we will overview some possibilities to generate new features here. So, let's start. In our data, we had purchases people made in Walmart visiting their shop in two weeks, and we had to classify them into 38 visiting trip types or classes. Let's take a quick look at features in the data. Trip type column represents the target, visit number represents ID which unites purchases made by one customer in one shopping trip. For example, a customer which made visit number seven, purchased two items which are located in the second and in the third lines of this data frame. Notice that all rows with the same visit number have the same trip type. An importa

nt moment, is that we have to predict a trip type for visit number, and not for each row in the train data. And as you can see, in the train we have around 647,000 rows and only 95,000 visits.

Back to the features, next feature is weekday which obviously represents the weekday of the visit. Next is UPC. UPC is an exact ID of a purchased item. Then, scan count. Scan count is the exact number of items purchased. Note that minus one here represents not the purchase but the return. Next features, department description, with 68 unique values is a broad category for an item.

And finally, fineline number, with around 5000 unique values, is a more refined category for an item. So, after we understood what this feature represents, let's recall that we have to make one prediction for each visit number. Let's take a look at the data for the visit number eight. We can see here that this particular visit has a lot of purchases in category paint and accessories, which means that trip type number 26 may represent a visit with most purchases in that category. Now, how to approach model train in here. Let's take another look at the data and assess our possibilities. Should we predict trip type for each item on the list or should we choose another way? Of course both of them are possible, but in the first one, we'll predict trip type for each row with each data set, we'll miss important interactions between items which belong to the same visit. For example, trip type may have a number of 26, if more than half of its items are from paint and accessories. But, if we will not account for interaction between these items, it can be quite hard to predict. So,

the second option of uniting all purchase in the visit and making a data set where each row represents a complete visit, seems more reasonable. And, as can be expected, this approach leads to more significant benefits in the competition. I'm going to show

you the easiest way to change the data format to the desired one. Let's choose the department description feature for the purpose of an example. First, let's group the data frame by visit number and calculate how many times each department description is present in a visit. Then, let's unstack last group by column so we will get a unique column for each department description value. Now, this is the format we wanted. Each row represents a visit and each column is a feature described in that visit. We can use this group by approach for other features besides department description. Also note that items in the visit are actually very

similar to words in a text. After our confirmation, each feature here represents counts, so we could apply ideas which usually works with text, for example, tf-idf transformation. As you can guess, a lot of possibilities emerge here. Great. After this is done and we process data in the desired format, let's move to choosing a model. Based on what we already have discussed, can you guess if we should expect the significant difference in scores between linear models and tree-based models here? Think about this a bit. For example, is there a reason why linear models will under perform in comparison to tree based-models? Yes, there is.

Again, I'm talking about interactions here. Indeed, tree-based models in neural network have significant superiority in quality in this competition for this very reason. But still, one can use linear models and TNN to produce useful method features here. Despite the fact that they didn't imply interactions, they were a valuable asset in my general staking scheme. I will not go int

o further details of staking here because we already covered most ideas in other videos about competitions. Instead, we'll talk a bit about feature generation. Except for interactions between items purchased in one visit, one could try to exploit interactions between features. The interesting and unexpected result here was that one fineline number can belong to multiple department descriptions, which means that fineline number is not a more detailed department description as you can think. Using this interaction, one can further improve his model. Another interesting feature generation idea was connected to the time structure of the data. Take a look at this plot, it represents the change in the weekday feature relative to the row number. It looks like the data is ordered by time here. And the data appears to consist of 31 days, but train test split wasn't time based. So, you could derive features like day number in the data set, number of a visit in a day, and the total amount of visits in a day. So, this is it. We just discussed the most interesting parts of this competition. Changing the data format to a more suitable, generating features while doing so, working with models while doing stacking. And finally, doing some for additional feature engineering. The challenge itself proved useful and interesting. And I would recommend you to check it out and try approaches we have talked about.

Hello everyone. Today, I will explain to you how, me and my team mate Gert, we won a very special Kaggle competition called the Acquired Valued Shoppers Challenge. First let me give you some background about the competition. It was a recommender's challenge. And when I say a recommender's challenge, I mean you have a list of products and a list of customers, and you try to bring them together. So we target them so we recommend which customer in order to increase sales loyalty or something else. There were around 1,000 teams and for back then, at least they were quite a lot. Now Kaggle has become much more popular. But back then, given the size of the data, which was quite big, I think this competition attracted a lot of attention. And as I said, we attained first place with my teammate Gert, for what it was, I think, a very clear win because we took a lead early in the competition, and we maintained it. And that solution was not actually very machine learning Kebbi, but it was focused on really trying to understand the problem well and find ways to validate properly. And in general, it was very, very focused on getting sensible results. And that's why I think it's really valuable to explain what we did. So what was the actual problem we tried to solve? I imagine you have 310,000 shoppers, almost equally split. I'm on a train and test. You have all their transactions from a point where they were given an offer, and these were about 350 million transactions. So, one year of all the transactions of all the customers involved in this challenge, and you have 37 different offers. When I say offer here, it is actually a coupon. So, a shopper is sent normally. It's a piece of paper that says, "If you buy this item, you can get a discount." So it recommends to certain item. Now, we don't know exactly what discount is. Maybe discounts, maybe it says, "You can buy another item for free." So, maybe a different promotion, but in principle is some sort of incentive for you to buy this item. I have mentioned items so far or products, but the notion of product was not actually present in this challenge, but we could infer it. We could say that a unique

the combination of a brand, category, and company that were present could form a product, at least for this competition. Let me give you a bit more details about the actual problem we are trying to solve. Imagine you have a timeline, starts from one year in the past until one year after. So, in the past, a customer makes a visit to the shop, buys certain items, leaves, comes back another day, makes another visit, buys more items. Then at some point, he or she is targeted with an offer, a coupon as I said. All transactional history for this customer stops there. You don't know anything else. Up to this point, you know everything for one year back up to this. The only thing you know is, you know that the customer has indeed redeemed that coupon. So he did buy the offer product. And for that training data only, you also have a special flag that tells you whether he or she bought that item again. And you can see why this is valuable. Because normally, when you target someone with a coupon, you give a discount, and you don't make that much profit actually, but you aim in establishing a long-term relationship with the customer. And that's why they were really interested in getting a model here that could predict which recommendation, which offer will have that effect, will make a customer develop a habit in buying this item. And what we were being tested on was AUC. By this point, I expect you roughly know what AUC is as a metric, but in principle, you're interested in how well your score discriminates between those that bought or can buy the item again and those that will not. So, when you have the highest score, you expect higher chance to buy the item, and a lower score, lower attempts to buy the item again. So, higher AUC means that this discrimination is stronger with your score. This was a challenging challenge. And it was challenging because, first of all, the datasets were quite big. As I said, 350 million transactions. Me and Gert, we didn't have crazy resources back then. I have to admit that I have personally improved my hardware since then, but actually back then, I was working only with a laptop that had 32-bit Windows and only four gigabytes of RAM. So, really small and mainly challenging that we had to deal with these client files. And then, we didn't have features. So, what we knew is this customer was given this offer, which is consistent by these three elements I mentioned before, category, brand, and company, and the time that this recommendation was made nothing else. Then, you had to go to the transactional history and try to generate features. And you know, anybody could create really anything they wanted. There was not a clear answer about which features would work best. So this is not your typical challenge where you're normally given the thesis. But it is quite difficult for the type of the recommender's challenge. And what really makes this competition difficult, interesting, and what I think at the end of the day gave us the win was the fact that the testing environment was very irregular. And we can define irregular, in this context, as an environment where the train data and the test data had different customers. So, no overlaps. Different customers, and one different in the other. Also, the training this data had in general different offers. It was showing you a graph that shows that the distribution of its offer and whether it appears in the train or in the test data or both. And you can see that most offers, either appear only in test or they appear only in train with minimal overlap. So, that

makes it a bit difficult because you basically have to make a model with soft products. They were offering the train, but in the test data, you have completely other offers. So you don't know how they would behave as these products have never been offered before. And the last element is, the test data is obviously in the future. That is expected. But given the other elements, this makes it more difficult, especially in some cases were well in the future. And some of it is not as important elements, but still crucial was that this challenge was focusing on acquisition. So, there is not that much history between the customer and the offered product. And I say this is irregular because grocery sales are in principle based on what the customer already like and has bought many times in the past. So we referred to these type of acquisition problem, where we don't have much history, as the cold start problem, and it is much more challenging because you don't have that direct link. That's, the customer really like this product I made an offer because we don't have a past history that can verify this or we don't have much history. And the last element is that if you actually see the propensity of an offer to be bought, again in the training data, the results were quite different. And here, I give you the offer by shortened propensity, and you can see some offers had much success to be bought again. It's like offer two that somehow this had much less. For example, 20 percent, and this is just a sample. There were some other offers that had around five percent. So, if you put now everything into the context, you have different customer, different offers, different buyer, different time periods. In principle, you don't have that much information about the customer and the offer product, and you know that the offers of the training data are actually quite different. It's really difficult to get a standard pattern here. And you know that the offers in the test data are going to be different. So, all this made it a difficult problem to solve or in irregular environment. How did we handle big data? We did it with indexing. And the way I did the indexing was, I saw that the data were already sorted by customer and time. So, I passed through these big data file of transactions, and every time I encountered a new customer, I created a new file. So I created a different file for each customer that contained all his or her transactions, and that made it really easy to generate features, because if I have to generate features for a specific customer, I would just access this file and create all the features I wanted at the will. This is also very scalable. So I could create threads to do this in parallel. So, access many customers in parallel. And I did this not only for every customer, but also for every category, brand, and company. So, every time I wanted to access information, I would just access the right category, the right brand, or the right customer, and I will get the information I wanted, and that made it very quick to handle all these big chunks of data. But what I think was the most crucial thing is how we handle this irregularity. I think at the end of the day, this is what determines our victory because once we got this right and we were able to try all sorts of things and we had the confidence that it will work in the test data. The first thing we tried to do and this is something that I want you to really understand, is how we can replicate internally what we are being tested on. That's really important. I'll give you the



room to try all these things. Try all different permutations and combinations of data techniques, anything you have you can put in mind, and really understand what works and what's not. So, we tried to do that. The first of attempt didn't go very well. So we try to randomly split the data between train and validation and we're trying to make certain that each offer is represented equally in each one of this train and validation data set proportionately equally. But was that correct? I mean, if you think about it. What we were doing there, we were saying I'm building a model with some offers and I'm validating in the same offers. That's good. Maybe we can do well here. But is this what we're really being tested on? No. Because in the test data, we'll have completely different offers. So, this didn't work very well. This was giving very nice internal results but not very good results in the test data. So, we tried something else. Can we leave one offer out? And I'm showing you roughly what this look like. So, for every offer, can we put an offer in the validation data and use all the cases of every other offer to train a model? So, if we were to predict offer 16, we will use all customers that received offer 1 to 15 and 17 to 24 to build the model and then we'll make predictions for all those customers that received offer 16. And you can see that this actually is quite close to what you're being tested on because you know you're building a model with some offers but, you're being tested on some other offer that is not there. And you can take the average of all these 24 users and I put 24 because this is how many offers you really have in the training data. You can take that average and that average may be much more close to the reality, close to what you were being tested on. And this was true. This gave better results, but we were still not there. And I'll show you why we were not there. Consider the following problem. Here, I'll give you a small sample of predictions for offer two and what was the actual target? What we see here is a perfect AUC score. Why? Because all our positive cases that are labeled with one and they have the green color, have higher score than the red ones, where the target is zero. So, the discrimination here is perfect. We have a point, a cutoff point. We can set 0.5 here where all cases that have score higher than this. We can safely say they are one and that is true and everything that has a score lower than this are zero. So, you see here one discrimination is perfect. Let's now take a sample from offer four. If you remember offer four, had in general lower propensity. Offer two had around 0.5 and offer four had around 0.2. So, it's mean we're center much lower and what you can see here is that, again, AUC is perfect for this sample because again, all the higher scores that are labeled with green have a target of one. And then the lower scores, everything that has a score less than 0.18 has a target of 0. The discrimination is perfect. We can find this cutoff point. We can say 0.8, where everything that has a score higher than this can safely be set to one. And that is always true. And vice versa, everything that's less than 18, then it's a 0. And that is always true. So, we have two scores. They discriminate really well between the good and the bad cases. However, we are not tested on the AUC of one offer. We are tested on the AUC of all offers together. So the test data have many offers. So, you are interested in the score that generalizes well against any offer. So, what happens if we try

to merge this table? AUC is no longer perfect and why this happens? Because some of the negative cases of the first offer had higher score than the positive cases, those that have a target equal to one from the second offer. So you can see, although the discrimination internally is really good, they don't blend that well. You lose something from that ability of your score to discriminate between ones and zeros. And the moment we saw this, we knew that just leaving one offer out was not going to be enough. We had to make certain that when we merge all those scores together, the score is still good. The ability of our model to discriminate is still powerful or it doesn't lose. And that's why we use a combination of the previous average AUC of all the offers and the AUC after doing this concatenation. So, the average of the two AUCs which really the metric we try to optimize because we thought that this is actually very close to what we were being tested on. And here I can show you the result of all our attempts and this is with a small subset of features because by that point, we were not interested to create the best features, we were interested to test which approach works best. So, you can see if you do it standard stratified K-fold, you can get much nicer results in internal cross-validation but in the test data, the relationship is almost opposite. So, highest score in cross-validation leads to worse results in the test data. And you can see why because you're not internally modeling or internally validating or on what you are actually being tested on. Doing the one-offer out keep obviously lower internal cross-validation results and better performance in the test data, but even better was doing this leave-one-offer plus one concatenation in the end. And this AUC was lower but actually had better test results. I believe we could get even better results if we made certain that we are also validating in new customers. But we didn't actually do this because we saw that this approach had already good results. But as a means to improve, we could have also made certain that we validate on different customers because this is what the test was like.