

```

import theano, theano.tensor as T
import numpy as np
import theano_lstm

from out_to_in_op import OutputFormToInputFormOp

from theano_lstm import Embedding, LSTM, RNN, StackedCells,
Layer, create_optimization_updates, masked_loss, MultiDropout

def has_hidden(layer):
    """
    Whether a layer has a trainable
    initial hidden state.
    """
    return hasattr(layer, 'initial_hidden_state')

def matrixify(vector, n):
    # Cast n to int32 if necessary to prevent error on 32 bit
    systems
    return T.repeat(T.shape_padleft(vector),
                    n if
    (theano.configdefaults.local_bitwidth() == 64) else
    T.cast(n, 'int32'),
            axis=0)

def initial_state(layer, dimensions = None):
    """
    Initalizes the recurrence relation with an initial hidden
    state
    if needed, else replaces with a "None" to tell Theano that
    the network will return something, but it does not
    need
    to send it to the next step of the recurrence
    """
    if dimensions is None:
        return layer.initial_hidden_state if
has_hidden(layer) else None
    else:
        return matrixify(layer.initial_hidden_state,
dimensions) if has_hidden(layer) else None

def initial_state_with_taps(layer, dimensions = None):
    """Optionally wrap tensor variable into a dict with
taps=[-1]"""
    state = initial_state(layer, dimensions)
    if state is not None:
        return dict(initial=state, taps=[-1])
    else:
        return None

```

```
class PassthroughLayer(Layer):
    """
    Empty "layer" used to get the final output of the LSTM
    """

    def __init__(self):
        self.is_recursive = False

    def create_variables(self):
        pass

    def activate(self, x):
        return x

    @property
    def params(self):
        return []

    @params.setter
    def params(self, param_list):
        pass

def get_last_layer(result):
    if isinstance(result, list):
        return result[-1]
    else:
        return result

def ensure_list(result):
    if isinstance(result, list):
        return result
    else:
        return [result]

class Model(object):

    def __init__(self, t_layer_sizes, p_layer_sizes,
                 dropout=0):

        self.t_layer_sizes = t_layer_sizes
        self.p_layer_sizes = p_layer_sizes

        # From our architecture definition, size of the
        notewise input
        self.t_input_size = 80

        # time network maps from notewise input size to
        various hidden sizes
```

```

        self.time_model = StackedCells( self.t_input_size,
celltype=LSTM, layers = t_layer_sizes)
        self.time_model.layers.append(PassthroughLayer())

        # pitch network takes last layer of time model and
state of last note, moving upward
        # and eventually ends with a two-element sigmoid layer
        p_input_size = t_layer_sizes[-1] + 2
        self.pitch_model = StackedCells( p_input_size,
celltype=LSTM, layers = p_layer_sizes)

self.pitch_model.layers.append(Layer(p_layer_sizes[-1], 2,
activation = T.nnet.sigmoid))

        self.dropout = dropout

        self.conservativity = T.fscalar()
        self.srng =
T.shared_randomstreams.RandomStreams(np.random.randint(0,
1024))

        self.setup_train()
        self.setup_predict()
        self.setup_slow_walk()

    @property
    def params(self):
        return self.time_model.params +
self.pitch_model.params

    @params.setter
    def params(self, param_list):
        ntimeparams = len(self.time_model.params)
        self.time_model.params = param_list[:ntimeparams]
        self.pitch_model.params = param_list[ntimeparams:]

    @property
    def learned_config(self):
        return [self.time_model.params,
self.pitch_model.params, [l.initial_hidden_state for mod in
(self.time_model, self.pitch_model) for l in mod.layers if
has_hidden(l)]]

    @learned_config.setter
    def learned_config(self, learned_list):
        self.time_model.params = learned_list[0]
        self.pitch_model.params = learned_list[1]
        for l, val in zip((l for mod in (self.time_model,
self.pitch_model) for l in mod.layers if has_hidden(l)),
learned_list[2]):

```

```

        l.initial_hidden_state.set_value(val.get_value())

    def setup_train(self):

        # dimensions: (batch, time, notes, input_data) with
        input_data as in architecture
        self.input_mat = T.btensor4()
        # dimensions: (batch, time, notes, onOrArtic) with
        0:on, 1:artic
        self.output_mat = T.btensor4()

        self.epsilon = np.spacing(np.float32(1.0))

        def step_time(in_data, *other):
            other = list(other)
            split = -len(self.t_layer_sizes) if self.dropout
            else len(other)
            hiddens = other[:split]
            masks = [None] + other[split:] if self.dropout
            else []
            new_states = self.time_model.forward(in_data,
            prev_hiddens=hiddens, dropout=masks)
            return new_states

        def step_note(in_data, *other):
            other = list(other)
            split = -len(self.p_layer_sizes) if self.dropout
            else len(other)
            hiddens = other[:split]
            masks = [None] + other[split:] if self.dropout
            else []
            new_states = self.pitch_model.forward(in_data,
            prev_hiddens=hiddens, dropout=masks)
            return new_states

        # We generate an output for each input, so it doesn't
        make sense to use the last output as an input.
        # Note that we assume the sentinel start value is
        already present
        # TEMP CHANGE: NO SENTINEL
        input_slice = self.input_mat[:,0:-1]
        n_batch, n_time, n_note, n_ipn = input_slice.shape

        # time_inputs is a matrix (time, batch/note,
        input_per_note)
        time_inputs =
        input_slice.transpose((1,0,2,3)).reshape((n_time,n_batch*n_note,n_ipn))
        num_time_parallel = time_inputs.shape[1]

        # apply dropout

```

```

        if self.dropout > 0:
            time_masks =
theano_lstm.MultiDropout( [(num_time_parallel, shape) for
shape in self.t_layer_sizes], self.dropout)
        else:
            time_masks = []

        time_outputs_info = [initial_state_with_taps(layer,
num_time_parallel) for layer in self.time_model.layers]
        time_result, _ = theano.scan(fn=step_time,
sequences=[time_inputs], non_sequences=time_masks,
outputs_info=time_outputs_info)

        self.time_thoughts = time_result

        # Now time_result is a list of matrix [layer](time,
batch/note, hidden_states) for each layer but we only care
about
        # the hidden state of the last layer.
        # Transpose to be (note, batch/time, hidden_states)
        last_layer = get_last_layer(time_result)
        n_hidden = last_layer.shape[2]
        time_final =
get_last_layer(time_result).reshape((n_time,n_batch,n_note,n_hidden)).tra

        # note_choices_inputs represents the last chosen
note. Starts with [0,0], doesn't include last note.
        # In (note, batch/time, 2) format
        # Shape of start is thus (1, N, 2), concatenated with
all but last element of output_mat transformed to (x, N, 2)
        start_note_values =
T.alloc(np.array(0,dtype=np.int8), 1, time_final.shape[1], 2 )
        correct_choices = self.output_mat[:,
1:,-1,:].transpose((2,0,1,3)).reshape((n_note-1,n_batch*n_time,
2))
        note_choices_inputs =
T.concatenate([start_note_values, correct_choices], axis=0)

        # Together, this and the output from the last LSTM
goes to the new LSTM, but rotated, so that the batches in
        # one direction are the steps in the other, and vice
versa.
        note_inputs = T.concatenate( [time_final,
note_choices_inputs], axis=2 )
        num_timebatch = note_inputs.shape[1]

        # apply dropout
        if self.dropout > 0:
            pitch_masks =
theano_lstm.MultiDropout( [(num_timebatch, shape) for shape

```

```

in self.p_layer_sizes], self.dropout)
    else:
        pitch_masks = []

        note_outputs_info = [initial_state_with_taps(layer,
num_timebatch) for layer in self.pitch_model.layers]
        note_result, _ = theano.scan(fn=step_note,
sequences=[note_inputs], non_sequences=pitch_masks,
outputs_info=note_outputs_info)

        self.note_thoughts = note_result

        # Now note_result is a list of matrix [layer](note,
batch/time, onOrArticProb) for each layer but we only care
about
        # the hidden state of the last layer.
        # Transpose to be (batch, time, note, onOrArticProb)
        note_final =
get_last_layer(note_result).reshape((n_note,n_batch,n_time,
2)).transpose(1,2,0,3)

        # The cost of the entire procedure is the negative
log likelihood of the events all happening.
        # For the purposes of training, if the ouputted
probability is P, then the likelihood of seeing a 1 is P, and
        # the likelihood of seeing 0 is (1-P). So the
likelihood is (1-P)(1-x) + Px = 2Px - P - x + 1
        # Since they are all binary decisions, and are all
probabilities given all previous decisions, we can just
        # multiply the likelihoods, or, since we are logging
them, add the logs.

        # Note that we mask out the articulations for those
notes that aren't played, because it doesn't matter
        # whether or not those are articulated.
        # The padright is there because self.output_mat[:, :, :,
0] -> 3D matrix with (b,x,y), but we need 3d tensor with
        # (b,x,y,1) instead
        active_notes = T.shape_padright(self.output_mat[:,
1: , :, 0])
        mask =
T.concatenate([T.ones_like(active_notes), active_notes],
axis=3)

        loglikelihoods = mask *
T.log( 2*note_final*self.output_mat[:,1:] - note_final -
self.output_mat[:,1:] + 1 + self.epsilon )
        self.cost = T.neg(T.sum(loglikelihoods))

        updates, _, _, _, _ =

```

```

create_optimization_updates(self.cost, self.params,
method="adadelat")
    self.update_fun = theano.function(
        inputs=[self.input_mat, self.output_mat],
        outputs=self.cost,
        updates=updates,
        allow_input_downcast=True)

    self.update_thought_fun = theano.function(
        inputs=[self.input_mat, self.output_mat],
        outputs= ensure_list(self.time_thoughts) +
ensure_list(self.note_thoughts) + [self.cost],
        allow_input_downcast=True)

    def _predict_step_note(self, in_data_from_time, *states):
        # States is [ *hiddens, last_note_choice ]
        hiddens = list(states[:-1])
        in_data_from_prev = states[-1]
        in_data = T.concatenate([in_data_from_time,
in_data_from_prev])

        # correct for dropout
        if self.dropout > 0:
            masks = [1 - self.dropout for layer in
self.pitch_model.layers]
            masks[0] = None
        else:
            masks = []

        new_states = self.pitch_model.forward(in_data,
prev_hiddens=hiddens, dropout=masks)

        # Now new_states is a per-layer set of activations.
        probabilities = get_last_layer(new_states)

        # Thus, probabilities is a vector of two
probabilities, P(play), and P(artic | play)

        shouldPlay = self.srng.uniform() < (probabilities[0]
** self.conservativity)
        shouldArtic = shouldPlay * (self.srng.uniform() <
probabilities[1])

        chosen = T.cast(T.stack(shouldPlay, shouldArtic),
"int8")

        return ensure_list(new_states) + [chosen]

    def setup_predict(self):
        # In prediction mode, note steps are contained in the

```

time steps. So the passing gets a little bit hairy.

```

self.predict_seed = T.bmatrix()
self.steps_to_simulate = T.iscalar()

def step_time(*states):
    # States is [ *hiddens, prev_result, time]
    hiddens = list(states[:-2])
    in_data = states[-2]
    time = states[-1]

    # correct for dropout
    if self.dropout > 0:
        masks = [1 - self.dropout for layer in
self.time_model.layers]
        masks[0] = None
    else:
        masks = []

    new_states = self.time_model.forward(in_data,
prev_hiddens=hiddens, dropout=masks)

    # Now new_states is a list of matrix [layer]
(notes, hidden_states) for each layer
    time_final = get_last_layer(new_states)

    start_note_values =
theano.tensor.alloc(np.array(0,dtype=np.int8), 2)

    # This gets a little bit complicated. In the
training case, we can pass in a combination of the
    # time net's activations with the known choices.
But in the prediction case, those choices don't
    # exist yet. So instead of iterating over the
combination, we iterate over only the activations,
    # and then combine in the previous outputs in the
step. And then since we are passing outputs to
    # previous inputs, we need an additional
outputs_info for the initial "previous" output of zero.
    note_outputs_info =
([ initial_state_with_taps(layer) for layer in
self.pitch_model.layers ] +

[ dict(initial=start_note_values, taps=[-1]) ])

    notes_result, updates =
theano.scan(fn=self._predict_step_note,
sequences=[time_final], outputs_info=note_outputs_info)

    # Now notes_result is a list of matrix [layer/
```



```

output](notes, onOrArtic)
    output = get_last_layer(notes_result)

    next_input = OutputFormToInputFormOp()(output,
time + 1) # TODO: Fix time
    #next_input = T.cast(T.alloc(0, 3, 4), 'int64')

    return (ensure_list(new_states) + [ next_input,
time + 1, output ]), updates

    # start_sentinel = startSentinel()
    num_notes = self.predict_seed.shape[0]

    time_outputs_info = ([ initial_state_with_taps(layer,
num_notes) for layer in self.time_model.layers ] +
[ dict(initial=self.predict_seed, taps=[-1]),
    dict(initial=0, taps=[-1]),
    None ])

    time_result, updates = theano.scan( fn=step_time,
outputs_info=time_outputs_info,
n_steps=self.steps_to_simulate )

    self.predict_thoughts = time_result

    self.predicted_output = time_result[-1]

    self.predict_fun = theano.function(
        inputs=[self.steps_to_simulate,
self.conservativity, self.predict_seed],
        outputs=self.predicted_output,
        updates=updates,
        allow_input_downcast=True)

    self.predict_thought_fun = theano.function(
        inputs=[self.steps_to_simulate,
self.conservativity, self.predict_seed],
        outputs=ensure_list(self.predict_thoughts),
        updates=updates,
        allow_input_downcast=True)

    def setup_slow_walk(self):

        self.walk_input = theano.shared(np.ones((2,2),
dtype='int8'))
        self.walk_time = theano.shared(np.array(0,
dtype='int64'))

```

```

        self.walk_hiddens = [theano.shared(np.ones((2,2),
dtype=theano.config.floatX)) for layer in
self.time_model.layers if has_hidden(layer)]

        # correct for dropout
        if self.dropout > 0:
            masks = [1 - self.dropout for layer in
self.time_model.layers]
            masks[0] = None
        else:
            masks = []

        new_states = self.time_model.forward(self.walk_input,
prev_hiddens=self.walk_hiddens, dropout=masks)

        # Now new_states is a list of matrix [layer](notes,
hidden_states) for each layer
        time_final = get_last_layer(new_states)

        start_note_values =
theano.tensor.alloc(np.array(0,dtype=np.int8), 2)
        note_outputs_info = ([ initial_state_with_taps(layer)
for layer in self.pitch_model.layers ] +
[ dict(initial=start_note_values, taps=[-1]) ])

        notes_result, updates =
theano.scan(fn=self._predict_step_note,
sequences=[time_final], outputs_info=note_outputs_info)

        # Now notes_result is a list of matrix [layer/output]
(notes, onOrArtic)
        output = get_last_layer(notes_result)

        next_input = OutputFormToInputFormOp()(output,
self.walk_time + 1) # TODO: Fix time
        #next_input = T.cast(T.alloc(0, 3, 4),'int64')

        slow_walk_results = (new_states[:-1] +
notes_result[:-1] + [ next_input, output ])

        updates.update({
            self.walk_time: self.walk_time+1,
            self.walk_input: next_input
        })

        updates.update({hidden:newstate for hidden, newstate,
layer in zip(self.walk_hiddens, new_states,
self.time_model.layers) if has_hidden(layer)})

```

```

        self.slow_walk_fun = theano.function(
            inputs=[self.conservativity],
            outputs=slow_walk_results,
            updates=updates,
            allow_input_downcast=True)

    def start_slow_walk(self, seed):
        seed = np.array(seed)
        num_notes = seed.shape[0]

        self.walk_time.set_value(0)
        self.walk_input.set_value(seed)
        for layer, hidden in zip((l for l in
self.time_model.layers if has_hidden(l)),self.walk_hiddens):
            hidden.set_value(np.repeat(np.reshape(layer.initial_hidden_state.get_value(
(1,-1))), num_notes, axis=0))

```

```

===- - -===

```

```

import os, random
from midi_to_statematrix import *
from data import *
#import cPickle as pickle
import pickle

import signal

batch_width = 10 # number of sequences in a batch
batch_len = 16*8 # length of each sequence
division_len = 16 # interval between possible start locations

def loadPieces(dirpath):

    pieces = {}

```

```

    for fname in os.listdir(dirpath):
        if fname[-4:] not in ('.mid', '.MID'):
            continue

        name = fname[:-4]

        outMatrix =
midiToNoteStateMatrix(os.path.join(dirpath, fname))
        if len(outMatrix) < batch_len:
            continue

        pieces[name] = outMatrix
        print("Loaded {}".format(name))

    return pieces

def getPieceSegment(pieces):
    piece_output = random.choice(pieces.values())
    start = random.randrange(0, len(piece_output) -
batch_len, division_len)
    # print "Range is {} {} {} ->
{}".format(0, len(piece_output) - batch_len, division_len, start)

    seg_out = piece_output[start:start+batch_len]
    seg_in = noteStateMatrixToInputForm(seg_out)

    return seg_in, seg_out

def getPieceBatch(pieces):
    i, o = zip(*[getPieceSegment(pieces) for _ in
range(batch_width)])
    return numpy.array(i), numpy.array(o)

def trainPiece(model, pieces, epochs, start=0):
    stopflag = [False]
    def signal_handler(signame, sf):
        stopflag[0] = True
    old_handler = signal.signal(signal.SIGINT, signal_handler)
    for i in range(start, start+epochs):
        if stopflag[0]:
            break
        error = model.update_fun(*getPieceBatch(pieces))
        if i % 100 == 0:
            print("epoch {}, error={}".format(i, error))
        if i % 500 == 0 or (i % 100 == 0 and i < 1000):
            xIpt, xOpt = map(numpy.array,
getPieceSegment(pieces))

    noteStateMatrixToMidi(numpy.concatenate((numpy.expand_dims(xOpt[0],
0), model.predict_fun(batch_len, 1, xIpt[0])),

```

```

axis=0),'output/sample{}'.format(i))
        pickle.dump(model.learned_config,open('output/
params{}.p'.format(i), 'wb'))
        signal.signal(signal.SIGINT, old_handler)
===----===

import theano, theano.tensor as T
import numpy as np

from data import noteStateSingleToInputForm

class OutputFormToInputFormOp(theano.Op):
    # Properties attribute
    __props__ = ()

    def make_node(self, state, time):
        state = T.as_tensor_variable(state)
        time = T.as_tensor_variable(time)
        return theano.Apply(self, [state, time],
[T.bmatrix()])

    # Python implementation:
    def perform(self, node, inputs_storage, output_storage):
        state, time = inputs_storage
        output_storage[0][0] =
np.array(noteStateSingleToInputForm(state, time),
dtype='int8')===----===

import midi, numpy

lowerBound = 24
upperBound = 102

def midiToNoteStateMatrix(midifile):

    pattern = midi.read_midifile(midifile)

    timeleft = [track[0].tick for track in pattern]

    posns = [0 for track in pattern]

    statematrix = []
    span = upperBound-lowerBound
    time = 0

    state = [[0,0] for x in range(span)]
    statematrix.append(state)
    while True:
        if time % (pattern.resolution / 4) ==

```

```

(pattern.resolution / 8):
    # Crossed a note boundary. Create a new state,
    defaulting to holding notes
    oldstate = state
    state = [[oldstate[x][0],0] for x in range(span)]
    statematrix.append(state)

    for i in range(len(timeleft)):
        while timeleft[i] == 0:
            track = pattern[i]
            pos = posns[i]

            evt = track[pos]
            if isinstance(evt, midi.NoteEvent):
                if (evt.pitch < lowerBound) or (evt.pitch
>= upperBound):
                    pass
                    # print "Note {} at time {} out of
bounds (ignoring)".format(evt.pitch, time)
                else:
                    if isinstance(evt, midi.NoteOffEvent)
or evt.velocity == 0:
                        state[evt.pitch-lowerBound] = [0,
0]
                    else:
                        state[evt.pitch-lowerBound] = [1,
1]
                    elif isinstance(evt, midi.TimeSignatureEvent):
                        if evt.numerator not in (2, 4):
                            # We don't want to worry about non-4
time signatures. Bail early!
                            # print "Found time signature event
{}. Bailing!".format(evt)
                            return statematrix

            try:
                timeleft[i] = track[pos + 1].tick
                posns[i] += 1
            except IndexError:
                timeleft[i] = None

            if timeleft[i] is not None:
                timeleft[i] -= 1

        if all(t is None for t in timeleft):
            break

        time += 1

    return statematrix

```

```

def noteStateMatrixToMidi(statematrix, name="example"):
    statematrix = numpy.asarray(statematrix)
    pattern = midi.Pattern()
    track = midi.Track()
    pattern.append(track)

    span = upperBound-lowerBound
    tickscale = 55

    lastcmdtime = 0
    prevstate = [[0,0] for x in range(span)]
    for time, state in enumerate(statematrix +
[prevstate[:] ]):
        offNotes = []
        onNotes = []
        for i in range(span):
            n = state[i]
            p = prevstate[i]
            if p[0] == 1:
                if n[0] == 0:
                    offNotes.append(i)
                elif n[1] == 1:
                    offNotes.append(i)
                    onNotes.append(i)
                elif n[0] == 1:
                    onNotes.append(i)
            for note in offNotes:
                track.append(midi.NoteOffEvent(tick=(time-
lastcmdtime)*tickscale, pitch=note+lowerBound))
                lastcmdtime = time
            for note in onNotes:
                track.append(midi.NoteOnEvent(tick=(time-
lastcmdtime)*tickscale, velocity=40, pitch=note+lowerBound))
                lastcmdtime = time

        prevstate = state

    eot = midi.EndOfTrackEvent(tick=1)
    track.append(eot)

    midi.write_midifile("{} .mid".format(name), pattern)===---
===

import itertools
from midi_to_statematrix import upperBound, lowerBound

def startSentinel():
    def noteSentinel(note):
        position = note

```

```

        part_position = [position]

        pitchclass = (note + lowerBound) % 12
        part_pitchclass = [int(i == pitchclass) for i in
range(12)]

        return part_position + part_pitchclass + [0]*66 + [1]
    return [noteSentinel(note) for note in range(upperBound-
lowerBound)]

def getOrDefault(l, i, d):
    try:
        return l[i]
    except IndexError:
        return d

def buildContext(state):
    context = [0]*12
    for note, notestate in enumerate(state):
        if notestate[0] == 1:
            pitchclass = (note + lowerBound) % 12
            context[pitchclass] += 1
    return context

def buildBeat(time):
    return [2*x-1 for x in [time%2, (time//2)%2, (time//4)%2,
(time//8)%2]]

def noteInputForm(note, state, context, beat):
    position = note
    part_position = [position]

    pitchclass = (note + lowerBound) % 12
    part_pitchclass = [int(i == pitchclass) for i in
range(12)]
    # Concatenate the note states for the previous vicinity
    part_prev_vicinity =
list(itertools.chain.from_iterable((getOrDefault(state,
note+i, [0,0]) for i in range(-12, 13))))

    part_context = context[pitchclass:] + context[:pitchclass]

    return part_position + part_pitchclass +
part_prev_vicinity + part_context + beat + [0]

def noteStateSingleToInputForm(state,time):
    beat = buildBeat(time)
    context = buildContext(state)
    return [noteInputForm(note, state, context, beat) for
note in range(len(state))]

```



```

def noteStateMatrixToInputForm(statematrix):
    # NOTE: May have to transpose this or transform it in
    some way to make Theano like it
    #[startSentinel()] +
    inputform = [ noteStateSingleToInputForm(state,time) for
    time,state in enumerate(statematrix) ]
    return inputform
===- - -===

import cPickle as pickle
import gzip
import numpy
from midi_to_statematrix import *

import multi_training
import model

def
gen_adaptive(m,pcs,times,keep_thoughts=False,name="final"):
    xIpt, xOpt = map(lambda x: numpy.array(x,
dtype='int8'), multi_training.getPieceSegment(pcs))
    all_outputs = [xOpt[0]]
    if keep_thoughts:
        all_thoughts = []
    m.start_slow_walk(xIpt[0])
    cons = 1
    for time in range(multi_training.batch_len*times):
        resdata = m.slow_walk_fun( cons )
        nnotes = numpy.sum(resdata[-1][:,0])
        if nnotes < 2:
            if cons > 1:
                cons = 1
            cons -= 0.02
        else:
            cons += (1 - cons)*0.3
        all_outputs.append(resdata[-1])
        if keep_thoughts:
            all_thoughts.append(resdata)

noteStateMatrixToMidi(numpy.array(all_outputs),'output/'+name)
    if keep_thoughts:
        pickle.dump(all_thoughts,
open('output/'+name+'.p','wb'))

def fetch_train_thoughts(m,pcs,batches,name="trainthoughts"):
    all_thoughts = []
    for i in range(batches):
        ipt, opt = multi_training.getPieceBatch(pcs)
        thoughts = m.update_thought_fun(ipt,opt)

```

```

        all_thoughts.append((ipt,opt,thoughts))
    pickle.dump(all_thoughts,
open('output/'+name+'.p','wb'))

if __name__ == '__main__':

    pcs = multi_training.loadPieces("music")

    m = model.Model([300,300],[100,50], dropout=0.5)

    multi_training.trainPiece(m, pcs, 10000)

    pickle.dump( m.learned_config, open( "output/
final_learned_config.p", "wb" ) )
===---===

import numpy as np

def sigmoid(x):
    return 1 / (1 + np.exp(-x))

def actToColor(memcell, activation):
    return [0, sigmoid(activation), sigmoid(memcell)]

def internalMatrixToImgArray(inmat):
    return np.array(
        [[actToColor(m,a) for m,a in
zip(row[:len(row)/2],row[len(row)/2:])]
        for row in inmat])

def probAndSuccessToImgArray(prob, succ, idx):
    return np.array([[pr[idx]]*3,[sr[idx],0,0]] for pr,
sr in zip(prob, succ)])

def thoughtsToImageArray(thoughts):

    spacer = np.zeros((thoughts[0].shape[0], 5, 3))

    sequence = [
        spacer,

probAndSuccessToImgArray(thoughts[4],thoughts[6], 0),
        spacer,

probAndSuccessToImgArray(thoughts[4],thoughts[6], 1)
    ]

    for thought in thoughts[:-3]:
        sequence = [ spacer,

```

```

internalMatrixToImgArray(thought) ] + sequence

    return (np.concatenate(sequence,
axis=1 )*255).astype('uint8')

def pastColor(prob, succ):
    return [prob[0], succ[0], succ[1]*succ[0]]

def drawPast(probs, succs):
    return np.array([
        [
            pastColor(probs[time][note_idx],
succs[time][note_idx])
            for time in range(len(probs))
        ]
        for note_idx in range(len(probs[0]))
    ])

def thoughtsAndPastToStackedArray(thoughts, probs, succs,
len_past):

    vert_spacer = np.zeros((thoughts[0].shape[0], 5, 3))

    past_out = drawPast(probs, succs)

    if len(probs) < len_past:
        past_out = np.pad(past_out, ((0,0),(len_past-
len(probs),0),(0,0)), mode='constant')

    def add_cur(ipt):
        return np.concatenate((
            ipt,
            vert_spacer,

probAndSuccessToImgArray(thoughts[-3],thoughts[-1], 0),
            vert_spacer,

probAndSuccessToImgArray(thoughts[-3],thoughts[-1], 1)),
axis=1)

    horiz_spacer = np.zeros((5, 1, 3))

    rows = [add_cur(past_out[-len_past:])]

    for thought in thoughts[:-3]:
        rows += [ horiz_spacer,
add_cur(internalMatrixToImgArray(thought)) ]

    maxlen = max([x.shape[1] for x in rows])
    rows = [np.pad(row, ((0,0),(maxlen-row.shape[1],0),

```

```
(0,0)), mode='constant') for row in rows]
```

```
        return (np.concatenate(rows,  
axis=0 )*255).astype('uint8')===---===
```