

```
79 audio.py
23 folds.py
0 __init__.py
80 padding.py
235 training.py
377 transforms.py
128 utils.py
922 total
359 apc.py
1249 classifiers.py
395 cpc.py
0 __init__.py
57 losses.py
2060 total
0 __init__.py
59 sound_dataset.py
2 wcl_datasets.txt
61 total
1 348 adversarial_test.py
2 32 create_class_map.py
3 1 dash.txt
4 0 datasets
5 154 evaluate_2d_cnn.py
6 439 finetune_hierarchical_cnn.py
7 201 LICENSE
8 133 linear_blend.py
9 0 networks
10 0 ops
11 124 predict_2d_cnn.py
12 220 README.md
13 190 relabel_noisy_data.py
14 100 requirements.txt
15 510 train_2d_cnn.py
16 278 train_apc.py
17 486 train_backbone_cnn.py
18 288 train_cpc.py
19 509 train_hierarchical_cnn.py
4013 total
# 3rd place solution to Freesound Audio Tagging 2019 Challenge
```

My approach is outlined below.

****Models****

I used two types of models, both are based on convolutions. The first type uses 2d convolutions and works on top of mel-scale spectrograms, while the second uses 1d-convolutions on top of raw STFT representations with relatively small window size like 256, so it's only 5 ms per frame or so. Both types of models are relatively shallow and consist of 10-12 convolutional layers (or 5-6 resnet blocks) with a small number of filters. I use a form of deep supervision by applying global max pooling after each block (typically starting from the first or second block) and then concatenating maxpool outputs from each layer to form the final feature vector which then goes to a 2-layer fully-connected classifier. I also tried using RNNs instead of a max pooling for some

models. It made results a bit worse, but RNN seemed to make different mistakes, so it turned out to be a good member of the final ensemble.

****Frequency encoding****

2d convolutions are position-invariant, so the output of a convolution would be the same regardless of where the feature is located. Spectrograms are not images, Y-axis corresponds to signal frequency, so it would be nice to assist a model by providing this sort of information. For this purpose, I used a linear frequency map going from -1 to 1 and concatenated it to input spectrogram as a second channel. It's hard to estimate now without retraining all the models how much gain I got from this little modification, but I can say It was no less than 0.005 in terms of local CV score.

****This is not really a classification task****

Most teams treated the problem as a multilabel classification and used a form of a binary loss such as binary cross entropy or focal loss. This approach is definitely valid, but in my experiments, it appeared to be a little suboptimal. The reason is the metric (lwlrap) is not a pure classification metric. Contrary to accuracy or f-score, it is based on **ranks**. So it wasn't really a surprise for me when I used a loss function based on ranks rather than on binary outputs, I got a huge improvement. Namely, I used something called LSEP (<https://arxiv.org/abs/1704.03135>) which is just a soft version of pairwise rank loss. It makes your model to score positive classes higher than negative ones, while a binary loss increases positive scores and decreases negative scores independently. When I switched to LSEP from BCE, I immediately got approximately 0.015 of improvement, and, as a nice bonus, my models started to converge much faster.

****Data augmentation****

I used two augmentation strategies. The first one is a modified MixUp. In contrast to the original approach, I used OR rule for mixing labels. I did so because a mix of two sounds still allows you to hear both. I tried the original approach with weighted targets on some point and my results got worse.

The second strategy is augmentations based on audio effects such as reverb, pitch, tempo and overdrive. I chose the parameters of these augmentations by carefully listening to augmented samples.

I have found augmentations to be very important for getting good results. I guess the total improvement I got from these two strategies is about 0.05 or so. I also tried several other approaches such as splitting the audio into several chunks and then shuffling them, replacing some parts of the original signals with silence and some other, but they didn't make my models better.

****Training****

I used quite large audio segments for training. For most of my models, I used segments from 8 to 12 seconds. I didn't use TTA for inference and used full-length audio instead.

****Noisy data****

I tried several unsupervised approaches such as [Contrastive Predicting Coding] (<https://arxiv.org/abs/1807.03748>), but never managed to get good results from it.

I ended up applying a form of iterative pseudolabeling. I predicted new labels for the noisy subset using a model trained on curated data only, chose best 1k in terms of the agreement between the predicted labels and actual labels and added these samples to the curated subset with the original labels. I repeated the procedure using top 2k labels this time. I applied this approach several times until I reached 5k best noisy samples. At that point, predictions generated by a model started to diverge significantly from the actual noisy labels. I decided to discard the labels of the remaining noisy samples and simply used model prediction as actual labels. In total, I trained approximately 20 models using different subsets of the noisy train set with different pseudolabeling strategies.

****Inference****

I got a great speed-up by computing both STFT spectrograms and mel spectrograms on a GPU. I also grouped samples with similar lengths together to avoid excessive padding. These two methods combined with relatively small models allowed me to predict the first stage test set in only 1 minute by any of my models (5 folds).

****Final ensemble****

For the final solution, I used a simple average of 11 models trained with slightly different architectures (1d/2d cnn, rnn/no-rnn), slightly different subsets of the noisy set (see "noisy data" section) and slightly different hyperparameters.

Project structure

Main training scripts are `train_2d_cnn.py` and `train_hierarchical_cnn.py`. All classification models are defined in `networks/classifiers`. All data augmentations are defined in `ops/transforms`.

Setting up the environment

I recommend using some environment manager such as conda or virtualenv in order to avoid potential conflicts between different versions of packages. To install all required packages, simply run `pip install -r requirements.txt`. This might take up to 15 minutes depending on your internet connection speed.

Preparing data

I place all the data into `data/` directory, please adjust the following code to match yours data location. Run

```
```bash
python create_class_map.py --train_df data/train_curated.csv --output_file data/classmap.json
```
```

This simply creates a JSON file with deterministic classname->label mapping used in all future experiments.

Running a basic 2d model

```
```bash
python train_2d_cnn.py \
 --train_df data/train_curated.csv \
 --train_data_dir data/train_curated/ \
 --classmap data/classmap.json \
 --device=cuda \
 --optimizer=adam \
 --folds 0 1 2 3 4 \
 --n_folds=5 \
 --log_interval=10 \
 --batch_size=20 \
 --epochs=20 \
 --accumulation_steps=1 \
 --save_every=20 \
 --num_conv_blocks=5 \
 --conv_base_depth=50 \
 --growth_rate=1.5 \
 --weight_decay=0.0 \
 --start_deep_supervision_on=1 \
 --aggregation_type=max \
 --lr=0.003 \
 --scheduler=1cycle_0.0001_0.005 \
 --test_data_dir data/test \
 --sample_submission data/sample_submission.csv \
 --num_workers=6 \
 --output_dropout=0.0 \
 --p_mixup=0.0 \
 --switch_off_augmentations_on=15 \
 --features=mel_2048_1024_128 \
 --max_audio_length=15 \
 --p_aug=0.0 \
 --label=basic_2d_cnn
```
```

Running a 2d model with augmentations

```
```bash
python train_2d_cnn.py \
 --train_df data/train_curated.csv \
 --train_data_dir data/train_curated/ \
```

```

--classmap data/classmap.json \
--device=cuda \
--optimizer=adam \
--folds 0 1 2 3 4 \
--n_folds=5 \
--log_interval=10 \
--batch_size=20 \
--epochs=100 \
--accumulation_steps=1 \
--save_every=20 \
--num_conv_blocks=5 \
--conv_base_depth=100 \
--growth_rate=1.5 \
--weight_decay=0.0 \
--start_deep_supervision_on=1 \
--aggregation_type=max \
--lr=0.003 \
--scheduler=1cycle_0.0001_0.005 \
--test_data_dir data/test \
--sample_submission data/sample_submission.csv \
--num_workers=16 \
--output_dropout=0.5 \
--p_mixup=0.5 \
--switch_off_augmentations_on=90 \
--features=mel_2048_1024_128 \
--max_audio_length=15 \
--p_aug=0.75 \
--label=2d_cnn
...

```

Note that each such run is followed by a creation of a new experiment subdirectory in the `experiments` folder. Each experiment has the following structure:

```

```bash
experiments/some_experiment/
??? checkpoints
??? command
??? commit_hash
??? config.json
??? log
??? predictions
??? results.json
??? summaries
```

```

### Using a clean model to select noisy samples

Create a new predictions directory:

```
```mkdir predictions/```
```

Then, running

```

```bash
python predict_2d_cnn.py \

```

```
--experiment=path_to_an_experiment (see above) \
--test_df=data/train_noisy.csv \
--test_data_dir=data/train_noisy/ \
--output_df=predictions/noisy_probabilities.csv \
--classmap=data/classmap.json \
--device=cuda
...
```

creates a new csv file in the predictions folder with the class probabilities for the noisy dataset.

Running

```
```bash  
python relabel_noisy_data.py \  
--noisy_df=data/train_noisy.csv \  
--noisy_predictions_df=predictions/noisy_probabilities.csv \  
--output_df=predictions/train_noisy_relabeled_1k.csv \  
--mode=scoring_1000  
...
```

creates a new noisy dataframe where only top 1k labels in terms of agreement between the model and the actual labels are kept.

Running a 2d model with noisy data

```
```bash  
python train_2d_cnn.py \
--train_df data/train_curated.csv \
--train_data_dir data/train_curated/ \
--noisy_train_df predictions/train_noisy_relabeled_1k.csv \
--noisy_train_data_dir data/train_noisy/ \
--classmap data/classmap.json \
--device=cuda \
--optimizer=adam \
--folds 0 1 2 3 4 \
--n_folds=5 \
--log_interval=10 \
--batch_size=20 \
--epochs=150 \
--accumulation_steps=1 \
--save_every=20 \
--num_conv_blocks=6 \
--conv_base_depth=100 \
--growth_rate=1.5 \
--weight_decay=0.0 \
--start_deep_supervision_on=1 \
--aggregation_type=max \
--lr=0.003 \
--scheduler=1cycle_0.0001_0.005 \
--test_data_dir data/test \
--sample_submission data/sample_submission.csv \
--num_workers=16 \
--output_dropout=0.7 \
...
```

```
--p_mixup=0.5 \
--switch_off_augmentations_on=140 \
--features=mel_2048_1024_128 \
--max_audio_length=15 \
--p_aug=0.75 \
--label=2d_cnn_noisy
...
```

Note that `relabel\_noisy\_data.py` script supports multiple relabeling strategies. I mostly followed "scoring" strategy (selecting top-k noisy samples based on the agreement between the model and the actual labels), but after 5k noisy samples I switched to "relabelall-replacenan" strategy which is just a pseudolabeling (usage of the old model outputs) where the samples without any predictions are discarded.

=====

```
abs1-py==0.7.1
astor==0.7.1
attrs==19.1.0
audioread==2.1.6
backcall==0.1.0
bleach==3.1.0
certifi==2019.3.9
cffi==1.12.2
chardet==3.0.4
cycller==0.10.0
decorator==4.4.0
defusedxml==0.5.0
entrypoints==0.3
gast==0.2.2
grpcio==1.19.0
h5py==2.9.0
idna==2.8
ipykernel==5.1.0
ipython==7.4.0
ipython-genutils==0.2.0
ipywidgets==7.4.2
iterative-stratification==0.1.6
jedi==0.13.3
Jinja2==2.10.1
joblib==0.13.2
jsonschema==3.0.1
jupyter==1.0.0
jupyter-client==5.2.4
jupyter-console==6.0.0
jupyter-core==4.4.0
kaggle==1.5.3
Keras-Applications==1.0.7
Keras-Preprocessing==1.0.9
kiwisolver==1.0.1
librosa==0.6.3
llvmlite==0.28.0
mag==0.1
Markdown==3.1
MarkupSafe==1.1.1
matplotlib==3.0.3
```

```
mistune==0.8.4
mock==2.0.0
munch==2.3.2
nbconvert==5.4.1
nbformat==4.4.0
notebook==5.7.8
numba==0.43.1
numpy==1.16.2
pandas==0.24.2
pandocfilters==1.4.2
parso==0.4.0
pbr==5.1.3
pexpect==4.7.0
pickleshare==0.7.5
Pillow==6.0.0
pkg-resources==0.0.0
pretrainedmodels==0.7.4
prometheus-client==0.6.0
prompt-toolkit==2.0.9
protobuf==3.7.1
ptyprocess==0.6.0
pycparser==2.19
Pygments==2.3.1
pyparsing==2.3.1
pysistent==0.14.11
pysndfx==0.3.6
python-dateutil==2.8.0
python-slugify==3.0.2
pytz==2018.9
pyzmq==18.0.1
qtconsole==4.4.3
requests==2.21.0
resampy==0.2.1
scikit-learn==0.20.3
scipy==1.2.1
Send2Trash==1.5.0
six==1.12.0
SoundFile==0.10.2
tensorboard==1.13.1
tensorboardX==1.6
tensorflow==1.13.1
tensorflow-estimator==1.13.0
termcolor==1.1.0
terminado==0.8.2
testpath==0.4.2
text-unidecode==1.2
torch==1.0.1.post2
torchcontrib==0.0.2
torchvision==0.2.2.post3
tornado==6.0.2
tqdm==4.31.1
traitlets==4.3.2
umap-learn==0.3.8
urllib3==1.24.1
wcwidth==0.1.7
webencodings==0.5.1
```



```
Werkzeug==0.15.2
widgetsnbextension==3.4.2

git+https://github.com/ex4sperans/mag
import os
import gc
import argparse
import json
import math
from functools import partial

import tqdm
import pandas as pd
import numpy as np
import torch
from matplotlib import pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score

from datasets.sound_dataset import SoundDataset
from networks.classifiers import HierarchicalCNNClassificationModel
from ops.folds import train_validation_data
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio)
from ops.utils import load_json, get_class_names_from_classmap,
lwrap
from ops.padding import make_collate_fn
from networks.classifiers import ResnetBlock

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)
parser.add_argument(
 "--train_data_dir", required=True, type=str,
 help="path to train data"
)
parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data"
)
parser.add_argument(
 "--test_df", required=True, type=str,
 help="path to train dataframe"
)
```

```
parser.add_argument(
 "--val_size", required=True, type=float,
 help="size of the validation set"
)
parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
parser.add_argument(
 "--epochs", type=int, default=100,
 help="number of epochs"
)
parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
parser.add_argument(
 "--features", type=str, required=True,
 help="feature descriptor"
)
parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are sampled"
)
parser.add_argument(
 "--batches_to_save", type=int, default=3,
 help="how many batches to save"
)
parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)

args = parser.parse_args()

train_df = pd.read_csv(args.train_df)
test_df = pd.read_csv(args.test_df)

if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index(drop=True)
 test_df = test_df.sample(args.max_samples).reset_index(drop=True)

all_train_fnames = [
 os.path.join(args.train_data_dir, fname) for fname in train_
```

```

df.fname.values]
all_test_fnames = [
 os.path.join(args.test_data_dir, fname) for fname in test_df
 .fname.values]

fnames = np.concatenate([all_train_fnames, all_test_fnames])
labels = np.concatenate([np.ones(len(train_df)), np.zeros(len(test_df))])

train_fnames, val_fnames, train_labels, val_labels = train_test_split(
 fnames, labels, test_size=args.val_size, shuffle=True)

audio_transform = AudioFeatures(args.features)

class Model(torch.nn.Module):

 def __init__(self):
 super().__init__()

 self.features = torch.nn.Sequential(
 torch.nn.BatchNorm1d(audio_transform.n_features),
 torch.nn.Conv1d(audio_transform.n_features, 32, kernel_size=1),
 ResnetBlock(32),
 torch.nn.MaxPool1d(kernel_size=2, stride=2),
 torch.nn.BatchNorm1d(32),
 torch.nn.Conv1d(32, 32, kernel_size=3),
 ResnetBlock(32),
 torch.nn.MaxPool1d(kernel_size=2, stride=2),
 torch.nn.BatchNorm1d(32),
 torch.nn.Conv1d(32, 64, kernel_size=3),
 ResnetBlock(64)
)

 self.pool = torch.nn.AdaptiveMaxPool1d(1)

 self.classifier = torch.nn.Sequential(
 torch.nn.BatchNorm1d(64),
 torch.nn.Conv1d(64, 1, kernel_size=1)
)

 def forward(self, x):

 x = x.permute(0, 2, 1)
 x = self.features(x)
 x = self.classifier(x)
 x = torch.sigmoid(x)
 nonpooled = x
 x = self.pool(x).squeeze(-1)

 return x.squeeze(1), nonpooled.squeeze(1)

train_loader = torch.utils.data.DataLoader(

```

```

 SoundDataset(
 audio_files=train_fnames,
 labels=train_labels,
 transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_length),
 audio_transform,
 RenameFields({"raw_labels": "labels"}),
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
])
),
 shuffle=True,
 drop_last=True,
 batch_size=args.batch_size,
 num_workers=4,
 collate_fn=make_collate_fn({"signal": audio_transform.paddin
g_value})),
)

validation_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=val_fnames,
 labels=val_labels,
 transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_length),
 audio_transform,
 RenameFields({"raw_labels": "labels"}),
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
])
),
 shuffle=False,
 drop_last=False,
 batch_size=args.batch_size,
 num_workers=4,
 collate_fn=make_collate_fn({"signal": audio_transform.paddin
g_value})),
)

model = Model().to(args.device)
optimizer = torch.optim.Adam(model.parameters(), args.lr)

for epoch in range(args.epochs):

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n"
 .format(epoch=epoch)
)

```

```

model.train()

with tqdm.tqdm(total=len(train_loader), ncols=80) as pb:
 for sample in train_loader:
 signal, labels = (
 sample["signal"].to(args.device),
 sample["labels"].to(args.device).float()
)

 probs, nonpooled = model(signal)

 optimizer.zero_grad()
 loss = torch.nn.functional.binary_cross_entropy(probs, labels)
 loss.backward()
 optimizer.step()

 pb.update()
 pb.set_description("Loss: {:.4f}".format(loss.item()))

model.eval()

val_probs = []
val_labels = []

with torch.no_grad():
 for sample in validation_loader:
 signal, labels = (
 sample["signal"].to(args.device),
 sample["labels"].to(args.device).float()
)

 probs, nonpooled = model(signal)

 val_probs.extend(probs.data.cpu().numpy())
 val_labels.extend(labels.data.cpu().numpy())

auc = roc_auc_score(val_labels, val_probs)

print("\nEpoch: {}, AUC: {}".format(epoch, auc))

model.eval()

plot probabilities
loader = iter(validation_loader)
directory = "plots/"
os.makedirs(directory, exist_ok=True)

for n in range(args.batches_to_save):
 with torch.no_grad():

```

```

 sample = next(loader)
 signal, labels = (
 sample["signal"].to(args.device),
 sample["labels"].to(args.device).float()
)

 probs, nonpooled = model(signal)

 nonpooled = nonpooled.data.cpu().numpy()
 signal = signal.data.cpu().numpy()
 labels = labels.data.cpu().numpy()

 for k in range(len(signal)):

 fig = plt.figure(figsize=(20, 7))
 fig.suptitle(str(labels[k]))
 ax = fig.add_subplot(211)
 ax.imshow(np.transpose(signal[k]))
 ax = fig.add_subplot(212)
 ax.plot(nonpooled[k])
 ax.set_ylim(0, 1)
 ax.set_xlim(0, len(nonpooled[k]) - 1)

 fig.savefig(os.path.join(directory, "plot_{}_{}.png".format(n, k)))
 plt.close()

compute average scores for classes
class_map = load_json(args.classmap)

names_with_labels = [
 fname for fname in val_fnames if fname in all_train_fnames]
labels = pd.DataFrame({
 "fname": [os.path.basename(fname) for fname in names_with_labels]
}).merge(
 train_df, on="fname", how="left").labels.values

loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=names_with_labels,
 labels=[item.split(",") for item in labels],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map),
 SampleLongAudio(max_length=args.max_audio_length),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 drop_last=False,
 batch_size=args.batch_size,
 num_workers=4,
 collate_fn=make_collate_fn({"signal": audio_transform.paddin

```

```

g_value)),
)

all_probs = []
all_labels = []

with torch.no_grad():
 for sample in loader:
 signal, labels = (
 sample["signal"].to(args.device),
 sample["labels"].to(args.device).float()
)

 probs, nonpooled = model(signal)

 all_probs.extend(probs.data.cpu().numpy())
 all_labels.extend(labels.data.cpu().numpy())

all_probs = np.array(all_probs)
all_labels = np.array(all_labels)

scores = all_labels * np.expand_dims(all_probs, -1)
mean_scores = scores.sum(axis=0) / all_labels.sum(axis=0)

classnames = get_class_names_from_classmap(class_map)

pd.options.display.max_rows = 100

print()
print(pd.DataFrame({"classname": classnames, "scores": mean_scores}))

=====
import json
import argparse

import pandas as pd

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)

parser.add_argument(
 "--output_file", type=str, required=True,
 help="where to save classmap"
)

args = parser.parse_args()

```

```
df = pd.read_csv(args.train_df)

all_labels = set()
for item in df.labels:
 all_labels.update(item.split(","))

classmap = dict((v, k) for k, v in enumerate(sorted(all_labels))
)

with open(args.output_file, "w") as file:
 json.dump(classmap, file, indent=4, sort_keys=True)=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
from mag.utils import green, bold
import mag

from datasets.sound_dataset import SoundDataset
from networks.classifiers import TwoDimensionalCNNClassification
Model
from ops.folds import train_validation_data_stratified
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation, FlipAudio, ShuffleAudio)
from ops.utils import load_json, get_class_names_from_classmap,
lwlrap
from ops.padding import make_collate_fn

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--experiment", type=str, required=True,
 help="path to an experiment"
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)

parser.add_argument(
 "--train_data_dir", required=True, type=str,
```



```

 help="path to train data"
)
 parser.add_argument(
 "--noisy_train_df", type=str,
 help="path to noisy train dataframe (optional)"
)
 parser.add_argument(
 "--noisy_train_data_dir", type=str,
 help="path to noisy train data (optional)"
)
 parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
 parser.add_argument(
 "--batch_size", type=int, default=32,
 help="batch size used for prediction"
)
 parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are sampled"
)
 parser.add_argument(
 "--n_tta", type=int, default=1,
 help="number of tta"
)
 parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
 parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)

args = parser.parse_args()

class_map = load_json(args.classmap)

train_df = pd.read_csv(args.train_df)

with Experiment(resume_from=args.experiment) as experiment:

 config = experiment.config

 audio_transform = AudioFeatures(config.data.features)

 splits = list(train_validation_data_stratified(
 train_df.fname, train_df.labels, class_map,
 config.data._n_folds, config.data._kfold_seed))

 all_labels = np.zeros(
 shape=(len(train_df), len(class_map)), dtype=np.float32)
 all_predictions = np.zeros(

```

```

 shape=(len(train_df), len(class_map)), dtype=np.float32)

 for fold in range(config.data._n_folds):

 print("\n\n ----- Fold {}\n".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
}
 if torch.cuda.is_available() else {})

 valid_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],
 labels=[item.split(",") for item in train_df.labels.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 SampleLongAudio(args.max_audio_length),
 ShuffleAudio(chunks_range=(12, 20), p=1.0),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
]),

),
 shuffle=False,
 batch_size=args.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
 **loader_kwargs
)

 model = TwoDimensionalCNNClassificationModel(
 experiment, device=args.device)
 model.load_best_model(fold)
 model.eval()

 val_preds = model.predict(valid_loader, n_tta=args.n_tta
)
 val_labels = np.array([item["labels"] for item in valid_loader.dataset])

 all_labels[valid] = val_labels
 all_predictions[valid] = val_preds

 metric = lwlraps(val_labels, val_preds)

```

```
 print("Fold metric:", metric)

 metric = lwlapr(all_labels, all_predictions)

 print("\nOverall metric:", green(bold(metric)))

=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
import mag
from sklearn.model_selection import train_test_split

from datasets.sound_dataset import SoundDataset
from networks.classifiers import HierarchicalCNNClassificationModel
del
from ops.folds import train_validation_data
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 STFT, MapLabels, RenameFields, MixUp)
from ops.utils import load_json, get_class_names_from_classmap,
lwlapr
from ops.padding import make_collate_fn

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)
parser.add_argument(
 "--train_data_dir", required=True, type=str,
 help="path to train data"
)
parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data"
)
parser.add_argument(
```

```
 "--sample_submission", required=True, type=str,
 help="path sample submission"
)
 parser.add_argument(
 "--pretrained_model", required=True, type=str,
 help="path to old experiment"
)
 parser.add_argument(
 "--pretrained_fold", required=True, type=int,
 help="pretrained fold"
)
 parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
 parser.add_argument(
 "--log_interval", default=10, type=int,
 help="how frequently to log batch metrics"
 "in terms of processed batches"
)
 parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
 parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
 parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
 parser.add_argument(
 "--holdout_size", type=float, default=0.0,
 help="size of holdout set"
)
 parser.add_argument(
 "--epochs", default=100, type=int,
 help="number of epochs to train"
)
 parser.add_argument(
 "--scheduler", type=str, default="step1r_1_0.5",
 help="scheduler type",
)
 parser.add_argument(
 "--accumulation_steps", type=int, default=1,
 help="number of gradient accumulation steps",
)
 parser.add_argument(
 "--save_every", type=int, default=1,
 help="how frequently to save a model",
)
 parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
```

```
)
parser.add_argument(
 "--weight_decay", type=float, default=1e-5,
 help="weight decay"
)
parser.add_argument(
 "--dropout", type=float, default=0.0,
 help="internal dropout"
)
parser.add_argument(
 "--output_dropout", type=float, default=0.0,
 help="output dropout"
)
parser.add_argument(
 "--p_mixup", type=float, default=0.0,
 help="probability of the mixup augmentation"
)
parser.add_argument(
 "--switch_off_augmentations_on", type=int, default=20,
 help="on which epoch to remove augmentations"
)
parser.add_argument(
 "--optimizer", type=str, required=True,
 help="which optimizer to use",
 choices=("adam", "momentum")
)
parser.add_argument(
 "--folds", type=int, required=True, nargs="+",
 help="which folds to use"
)
parser.add_argument(
 "--n_folds", type=int, default=4,
 help="number of folds"
)
parser.add_argument(
 "--kfold_seed", type=int, default=42,
 help="kfold seed"
)
parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)
parser.add_argument(
 "--label", type=str, default="finetuned_hierarchical_cnn_classifier",
 help="optional label",
)
args = parser.parse_args()

class_map = load_json(args.classmap)

pretrained = Experiment(resume_from=args.pretrained_model)

with Experiment({
 "network": {
 "num_conv_blocks": pretrained.config.network.num_conv_bl
```

```

ocks,
 "start_deep_supervision_on": pretrained.config.network.s
tart_deep_supervision_on,
 "conv_base_depth": pretrained.config.network.conv_base_d
epth,
 "growth_rate": pretrained.config.network.growth_rate,
 "dropout": args.dropout,
 "output_dropout": args.output_dropout,
},
"data": {
 "_n_folds": args.n_folds,
 "_kfold_seed": args.kfold_seed,
 "n_fft": pretrained.config.data.n_fft,
 "hop_size": pretrained.config.data.hop_size,
 "_input_dim": pretrained.config.data.n_fft // 2 + 1,
 "_n_classes": len(class_map),
 "_holdout_size": args.holdout_size,
 "p_mixup": args.p_mixup
},
"train": {
 "accumulation_steps": args.accumulation_steps,
 "batch_size": args.batch_size,
 "learning_rate": args.lr,
 "scheduler": args.scheduler,
 "optimizer": args.optimizer,
 "epochs": args.epochs,
 "_save_every": args.save_every,
 "weight_decay": args.weight_decay,
 "switch_off_augmentations_on": args.switch_off_augmentat
ions_on,
 "_pretrained_experiment": args.pretrained_model,
 "_pretrained_fold": args.pretrained_fold,
},
"label": args.label
}) as experiment:

 config = experiment.config
 print()
 print(" ////////// CONFIG //////////")
 print(experiment.config)

 train_df = pd.read_csv(args.train_df)
 test_df = pd.read_csv(args.sample_submission)

 if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index
(drop=True)
 test_df = test_df.sample(
 min(args.max_samples, len(test_df))).reset_index(dro
p=True)

 if args.holdout_size:
 keep, holdout = train_test_split(
 np.arange(len(train_df)), test_size=args.holdout_siz
e,
 random_state=args.kfold_seed)

```

```

e) holdout_df = train_df.iloc[holdout].reset_index(drop=True)

train_df = train_df.iloc[keep].reset_index(drop=True)

splits = list(train_validation_data(
 train_df.fname, train_df.labels,
 config.data._n_folds, config.data._kfold_seed))

for fold in args.folds:

 print("\n\n ----- Fold {}\n".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
 }

 if torch.cuda.is_available() else {})

 experiment.register_directory("checkpoints")
 experiment.register_directory("predictions")

 train_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[train]],
 labels=[item.split(",") for item in train_df.labels.values[train]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 MixUp(p=args.p_mixup),
 STFT(n_fft=config.data.n_fft, hop_size=config.data.hop_size),
 DropFields(("audio", "filename", "sr")),
 RenameFields({"stft": "signal"})
]),
 clean_transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
]),
 shuffle=True,
 drop_last=True,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": math.log(STFT.
eps)})),
 **loader_kwargs
)

 valid_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],

```

```

 labels=[item.split(",") for item in train_df.labels.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 STFT(n_fft=config.data.n_fft, hop_size=config.data.hop_size),
 DropFields(("audio", "filename", "sr")),
 RenameFields({"stft": "signal"})
]),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": math.log(STFT.
eps))),
 **loader_kwargs
)

 model = HierarchicalCNNClassificationModel(experiment, device=args.device)
 # load pretrained model
 model.load_state_dict(
 torch.load(
 os.path.join(
 pretrained.checkpoints,
 "fold_{}".format(args.pretrained_fold),
 "best_model.pth"
)
)
)

 scores = model.fit_validate(
 train_loader, valid_loader,
 epochs=experiment.config.train.epochs, fold=fold,
 log_interval=args.log_interval
)

 best_metric = max(scores)
 experiment.register_result("fold{}.metric".format(fold),
best_metric)

 torch.save(
 model.state_dict(),
 os.path.join(
 experiment.checkpoints,
 "fold_{}".format(fold),
 "final_model.pth"
)
)

 # predictions
 model.load_best_model(fold)

 # validation

 val_preds = model.predict(valid_loader)
 val_predictions_df = pd.DataFrame(

```



```

 val_preds, columns=get_class_names_from_classmap(class_map))
 val_predictions_df["fname"] = train_df.fname[valid].values
 val_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "val_preds_fold_{}.csv".format(fold)
),
 index=False
)
 del val_predictions_df

 # test
 test_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.test_data_dir, fname)
 for fname in test_df.fname.values],
 transform=Compose([
 LoadAudio(),
 STFT(n_fft=config.data.n_fft, hop_size=config.data.hop_size),
 DropFields(("audio", "filename", "sr")),
 RenameFields({"stft": "signal"})
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": math.log(STFT.
eps))),
 **loader_kwargs
)

 test_preds = model.predict(test_loader)
 test_predictions_df = pd.DataFrame(
 test_preds, columns=get_class_names_from_classmap(class_map))
 test_predictions_df["fname"] = test_df.fname
 test_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
),
 index=False
)
 del test_predictions_df

 # holdout
 if args.holdout_size:
 holdout_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in holdout_df.fname.values],
 labels=[item.split(",") for item in holdout_

```

```

df.labels.values],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map),
 STFT(n_fft=config.data.n_fft, hop_size=c
onfig.data.hop_size),
 DropFields(("audio", "filename", "sr")),
 RenameFields({"stft": "signal"})
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": math.log(S
TFT.eps)})),
 **loader_kwargs
)

holdout_metric = model.evaluate(holdout_loader)
experiment.register_result(
 "fold{}.holdout_metric".format(fold), holdout_me
tric)

print("\nHoldout metric: {:.4f}".format(holdout_metr
ic))

if args.device == "cuda":
 torch.cuda.empty_cache()

global metric
if all(
 "fold{}".format(k) in experiment.results.to_dict()
 for k in range(config.data._n_folds)):
 val_df_files = [
 os.path.join(
 experiment.predictions,
 "val_preds_fold_{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

 val_predictions_df = pd.concat([
 pd.read_csv(file) for file in val_df_files]).reset_i
ndex(drop=True)

 labels = np.asarray([
 item["labels"] for item in SoundDataset(
 audio_files=train_df.fname.tolist(),
 labels=[item.split(",") for item in train_df.lab
els.values],
 transform=MapLabels(class_map)
)
])

 val_labels_df = pd.DataFrame(

```

```

 labels, columns=get_class_names_from_classmap(class_
map))
 val_labels_df["fname"] = train_df.fname

 assert set(val_predictions_df.fname) == set(val_labels_d
f.fname)

 val_predictions_df.sort_values(by="fname", inplace=True)
 val_labels_df.sort_values(by="fname", inplace=True)

 metric = lwlrapp(
 val_labels_df.drop("fname", axis=1).values,
 val_predictions_df.drop("fname", axis=1).values
)

 experiment.register_result("metric", metric)

submission

test_df_files = [
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

if all(os.path.isfile for file in test_df_files):
 test_dfs = [pd.read_csv(file) for file in test_df_files]
 submission_df = pd.DataFrame({"fname": test_dfs[0].fname
.values})
 for c in get_class_names_from_classmap(class_map):
 submission_df[c] = np.mean([d[c].values for d in tes
t_dfs], axis=0)
 submission_df.to_csv(
 os.path.join(experiment.predictions, "submission.csv
"), index=False)=====

```

Apache License  
 Version 2.0, January 2004  
<http://www.apache.org/licenses/>

## TERMS AND CONDITIONS FOR USE, REPRODUCTION, AND DISTRIBUTION

### 1. Definitions.

"License" shall mean the terms and conditions for use, reproduction, and distribution as defined by Sections 1 through 9 of this document.

"Licensor" shall mean the copyright owner or entity authorized by the copyright owner that is granting the License.

"Legal Entity" shall mean the union of the acting entity and all

other entities that control, are controlled by, or are under common control with that entity. For the purposes of this definition, "control" means (i) the power, direct or indirect, to cause the direction or management of such entity, whether by contract or otherwise, or (ii) ownership of fifty percent (50%) or more of the outstanding shares, or (iii) beneficial ownership of such entity.

"You" (or "Your") shall mean an individual or Legal Entity exercising permissions granted by this License.

"Source" form shall mean the preferred form for making modifications, including but not limited to software source code, documentation source, and configuration files.

"Object" form shall mean any form resulting from mechanical transformation or translation of a Source form, including but not limited to compiled object code, generated documentation, and conversions to other media types.

"Work" shall mean the work of authorship, whether in Source or Object form, made available under the License, as indicated by a copyright notice that is included in or attached to the work (an example is provided in the Appendix below).

"Derivative Works" shall mean any work, whether in Source or Object form, that is based on (or derived from) the Work and for which the editorial revisions, annotations, elaborations, or other modifications represent, as a whole, an original work of authorship. For the purposes of this License, Derivative Works shall not include works that remain separable from, or merely link (or bind by name) to the interfaces of, the Work and Derivative Works thereof.

"Contribution" shall mean any work of authorship, including the original version of the Work and any modifications or additions

to that Work or Derivative Works thereof, that is intentionally submitted to Licensor for inclusion in the Work by the copyright owner or by an individual or Legal Entity authorized to submit on behalf of the copyright owner. For the purposes of this definition, "submitted" means any form of electronic, verbal, or written communication sent to the Licensor or its representatives, including but not limited to communication on electronic mailing lists, source code control systems, and issue tracking systems that are managed by, or on behalf of, the Licensor for the purpose of discussing and improving the Work, but excluding communication that is conspicuously marked or otherwise designated in writing by the copyright owner as "Not a Contribution."

"Contributor" shall mean Licensor and any individual or Legal Entity on behalf of whom a Contribution has been received by Licensor and subsequently incorporated within the Work.

2. Grant of Copyright License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable copyright license to reproduce, prepare Derivative Works of, publicly display, publicly perform, sublicense, and distribute the Work and such Derivative Works in Source or Object form.

3. Grant of Patent License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable (except as stated in this section) patent license to make, have made, use, offer to sell, sell, import, and otherwise transfer the Work, where such license applies only to those patent claims licensable by such Contributor that are necessarily infringed by their Contribution(s) alone or by combination of their Contribut

ion(s)  
with the Work to which such Contribution(s) was submitted.  
If You  
institute patent litigation against any entity (including  
a  
cross-claim or counterclaim in a lawsuit) alleging that th  
e Work  
or a Contribution incorporated within the Work constitutes  
direct  
or contributory patent infringement, then any patent licen  
ses  
granted to You under this License for that Work shall term  
inate  
as of the date such litigation is filed.

4. Redistribution. You may reproduce and distribute copies of  
the  
Work or Derivative Works thereof in any medium, with or wi  
thout  
modifications, and in Source or Object form, provided that  
You  
meet the following conditions:

- (a) You must give any other recipients of the Work or  
Derivative Works a copy of this License; and
- (b) You must cause any modified files to carry prominent n  
otices  
stating that You changed the files; and
- (c) You must retain, in the Source form of any Derivative  
Works  
that You distribute, all copyright, patent, trademark,  
and  
attribution notices from the Source form of the Work,  
excluding those notices that do not pertain to any par  
t of  
the Derivative Works; and
- (d) If the Work includes a "NOTICE" text file as part of i  
ts  
distribution, then any Derivative Works that You distr  
ibute must  
include a readable copy of the attribution notices con  
tained  
within such NOTICE file, excluding those notices that  
do not  
pertain to any part of the Derivative Works, in at lea  
st one  
of the following places: within a NOTICE text file dis  
tributed  
as part of the Derivative Works; within the Source for  
m or  
documentation, if provided along with the Derivative W  
orks; or,  
within a display generated by the Derivative Works, if

and  
contents  
and  
ution  
longside  
ovided  
strued

wherever such third-party notices normally appear. The  
of the NOTICE file are for informational purposes only  
do not modify the License. You may add Your own attrib  
notices within Derivative Works that You distribute, a  
or as an addendum to the NOTICE text from the Work, pr  
that such additional attribution notices cannot be con  
as modifying the License.

ions and  
itions  
ns, or  
e,  
ies with

You may add Your own copyright statement to Your modificat  
may provide additional or different license terms and cond  
for use, reproduction, or distribution of Your modificatio  
for any such Derivative Works as a whole, provided Your us  
reproduction, and distribution of the Work otherwise compl  
the conditions stated in this License.

5. Submission of Contributions. Unless You explicitly state o  
therwise,  
the Work  
ions of

any Contribution intentionally submitted for inclusion in  
by You to the Licensor shall be under the terms and condit  
this License, without any additional terms or conditions.  
Notwithstanding the above, nothing herein shall supersede  
or modify  
the terms of any separate license agreement you may have e  
xecuted  
with Licensor regarding such Contributions.

6. Trademarks. This License does not grant permission to use  
the trade  
Licensor,  
cribing the  
CE file.

names, trademarks, service marks, or product names of the  
except as required for reasonable and customary use in des  
origin of the Work and reproducing the content of the NOTI

7. Disclaimer of Warranty. Unless required by applicable law  
or  
S,  
ss or

agreed to in writing, Licensor provides the Work (and each  
Contributor provides its Contributions) on an "AS IS" BASI  
WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either expre

implied, including, without limitation, any warranties or conditions of TITLE, NON-INFRINGEMENT, MERCHANTABILITY, or FITNESS FOR A PARTICULAR PURPOSE. You are solely responsible for determining the appropriateness of using or redistributing the Work and assume any risks associated with Your exercise of permissions under this License.

8. Limitation of Liability. In no event and under no legal theory, whether in tort (including negligence), contract, or otherwise, unless required by applicable law (such as deliberate and grossly negligent acts) or agreed to in writing, shall any Contributor be liable to You for damages, including any direct, indirect, special, incidental, or consequential damages of any character arising as a result of this License or out of the use or inability to use the Work (including but not limited to damages for loss of goodwill, work stoppage, computer failure or malfunction, or any and all other commercial damages or losses), even if such Contributor has been advised of the possibility of such damages.

9. Accepting Warranty or Additional Liability. While redistributing the Work or Derivative Works thereof, You may choose to offer, and charge a fee for, acceptance of support, warranty, indemnity, or other liability obligations and/or rights consistent with this License. However, in accepting such obligations, You may act only on Your own behalf and on Your sole responsibility, not on behalf of any other Contributor, and only if You agree to indemnify, defend, and hold each Contributor harmless for any liability incurred by, or claims asserted against, such Contributor by reason of your accepting any such warranty or additional liability.

END OF TERMS AND CONDITIONS



APPENDIX: How to apply the Apache License to your work.

To apply the Apache License to your work, attach the following boilerplate notice, with the fields enclosed by brackets "[]" replaced with your own identifying information. (Don't include the brackets!) The text should be enclosed in the appropriate comment syntax for the file format. We also recommend that a file or class name and description of purpose be included on the same "printed page" as the copyright notice for easier identification within third-party archives.

Copyright [yyyy] [name of copyright owner]

Licensed under the Apache License, Version 2.0 (the "License");  
you may not use this file except in compliance with the License.  
You may obtain a copy of the License at

<http://www.apache.org/licenses/LICENSE-2.0>

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.  
See the License for the specific language governing permissions and limitations under the License.

=====

```
import argparse
import glob
from pathlib import Path
```

```
import pandas as pd
import numpy as np
import scipy.optimize
from scipy.stats import rankdata
from mag.utils import blue, green, bold
```

```
from ops.utils import lwlap
```

```
parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)
```

```
parser.add_argument(
 "--experiments", type=str, required=True, nargs="+",
```

```

 help="experiments to blend"
)
 parser.add_argument(
 "--train_df", type=str, required=True,
 help="path to train df"
)
 parser.add_argument(
 "--rankdata", action="store_true", default=False,
 help="whether to use ranks instead of raw scores"
)
 parser.add_argument(
 "--output_df", type=str, required=True,
 help="where to save test submission"
)
)

args = parser.parse_args()

n = len(args.experiments)

def load_predictions(experiment):
 prediction_files = (
 "experiments" / Path(experiment) / "predictions").glob("
val_preds*")
 dfs = [pd.read_csv(f) for f in prediction_files]
 df = pd.concat(dfs).reset_index(drop=True)
 df = df.sort_values(by="fname")
 df = df[sorted(df.columns.tolist())]
 return df

def to_ranks(values):
 return np.array([rankdata(r) for r in values])

predictions = [load_predictions(exp) for exp in args.experiments
]
class_cols = predictions[0].columns.drop("fname")
prediction_values = [p[class_cols].values for p in predictions]
if args.rankdata:
 prediction_values = [to_ranks(p) for p in prediction_values]

train_df = pd.read_csv(args.train_df)

def make_actual_labels(train_df):
 classname_to_idx = dict((c, i) for i, c in enumerate(class_c
ols))
 actual_labels = np.zeros((len(train_df), len(class_cols)), d
type=np.float32)
 for k in range(train_df.labels.values.size):
 for label in str(train_df.labels.values[k]).split(","):
 actual_labels[k, classname_to_idx[label]] = 1

```

```

 return actual_labels

actual_labels = make_actual_labels(train_df)

def constraints():
 A = np.ones(n)
 yield scipy.optimize.LinearConstraint(A=A, lb=0.01, ub=0.99)
 for k in range(n):
 A = np.zeros(n)
 A[k] = 1
 yield scipy.optimize.LinearConstraint(A=A, lb=0, ub=1)

def initial():
 return np.ones(n) / n

def target(alphas, *args):
 prediction = np.sum([a * p for a, p in zip(alphas, prediction_values)], axis=0)
 return -lwrap(actual_labels, prediction)

alphas = scipy.optimize.minimize(
 target,
 initial(),
 constraints=list(constraints()),
 method="COBYLA").x

print()
for experiment, alpha in zip(args.experiments, alphas):
 print("{}: {}".format(green(bold(experiment)), blue(bold(alpha))))

print()
print("Final lwrap:", bold(green(-target(alphas))))

def load_test_predictions(experiment):
 prediction_files = (
 "experiments" / Path(experiment) / "predictions").glob("test_preds*")
 dfs = [pd.read_csv(f) for f in prediction_files]
 dfs = [df.sort_values(by="fname") for df in dfs]
 return dfs

test_preds = []

for alpha, exp in zip(alphas, args.experiments):
 experiment_test_predictions = load_test_predictions(experiment)
 for p in experiment_test_predictions:

```

```

 if args.rankdata:
 test_preds.append(to_ranks(p[class_cols].values) * alpha)
 else:
 test_preds.append(p[class_cols].values * alpha)

test_preds = np.sum(test_preds, 0)

sub = pd.DataFrame(test_preds, columns=class_cols)
sub["fname"] = p.fname

sub.to_csv(args.output_df, index=False)=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
from mag.utils import green, bold
import mag

from datasets.sound_dataset import SoundDataset
from networks.classifiers import TwoDimensionalCNNClassification
Model
from ops.folds import train_validation_data_stratified
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation, FlipAudio, ShuffleAudio)
from ops.utils import load_json, get_class_names_from_classmap,
lwrap
from ops.padding import make_collate_fn

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--experiment", type=str, required=True,
 help="path to an experiment"
)

parser.add_argument(
 "--test_df", required=True, type=str,
 help="path to test dataframe"
)

parser.add_argument(
 "--output_df", required=True, type=str,
 help="where to save resulting dataframe"

```

```

)
parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data directory"
)
parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
parser.add_argument(
 "--batch_size", type=int, default=32,
 help="batch size used for prediction"
)
parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)

args = parser.parse_args()

class_map = load_json(args.classmap)
test_df = pd.read_csv(args.test_df)

with Experiment(resume_from=args.experiment) as experiment:
 config = experiment.config
 audio_transform = AudioFeatures(config.data.features)
 all_predictions = np.zeros(
 shape=(len(test_df), len(class_map)), dtype=np.float32)
 for fold in range(config.data._n_folds):
 print("\n\n ----- Fold {}\n".format(fold))
 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
 }
 if torch.cuda.is_available() else {})
 test_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.test_data_dir, fname)
 for fname in test_df.fname.values],
 labels=None,
 transform=Compose([
 LoadAudio(),
 audio_transform,

```

```

 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
]),
),
shuffle=False,
batch_size=args.batch_size,
collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
**loader_kwargs
)

model = TwoDimensionalCNNClassificationModel(
 experiment, device=args.device)
model.load_best_model(fold)
model.eval()

val_preds = model.predict(test_loader)

all_predictions += val_preds / config.data._n_folds

result = pd.DataFrame(
 all_predictions, columns=get_class_names_from_classmap(class_map))
result["fname"] = test_df.fname

result.to_csv(args.output_df, index=False)=====
import os
import gc
import argparse
import json
import math
from functools import partial

from scipy.sparse import csr_matrix
from scipy.stats import rankdata

import pandas as pd
import numpy as np

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--noisy_df", required=True, type=str,
 help="path to noisy dataframe"
)

parser.add_argument(
 "--noisy_predictions_df", required=True, type=str,
 help="path to noisy predictions"
)

parser.add_argument(

```

```

 "--output_df", required=True, type=str,
 help="where to save relabeled dataframe"
)
 parser.add_argument(
 "--mode", required=True, type=str,
 help="relabeling strategy"
)

 args = parser.parse_args()

 noisy_df = pd.read_csv(args.noisy_df)
 noisy_predictions_df = pd.read_csv(args.noisy_predictions_df)

 noisy_df.sort_values(by="fname", inplace=True)
 noisy_predictions_df.sort_values(by="fname", inplace=True)

 mode, *params = args.mode.split("_")

 class_cols = noisy_predictions_df.columns.drop("fname").values
 classname_to_idx = dict((c, i) for i, c in enumerate(class_cols))
 idx_to_classname = dict(enumerate(class_cols))
 noisy_labels = np.zeros((len(noisy_df), len(class_cols)), dtype=
np.float32)
 for k in range(noisy_df.labels.values.size):
 for label in str(noisy_df.labels.values[k]).split(","):
 noisy_labels[k, classname_to_idx[label]] = 1

def binary_to_labels(binary):
 labels = []
 for row in binary:
 labels.append(",".join(idx_to_classname[k] for k in nonz
ero(row)))

 return labels

def find_threshold(probs, expected_classes_per_sample):

 thresholds = np.linspace(0, 1, 10000)
 classes_per_sample = np.zeros_like(thresholds)

 for k in range(thresholds.size):
 c = (probs > thresholds[k]).sum(-1).mean()
 classes_per_sample[k] = c

 k = np.argmin(np.abs(classes_per_sample - expected_class
es_per_sample))

 return thresholds[k]

def nonzero(x):
 return np.nonzero(x)[0]

```

```

def merge_labels(first, second):
 merged = []
 for f, s in zip(first, second):
 m = set(f.split(",") | set(s.split(",")))
 if "" in m:
 m.remove("")
 merged.append(", ".join(m))

 return merged

def score_samples(y_true, y_score):
 scores = []

 y_true = csr_matrix(y_true)
 y_score = -y_score

 n_samples, n_labels = y_true.shape

 for i, (start, stop) in enumerate(zip(y_true.indptr, y_true.
indptr[1:])):
 relevant = y_true.indices[start:stop]

 if (relevant.size == 0 or relevant.size == n_labels):
 # If all labels are relevant or irrelevant, the score
e is also
 # equal to 1. The label ranking has no meaning.
 aux = 1.
 else:
 scores_i = y_score[i]
 rank = rankdata(scores_i, 'max')[relevant]
 L = rankdata(scores_i[relevant], 'max')
 aux = (L / rank).mean()

 scores.append(aux)

 return np.array(scores)

if mode == "fullmatch":
 expected_classes_per_sample, = params
 expected_classes_per_sample = float(expected_classes_per_sample)

 probs = noisy_predictions_df[class_cols].values
 threshold = find_threshold(probs, expected_classes_per_sample)
 binary = probs > threshold

 match = (binary == noisy_labels).all(-1)

 relabeled = noisy_df[match]

elif mode == "relabelall":

```



```

 expected_classes_per_sample, = params
 expected_classes_per_sample = float(expected_classes_per_sam
ple)

 probs = noisy_predictions_df[class_cols].values
 threshold = find_threshold(probs, expected_classes_per_sampl
e)
 binary = probs > threshold

 new_labels = binary_to_labels(binary)

 noisy_df.labels = new_labels
 noisy_df = noisy_df[noisy_df.labels != ""]

 relabeled = noisy_df

elif mode == "relabelall-replacen":

 expected_classes_per_sample, = params
 expected_classes_per_sample = float(expected_classes_per_sam
ple)

 probs = noisy_predictions_df[class_cols].values
 threshold = find_threshold(probs, expected_classes_per_sampl
e)
 binary = probs > threshold

 new_labels = pd.Series(binary_to_labels(binary))
 where_non_empty = (new_labels != "")
 noisy_df = noisy_df[where_non_empty]
 noisy_df.labels = new_labels[where_non_empty]

 relabeled = noisy_df

elif mode == "relabelall-merge":

 expected_classes_per_sample, = params
 expected_classes_per_sample = float(expected_classes_per_sam
ple)

 probs = noisy_predictions_df[class_cols].values
 threshold = find_threshold(probs, expected_classes_per_sampl
e)
 binary = probs > threshold

 new_labels = binary_to_labels(binary)
 noisy_df.labels = merge_labels(noisy_df.labels.values, new_l
abels)

 relabeled = noisy_df

elif mode == "scoring":

 topk, = params
 topk = int(topk)

```

```

probs = noisy_predictions_df[class_cols].values
scores = score_samples(noisy_labels, probs)

selection = np.argsort(-scores)[:topk]

relabeled = noisy_df.iloc[selection]

print("Relabeled df shape:", relabeled.shape)

relabeled.to_csv(args.output_df, index=False)
=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
import mag
from sklearn.model_selection import train_test_split

from datasets.sound_dataset import SoundDataset
from networks.classifiers import TwoDimensionalCNNClassification
Model
from ops.folds import train_validation_data, train_validation_data_stratified
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation, ShuffleAudio, CutOut, Identity)
from ops.utils import load_json, get_class_names_from_classmap,
lwlrap
from ops.padding import make_collate_fn

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)

parser.add_argument(
 "--train_data_dir", required=True, type=str,

```

```
 help="path to train data"
)
 parser.add_argument(
 "--noisy_train_df", type=str,
 help="path to noisy train dataframe (optional)"
)
 parser.add_argument(
 "--noisy_train_data_dir", type=str,
 help="path to noisy train data (optional)"
)
 parser.add_argument(
 "--share_noisy", action="store_true", default=False,
 help="whether to share noisy files across folds"
)
 parser.add_argument(
 "--resume", action="store_true", default=False,
 help="allow resuming even if experiment exists"
)
 parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data"
)
 parser.add_argument(
 "--sample_submission", required=True, type=str,
 help="path sample submission"
)
 parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
 parser.add_argument(
 "--log_interval", default=10, type=int,
 help="how frequently to log batch metrics"
 "in terms of processed batches"
)
 parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
 parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are samp
led"
)
 parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
 parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
 parser.add_argument(
 "--holdout_size", type=float, default=0.0,
 help="size of holdout set"
)
)
```

```
parser.add_argument(
 "--epochs", default=100, type=int,
 help="number of epochs to train"
)
parser.add_argument(
 "--scheduler", type=str, default="step1r_1_0.5",
 help="scheduler type",
)
parser.add_argument(
 "--accumulation_steps", type=int, default=1,
 help="number of gradient accumulation steps",
)
parser.add_argument(
 "--save_every", type=int, default=1,
 help="how frequently to save a model",
)
parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
parser.add_argument(
 "--aggregation_type", type=str, required=True,
 help="how to aggregate outputs",
 choices=("max", "rnn")
)
parser.add_argument(
 "--num_conv_blocks", type=int, default=5,
 help="number of conv blocks"
)
parser.add_argument(
 "--start_deep_supervision_on", type=int, default=2,
 help="from which layer to start aggregating features for cla
ssification"
)
parser.add_argument(
 "--conv_base_depth", type=int, default=64,
 help="base depth for conv layers"
)
parser.add_argument(
 "--growth_rate", type=float, default=2,
 help="how quickly to increase the number of units as a funct
ion of layer"
)
parser.add_argument(
 "--weight_decay", type=float, default=1e-5,
 help="weight decay"
)
parser.add_argument(
 "--output_dropout", type=float, default=0.0,
 help="output dropout"
)
parser.add_argument(
 "--p_mixup", type=float, default=0.0,
 help="probability of the mixup augmentation"
)
```

```
parser.add_argument(
 "--p_aug", type=float, default=0.0,
 help="probability of audio augmentation"
)
parser.add_argument(
 "--switch_off_augmentations_on", type=int, default=20,
 help="on which epoch to remove augmentations"
)
parser.add_argument(
 "--features", type=str, required=True,
 help="feature descriptor"
)
parser.add_argument(
 "--optimizer", type=str, required=True,
 help="which optimizer to use",
 choices=("adam", "momentum")
)
parser.add_argument(
 "--folds", type=int, required=True, nargs="+",
 help="which folds to use"
)
parser.add_argument(
 "--n_folds", type=int, default=4,
 help="number of folds"
)
parser.add_argument(
 "--kfold_seed", type=int, default=42,
 help="kfold seed"
)
parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)
parser.add_argument(
 "--label", type=str, default="2d_cnn",
 help="optional label",
)
args = parser.parse_args()

class_map = load_json(args.classmap)

audio_transform = AudioFeatures(args.features)

with Experiment({
 "network": {
 "num_conv_blocks": args.num_conv_blocks,
 "start_deep_supervision_on": args.start_deep_supervision
_on,
 "conv_base_depth": args.conv_base_depth,
 "growth_rate": args.growth_rate,
 "output_dropout": args.output_dropout,
 "aggregation_type": args.aggregation_type
 },
 "data": {
 "features": args.features,
 "_n_folds": args.n_folds,
```

```

 "_kfold_seed": args.kfold_seed,
 "_input_dim": audio_transform.n_features,
 "_n_classes": len(class_map),
 "_holdout_size": args.holdout_size,
 "p_mixup": args.p_mixup,
 "p_aug": args.p_aug,
 "max_audio_length": args.max_audio_length,
 "noisy": args.noisy_train_df is not None,
 "_train_df": args.train_df,
 "_train_data_dir": args.train_data_dir,
 "_noisy_train_df": args.noisy_train_df,
 "_noisy_train_data_dir": args.noisy_train_data_dir,
 "_share_noisy": args.share_noisy
 },
 "train": {
 "accumulation_steps": args.accumulation_steps,
 "batch_size": args.batch_size,
 "learning_rate": args.lr,
 "scheduler": args.scheduler,
 "optimizer": args.optimizer,
 "epochs": args.epochs,
 "_save_every": args.save_every,
 "weight_decay": args.weight_decay,
 "switch_off_augmentations_on": args.switch_off_augmentations_on
 },
 "label": args.label
}, implicit_resuming=args.resume) as experiment:

 config = experiment.config
 print()
 print(" ////////// CONFIG //////////")
 print(experiment.config)

 train_df = pd.read_csv(args.train_df)
 test_df = pd.read_csv(args.sample_submission)

 if args.noisy_train_df:
 noisy_train_df = pd.read_csv(args.noisy_train_df)

 if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index(drop=True)
 test_df = test_df.sample(
 min(args.max_samples, len(test_df))).reset_index(drop=True)

 if args.holdout_size:
 keep, holdout = train_test_split(
 np.arange(len(train_df)), test_size=args.holdout_size,
 random_state=args.kfold_seed)
 holdout_df = train_df.iloc[holdout].reset_index(drop=True)
 train_df = train_df.iloc[keep].reset_index(drop=True)

```

```

splits = list(train_validation_data_stratified(
 train_df.fname, train_df.labels, class_map,
 config.data._n_folds, config.data._kfold_seed))

if args.noisy_train_df:
 noisy_splits = list(train_validation_data(
 noisy_train_df.fname, noisy_train_df.labels,
 config.data._n_folds, config.data._kfold_seed))

for fold in args.folds:

 print("\n\n ----- Fold {}\n".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
 }

 if torch.cuda.is_available() else {})

 experiment.register_directory("checkpoints")
 experiment.register_directory("predictions")

 if args.noisy_train_df:

 noisy_train, noisy_valid = noisy_splits[fold]

 if config.data._share_noisy:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname
e)
 for fname in noisy_train_df.fname.values]
 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values]
 else:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname
e)
 for fname in noisy_train_df.fname.values[noi
sy_valid]]
 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values[noisy_valid]]
 else:
 noisy_audio_files = []
 noisy_labels = []

 train_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[train]] +
noisy_audio_files,
 labels=[
 item.split(",") for item in

```

```

 train_df.labels.values[train]] + noisy_label
s,
 is_noisy=[0] * len(train) + [1] * len(noisy_labels),
 transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_length),
 MapLabels(class_map=class_map),
 (
 ShuffleAudio(chunk_length=0.5, p=0.5)
 if config.network.aggregation_type != "random"
)
 else Identity()
),
 MixUp(p=args.p_mixup),
 AudioAugmentation(p=args.p_aug),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_length),
 MapLabels(class_map=class_map),
]),
 shuffle=True,
 drop_last=True,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform, "padding_value": m.padding_value}),
 **loader_kwargs
)

valid_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],
 labels=[item.split(",") for item in train_df.labels.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform, "padding_value": m.padding_value}),
 **loader_kwargs
)

model = TwoDimensionalCNNClassificationModel(

```



```

 experiment, device=args.device)

 scores = model.fit_validate(
 train_loader, valid_loader,
 epochs=experiment.config.train.epochs, fold=fold,
 log_interval=args.log_interval
)

 best_metric = max(scores)
 experiment.register_result("fold{}.metric".format(fold),
 best_metric)

 torch.save(
 model.state_dict(),
 os.path.join(
 experiment.checkpoints,
 "fold_{}".format(fold),
 "final_model.pth")
)

 # predictions
 model.load_best_model(fold)

 # validation

 val_preds = model.predict(valid_loader)
 val_predictions_df = pd.DataFrame(
 val_preds, columns=get_class_names_from_classmap(class_map))
 val_predictions_df["fname"] = train_df.fname[valid].values

 val_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "val_preds_fold{}.csv".format(fold)
),
 index=False
)
 del val_predictions_df

 # test
 test_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.test_data_dir, fname)
 for fname in test_df.fname.values],
 transform=Compose([
 LoadAudio(),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),

```

```

 **loader_kwargs
)

 test_preds = model.predict(test_loader)
 test_predictions_df = pd.DataFrame(
 test_preds, columns=get_class_names_from_classmap(class_map))
 test_predictions_df["fname"] = test_df.fname
 test_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "test_preds_fold{}.csv".format(fold)
),
 index=False
)
 del test_predictions_df

 # holdout
 if args.holdout_size:
 holdout_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in holdout_df.fname.values],
 labels=[item.split(",") for item in holdout_
df.labels.values],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
 **loader_kwargs
)

 holdout_metric = model.evaluate(holdout_loader)
 experiment.register_result(
 "fold{}.holdout_metric".format(fold), holdout_metric)

 print("\nHoldout metric: {:.4f}".format(holdout_metric))

 if args.device == "cuda":
 torch.cuda.empty_cache()

 # global metric

 if all(
 "fold{}".format(k) in experiment.results.to_dict()
 for k in range(config.data._n_folds)):

```

```

val_df_files = [
 os.path.join(
 experiment.predictions,
 "val_preds_fold_{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

val_predictions_df = pd.concat([
 pd.read_csv(file) for file in val_df_files]).reset_index(drop=True)

labels = np.asarray([
 item["labels"] for item in SoundDataset(
 audio_files=train_df.fname.tolist(),
 labels=[item.split(",") for item in train_df.labels.values],
 transform=MapLabels(class_map)
)
])

val_labels_df = pd.DataFrame(
 labels, columns=get_class_names_from_classmap(class_map))

val_labels_df["fname"] = train_df.fname

assert set(val_predictions_df.fname) == set(val_labels_df.fname)

val_predictions_df.sort_values(by="fname", inplace=True)
val_labels_df.sort_values(by="fname", inplace=True)

metric = lwlrap(
 val_labels_df.drop("fname", axis=1).values,
 val_predictions_df.drop("fname", axis=1).values
)

experiment.register_result("metric", metric)

submission

test_df_files = [
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

if all(os.path.isfile for file in test_df_files):
 test_dfs = [pd.read_csv(file) for file in test_df_files]
 submission_df = pd.DataFrame({"fname": test_dfs[0].fname.values})
 for c in get_class_names_from_classmap(class_map):
 submission_df[c] = np.mean([d[c].values for d in test_dfs])

```

```

t_dfs], axis=0)
 submission_df.to_csv(
 os.path.join(experiment.predictions, "submission.csv
"), index=False)=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
import mag
from sklearn.model_selection import train_test_split

from datasets.sound_dataset import SoundDataset
from networks.classifiers import CNNBackboneClassificationModel
from ops.folds import train_validation_data, train_validation_data_stratified
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation, ShuffleAudio, CutOut, Identity)
from ops.utils import load_json, get_class_names_from_classmap,
lwrap
from ops.padding import make_collate_fn

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)

parser.add_argument(
 "--train_data_dir", required=True, type=str,
 help="path to train data"
)

parser.add_argument(
 "--noisy_train_df", type=str,
 help="path to noisy train dataframe (optional)"
)

parser.add_argument(
 "--noisy_train_data_dir", type=str,
 help="path to noisy train data (optional)"

```

```
)
parser.add_argument(
 "--share_noisy", action="store_true", default=False,
 help="whether to share noisy files across folds"
)
parser.add_argument(
 "--resume", action="store_true", default=False,
 help="allow resuming even if experiment exists"
)
parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data"
)
parser.add_argument(
 "--sample_submission", required=True, type=str,
 help="path sample submission"
)
parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
parser.add_argument(
 "--log_interval", default=10, type=int,
 help="how frequently to log batch metrics"
 "in terms of processed batches"
)
parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are samp
led"
)
parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
parser.add_argument(
 "--holdout_size", type=float, default=0.0,
 help="size of holdout set"
)
parser.add_argument(
 "--epochs", default=100, type=int,
 help="number of epochs to train"
)
parser.add_argument(
 "--scheduler", type=str, default="steplr_1_0.5",
 help="scheduler type",
)
parser.add_argument(
```

```
 "--accumulation_steps", type=int, default=1,
 help="number of gradient accumulation steps",
)
 parser.add_argument(
 "--save_every", type=int, default=1,
 help="how frequently to save a model",
)
 parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
 parser.add_argument(
 "--backbone", type=str, required=True,
 help="which backbone to use",
 choices=("resnet18", "resnet34")
)
 parser.add_argument(
 "--weight_decay", type=float, default=1e-5,
 help="weight decay"
)
 parser.add_argument(
 "--output_dropout", type=float, default=0.0,
 help="output dropout"
)
 parser.add_argument(
 "--p_mixup", type=float, default=0.0,
 help="probability of the mixup augmentation"
)
 parser.add_argument(
 "--p_aug", type=float, default=0.0,
 help="probability of audio augmentation"
)
 parser.add_argument(
 "--switch_off_augmentations_on", type=int, default=20,
 help="on which epoch to remove augmentations"
)
 parser.add_argument(
 "--features", type=str, required=True,
 help="feature descriptor"
)
 parser.add_argument(
 "--optimizer", type=str, required=True,
 help="which optimizer to use",
 choices=("adam", "momentum")
)
 parser.add_argument(
 "--folds", type=int, required=True, nargs="+",
 help="which folds to use"
)
 parser.add_argument(
 "--n_folds", type=int, default=4,
 help="number of folds"
)
 parser.add_argument(
 "--kfold_seed", type=int, default=42,
```

```

 help="kfold seed"
)
 parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)
 parser.add_argument(
 "--label", type=str, default="backbone",
 help="optional label",
)
 args = parser.parse_args()

 class_map = load_json(args.classmap)

 audio_transform = AudioFeatures(args.features)

 with Experiment({
 "network": {
 "backbone": args.backbone,
 "output_dropout": args.output_dropout,
 },
 "data": {
 "features": args.features,
 "_n_folds": args.n_folds,
 "_kfold_seed": args.kfold_seed,
 "_input_dim": audio_transform.n_features,
 "_n_classes": len(class_map),
 "_holdout_size": args.holdout_size,
 "p_mixup": args.p_mixup,
 "p_aug": args.p_aug,
 "max_audio_length": args.max_audio_length,
 "noisy": args.noisy_train_df is not None,
 "_train_df": args.train_df,
 "_train_data_dir": args.train_data_dir,
 "_noisy_train_df": args.noisy_train_df,
 "_noisy_train_data_dir": args.noisy_train_data_dir,
 "_share_noisy": args.share_noisy
 },
 "train": {
 "accumulation_steps": args.accumulation_steps,
 "batch_size": args.batch_size,
 "learning_rate": args.lr,
 "scheduler": args.scheduler,
 "optimizer": args.optimizer,
 "epochs": args.epochs,
 "_save_every": args.save_every,
 "weight_decay": args.weight_decay,
 "switch_off_augmentations_on": args.switch_off_augmentations_on
 },
 "label": args.label
 }, implicit_resuming=args.resume) as experiment:

 config = experiment.config
 print()
 print(" ////////// CONFIG //////////")

```

```

print(experiment.config)

train_df = pd.read_csv(args.train_df)
test_df = pd.read_csv(args.sample_submission)

if args.noisy_train_df:
 noisy_train_df = pd.read_csv(args.noisy_train_df)

if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index(
(drop=True)
 test_df = test_df.sample(
 min(args.max_samples, len(test_df))).reset_index(drop=True)

if args.holdout_size:
 keep, holdout = train_test_split(
 np.arange(len(train_df)), test_size=args.holdout_size,
 random_state=args.kfold_seed)
 holdout_df = train_df.iloc[holdout].reset_index(drop=True)
 train_df = train_df.iloc[keep].reset_index(drop=True)

splits = list(train_validation_data_stratified(
 train_df.fname, train_df.labels, class_map,
 config.data._n_folds, config.data._kfold_seed))

if args.noisy_train_df:
 noisy_splits = list(train_validation_data(
 noisy_train_df.fname, noisy_train_df.labels,
 config.data._n_folds, config.data._kfold_seed))

for fold in args.folds:
 print("\n\n ----- Fold {}\n".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
 }
 if torch.cuda.is_available() else {})

 experiment.register_directory("checkpoints")
 experiment.register_directory("predictions")

 if args.noisy_train_df:
 noisy_train, noisy_valid = noisy_splits[fold]

 if config.data._share_noisy:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname)
 for fname in noisy_train_df.fname.values]

```



```

 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values]
 else:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname
e)
 for fname in noisy_train_df.fname.values[noi
sy_valid]]
 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values[noisy_valid]]
 else:
 noisy_audio_files = []
 noisy_labels = []

 train_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[train]] +
noisy_audio_files,
 labels=[
 item.split(",") for item in
 train_df.labels.values[train]] + noisy_label
s,
 is_noisy=[0] * len(train) + [1] * len(noisy_labe
ls),
 transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_le
ngth),
 MapLabels(class_map=class_map),
 ShuffleAudio(chunk_length=0.5, p=0.5),
 MixUp(p=args.p_mixup),
 AudioAugmentation(p=args.p_aug),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_le
ngth),
 MapLabels(class_map=class_map),
]),
 shuffle=True,
 drop_last=True,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transfor
m.padding_value})),
 **loader_kwargs
)

 valid_loader = torch.utils.data.DataLoader(
 SoundDataset(

```

```

 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],
 labels=[item.split(",") for item in train_df.labels.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
 **loader_kwargs
)

model = CNNBackboneClassificationModel(experiment, device=args.device)

scores = model.fit_validate(
 train_loader, valid_loader,
 epochs=experiment.config.train.epochs, fold=fold,
 log_interval=args.log_interval
)

best_metric = max(scores)
experiment.register_result("fold{}.metric".format(fold),
best_metric)

torch.save(
 model.state_dict(),
 os.path.join(
 experiment.checkpoints,
 "fold_{}".format(fold),
 "final_model.pth"
)
)

predictions
model.load_best_model(fold)

validation

val_preds = model.predict(valid_loader)
val_predictions_df = pd.DataFrame(
 val_preds, columns=get_class_names_from_classmap(class_map))
val_predictions_df["fname"] = train_df.fname[valid].values

val_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "val_preds_fold{}.csv".format(fold)
),

```

```

 index=False
)
 del val_predictions_df

 # test
 test_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.test_data_dir, fname)
 for fname in test_df.fname.values],
 transform=Compose([
 LoadAudio(),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.
padding_value}),
 **loader_kwargs
)

 test_preds = model.predict(test_loader)
 test_predictions_df = pd.DataFrame(
 test_preds, columns=get_class_names_from_classmap(class_map))
 test_predictions_df["fname"] = test_df.fname
 test_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
),
 index=False
)
 del test_predictions_df

 # holdout
 if args.holdout_size:
 holdout_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in holdout_df.fname.values],
 labels=[item.split(",") for item in holdout_
df.labels.values],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_tran

```

```

sform.padding_value)),
 **loader_kwargs
)

 holdout_metric = model.evaluate(holdout_loader)
 experiment.register_result(
 "fold{}.holdout_metric".format(fold), holdout_me
tric)

 print("\nHoldout metric: {:.4f}".format(holdout_metr
ic))

 if args.device == "cuda":
 torch.cuda.empty_cache()

global metric

if all(
 "fold{}".format(k) in experiment.results.to_dict()
 for k in range(config.data._n_folds)):

 val_df_files = [
 os.path.join(
 experiment.predictions,
 "val_preds_fold{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

 val_predictions_df = pd.concat([
 pd.read_csv(file) for file in val_df_files]).reset_i
ndex(drop=True)

 labels = np.asarray([
 item["labels"] for item in SoundDataset(
 audio_files=train_df.fname.tolist(),
 labels=[item.split(",") for item in train_df.lab
els.values],
 transform=MapLabels(class_map)
)
])

 val_labels_df = pd.DataFrame(
 labels, columns=get_class_names_from_classmap(class_
map))
 val_labels_df["fname"] = train_df.fname

 assert set(val_predictions_df.fname) == set(val_labels_d
f.fname)

 val_predictions_df.sort_values(by="fname", inplace=True)
 val_labels_df.sort_values(by="fname", inplace=True)

 metric = lwlap(
 val_labels_df.drop("fname", axis=1).values,
 val_predictions_df.drop("fname", axis=1).values

```

```

)

 experiment.register_result("metric", metric)

submission

test_df_files = [
 os.path.join(
 experiment.predictions,
 "test_preds_fold{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

if all(os.path.isfile for file in test_df_files):
 test_dfs = [pd.read_csv(file) for file in test_df_files]
 submission_df = pd.DataFrame({"fname": test_dfs[0].fname
 .values})
 for c in get_class_names_from_classmap(class_map):
 submission_df[c] = np.mean([d[c].values for d in tes
t_dfs], axis=0)
 submission_df.to_csv(
 os.path.join(experiment.predictions, "submission.csv
"), index=False)=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
import mag
from sklearn.model_selection import train_test_split

from ops.utils import load_json, get_class_names_from_classmap
from datasets.sound_dataset import SoundDataset
from networks.cpc import CPCModel
from ops.folds import train_validation_data
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation)
from ops.padding import make_collate_fn

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(

```

```
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

 parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)
 parser.add_argument(
 "--train_data_dir", required=True, type=str,
 help="path to train data"
)
 parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
 parser.add_argument(
 "--log_interval", default=10, type=int,
 help="how frequently to log batch metrics"
 "in terms of processed batches"
)
 parser.add_argument(
 "--proj_interval", default=10, type=int,
 help="how frequently to make projection in terms of epochs"
)
 parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
 parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are samp
led"
)
 parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
 parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
 parser.add_argument(
 "--epochs", default=100, type=int,
 help="number of epochs to train"
)
 parser.add_argument(
 "--scheduler", type=str, default="steplr_1_0.5",
 help="scheduler type",
)
 parser.add_argument(
 "--accumulation_steps", type=int, default=1,
 help="number of gradient accumulation steps",
)
 parser.add_argument(
 "--save_every", type=int, default=1,
 help="how frequently to save a model",
```

```
)
parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
parser.add_argument(
 "--n_encoder_layers", type=int, default=5,
 help="number of encoder layers"
)
parser.add_argument(
 "--conv_base_depth", type=int, default=64,
 help="base depth for conv layers"
)
parser.add_argument(
 "--context_size", type=int, default=64,
 help="context size for c network"
)
parser.add_argument(
 "--growth_rate", type=float, default=2,
 help="how quickly to increase the number of units as a function of layer"
)
parser.add_argument(
 "--prediction_steps", type=int, default=10,
 help="how many steps to predict in the future"
)
parser.add_argument(
 "--weight_decay", type=float, default=1e-5,
 help="weight decay"
)
parser.add_argument(
 "--p_aug", type=float, default=0.0,
 help="probability of audio augmentation"
)
parser.add_argument(
 "--switch_off_augmentations_on", type=int, default=20,
 help="on which epoch to remove augmentations"
)
parser.add_argument(
 "--features", type=str, required=True,
 help="feature descriptor"
)
parser.add_argument(
 "--optimizer", type=str, required=True,
 help="which optimizer to use",
 choices=("adam", "momentum")
)
parser.add_argument(
 "--folds", type=int, required=True, nargs="+",
 help="which folds to use"
)
parser.add_argument(
 "--n_folds", type=int, default=4,
 help="number of folds"
)
```

```

parser.add_argument(
 "--kfold_seed", type=int, default=42,
 help="kfold seed"
)
parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)
parser.add_argument(
 "--label", type=str, default="cpc",
 help="optional label",
)
args = parser.parse_args()

class_map = load_json(args.classmap)

audio_transform = AudioFeatures(args.features)

with Experiment({
 "network": {
 "n_encoder_layers": args.n_encoder_layers,
 "conv_base_depth": args.conv_base_depth,
 "growth_rate": args.growth_rate,
 "prediction_steps": args.prediction_steps,
 "context_size": args.context_size
 },
 "data": {
 "features": args.features,
 "_n_folds": args.n_folds,
 "_kfold_seed": args.kfold_seed,
 "_input_dim": audio_transform.n_features,
 "p_aug": args.p_aug,
 "max_audio_length": args.max_audio_length
 },
 "train": {
 "_proj_interval": args.proj_interval,
 "accumulation_steps": args.accumulation_steps,
 "batch_size": args.batch_size,
 "learning_rate": args.lr,
 "scheduler": args.scheduler,
 "optimizer": args.optimizer,
 "epochs": args.epochs,
 "_save_every": args.save_every,
 "weight_decay": args.weight_decay,
 "switch_off_augmentations_on": args.switch_off_augmentat
ions_on
 },
 "label": args.label
}) as experiment:

 config = experiment.config
 print()
 print(" ////////// CONFIG //////////")
 print(experiment.config)

 train_df = pd.read_csv(args.train_df)

```



```

 if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index(
(drop=True)

 splits = list(train_validation_data(
 train_df.fname, train_df.labels,
 config.data._n_folds, config.data._kfold_seed))

 for fold in args.folds:

 print("\n\n ----- Fold {}".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
}
 if torch.cuda.is_available() else {})

 experiment.register_directory("checkpoints")
 experiment.register_directory("predictions")

 train_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[train]],
 labels=[
 item.split(",") for item in
 train_df.labels.values[train]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 SampleLongAudio(max_length=args.max_audio_le
ngth),
 AudioAugmentation(p=args.p_aug),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=True,
 drop_last=True,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transfor
m.padding_value}),
 **loader_kwargs
)

 valid_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],
 labels=[
 item.split(",") for item in

```

```

 train_df.labels.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 SampleLongAudio(max_length=args.max_audio_le
length),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transfor
m.padding_value}),
 **loader_kwargs
)

model = CPCModel(experiment, device=args.device)

scores = model.fit_validate(
 train_loader, valid_loader,
 epochs=experiment.config.train.epochs, fold=fold,
 log_interval=args.log_interval
)

best_metric = max(scores)
experiment.register_result("fold{}.metric".format(fold),
best_metric)

torch.save(
 model.state_dict(),
 os.path.join(
 experiment.checkpoints,
 "fold_{}".format(fold),
 "final_model.pth")
)

predictions
model.load_best_model(fold)

if args.device == "cuda":
 torch.cuda.empty_cache()
=====
import os
import gc
import argparse
import json
import math
from functools import partial

import pandas as pd
import numpy as np
import torch
from mag.experiment import Experiment
import mag
from sklearn.model_selection import train_test_split

```

```
from datasets.sound_dataset import SoundDataset
from networks.classifiers import HierarchicalCNNClassificationModel
from ops.folds import train_validation_data, train_validation_data_stratified
from ops.transforms import (
 Compose, DropFields, LoadAudio,
 AudioFeatures, MapLabels, RenameFields,
 MixUp, SampleSegment, SampleLongAudio,
 AudioAugmentation, ShuffleAudio, CutOut, Identity)
from ops.utils import load_json, get_class_names_from_classmap,
lwlrap
from ops.padding import make_collate_fn

torch.manual_seed(42)
if torch.cuda.is_available():
 torch.cuda.manual_seed_all(42)

mag.use_custom_separator("-")

parser = argparse.ArgumentParser(
 formatter_class=argparse.ArgumentDefaultsHelpFormatter
)

parser.add_argument(
 "--train_df", required=True, type=str,
 help="path to train dataframe"
)
parser.add_argument(
 "--train_data_dir", required=True, type=str,
 help="path to train data"
)
parser.add_argument(
 "--noisy_train_df", type=str,
 help="path to noisy train dataframe (optional)"
)
parser.add_argument(
 "--noisy_train_data_dir", type=str,
 help="path to noisy train data (optional)"
)
parser.add_argument(
 "--share_noisy", action="store_true", default=False,
 help="whether to share noisy files across folds"
)
parser.add_argument(
 "--resume", action="store_true", default=False,
 help="allow resuming even if experiment exists"
)
parser.add_argument(
 "--test_data_dir", required=True, type=str,
 help="path to test data"
)
parser.add_argument(
 "--sample_submission", required=True, type=str,
 help="path sample submission"
```

```
)
parser.add_argument(
 "--classmap", required=True, type=str,
 help="path to class map json"
)
parser.add_argument(
 "--log_interval", default=10, type=int,
 help="how frequently to log batch metrics"
 "in terms of processed batches"
)
parser.add_argument(
 "--batch_size", type=int, default=64,
 help="minibatch size"
)
parser.add_argument(
 "--max_audio_length", type=int, default=10,
 help="max audio length in seconds. For longer clips are samp
led"
)
parser.add_argument(
 "--lr", default=0.01, type=float,
 help="starting learning rate"
)
parser.add_argument(
 "--max_samples", type=int,
 help="maximum number of samples to use"
)
parser.add_argument(
 "--holdout_size", type=float, default=0.0,
 help="size of holdout set"
)
parser.add_argument(
 "--epochs", default=100, type=int,
 help="number of epochs to train"
)
parser.add_argument(
 "--scheduler", type=str, default="steplr_1_0.5",
 help="scheduler type",
)
parser.add_argument(
 "--accumulation_steps", type=int, default=1,
 help="number of gradient accumulation steps",
)
parser.add_argument(
 "--save_every", type=int, default=1,
 help="how frequently to save a model",
)
parser.add_argument(
 "--device", type=str, required=True,
 help="whether to train on cuda or cpu",
 choices=("cuda", "cpu")
)
parser.add_argument(
 "--aggregation_type", type=str, required=True,
 help="how to aggregate outputs",
 choices=("max", "rnn")
)
```

```
)
parser.add_argument(
 "--num_conv_blocks", type=int, default=5,
 help="number of conv blocks"
)
parser.add_argument(
 "--start_deep_supervision_on", type=int, default=2,
 help="from which layer to start aggregating features for cla
ssification"
)
parser.add_argument(
 "--conv_base_depth", type=int, default=64,
 help="base depth for conv layers"
)
parser.add_argument(
 "--growth_rate", type=float, default=2,
 help="how quickly to increase the number of units as a funct
ion of layer"
)
parser.add_argument(
 "--weight_decay", type=float, default=1e-5,
 help="weight decay"
)
parser.add_argument(
 "--output_dropout", type=float, default=0.0,
 help="output dropout"
)
parser.add_argument(
 "--p_mixup", type=float, default=0.0,
 help="probability of the mixup augmentation"
)
parser.add_argument(
 "--p_aug", type=float, default=0.0,
 help="probability of audio augmentation"
)
parser.add_argument(
 "--switch_off_augmentations_on", type=int, default=20,
 help="on which epoch to remove augmentations"
)
parser.add_argument(
 "--features", type=str, required=True,
 help="feature descriptor"
)
parser.add_argument(
 "--optimizer", type=str, required=True,
 help="which optimizer to use",
 choices=("adam", "momentum")
)
parser.add_argument(
 "--folds", type=int, required=True, nargs="+",
 help="which folds to use"
)
parser.add_argument(
 "--n_folds", type=int, default=4,
 help="number of folds"
)
```

```

parser.add_argument(
 "--kfold_seed", type=int, default=42,
 help="kfold seed"
)
parser.add_argument(
 "--num_workers", type=int, default=4,
 help="number of workers for data loader",
)
parser.add_argument(
 "--label", type=str, default="1d_cnn",
 help="optional label",
)
args = parser.parse_args()

class_map = load_json(args.classmap)

audio_transform = AudioFeatures(args.features)

with Experiment({
 "network": {
 "num_conv_blocks": args.num_conv_blocks,
 "start_deep_supervision_on": args.start_deep_supervision
_on,
 "conv_base_depth": args.conv_base_depth,
 "growth_rate": args.growth_rate,
 "output_dropout": args.output_dropout,
 "aggregation_type": args.aggregation_type
 },
 "data": {
 "features": args.features,
 "_n_folds": args.n_folds,
 "_kfold_seed": args.kfold_seed,
 "_input_dim": audio_transform.n_features,
 "_n_classes": len(class_map),
 "_holdout_size": args.holdout_size,
 "p_mixup": args.p_mixup,
 "p_aug": args.p_aug,
 "max_audio_length": args.max_audio_length,
 "noisy": args.noisy_train_df is not None,
 "_train_df": args.train_df,
 "_train_data_dir": args.train_data_dir,
 "_noisy_train_df": args.noisy_train_df,
 "_noisy_train_data_dir": args.noisy_train_data_dir,
 "_share_noisy": args.share_noisy
 },
 "train": {
 "accumulation_steps": args.accumulation_steps,
 "batch_size": args.batch_size,
 "learning_rate": args.lr,
 "scheduler": args.scheduler,
 "optimizer": args.optimizer,
 "epochs": args.epochs,
 "_save_every": args.save_every,
 "weight_decay": args.weight_decay,
 "switch_off_augmentations_on": args.switch_off_augmentat
ions_on

```

```

 },
 "label": args.label
}, implicit_resuming=args.resume) as experiment:

 config = experiment.config
 print()
 print(" ////////// CONFIG //////////")
 print(experiment.config)

 train_df = pd.read_csv(args.train_df)
 test_df = pd.read_csv(args.sample_submission)

 if args.noisy_train_df:
 noisy_train_df = pd.read_csv(args.noisy_train_df)

 if args.max_samples:
 train_df = train_df.sample(args.max_samples).reset_index(
(drop=True)
 test_df = test_df.sample(
 min(args.max_samples, len(test_df))).reset_index(drop
p=True)

 if args.holdout_size:
 keep, holdout = train_test_split(
 np.arange(len(train_df)), test_size=args.holdout_siz
e,
 random_state=args.kfold_seed)
 holdout_df = train_df.iloc[holdout].reset_index(drop=Tru
e)
 train_df = train_df.iloc[keep].reset_index(drop=True)

 splits = list(train_validation_data_stratified(
 train_df.fname, train_df.labels, class_map,
 config.data._n_folds, config.data._kfold_seed))

 if args.noisy_train_df:
 noisy_splits = list(train_validation_data(
 noisy_train_df.fname, noisy_train_df.labels,
 config.data._n_folds, config.data._kfold_seed))

 for fold in args.folds:

 print("\n\n ----- Fold {} \n".format(fold))

 train, valid = splits[fold]

 loader_kwargs = (
 {"num_workers": args.num_workers, "pin_memory": True
}
 if torch.cuda.is_available() else {})

 experiment.register_directory("checkpoints")
 experiment.register_directory("predictions")

 if args.noisy_train_df:

```

```

noisy_train, noisy_valid = noisy_splits[fold]

if config.data._share_noisy:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname
e)
 for fname in noisy_train_df.fname.values]
 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values]
else:
 noisy_audio_files = [
 os.path.join(args.noisy_train_data_dir, fname
e)
 for fname in noisy_train_df.fname.values[noi
sy_valid]]
 noisy_labels = [
 item.split(",") for item in
 noisy_train_df.labels.values[noisy_valid]]
else:
 noisy_audio_files = []
 noisy_labels = []

train_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[train]] +
noisy_audio_files,
 labels=[
 item.split(",") for item in
 train_df.labels.values[train]] + noisy_label
s,
 is_noisy=[0] * len(train) + [1] * len(noisy_labe
ls),
 transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_le
ngth),
 MapLabels(class_map=class_map),
 (
 ShuffleAudio(chunk_length=0.5, p=0.5)
 if config.network.aggregation_type != "r
nn" else Identity()
),
 MixUp(p=args.p_mixup),
 AudioAugmentation(p=args.p_aug),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 clean_transform=Compose([
 LoadAudio(),
 SampleLongAudio(max_length=args.max_audio_le
ngth),
 MapLabels(class_map=class_map),
])
)

```



```

),
 shuffle=True,
 drop_last=True,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.
m.padding_value}),
 **loader_kwargs
)

 valid_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in train_df.fname.values[valid]],
 labels=[item.split(",") for item in train_df.labels.
els.values[valid]],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map=class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.
m.padding_value}),
 **loader_kwargs
)

 model = HierarchicalCNNClassificationModel(experiment, device=args.device)

 scores = model.fit_validate(
 train_loader, valid_loader,
 epochs=experiment.config.train.epochs, fold=fold,
 log_interval=args.log_interval
)

 best_metric = max(scores)
 experiment.register_result("fold{}.metric".format(fold),
best_metric)

 torch.save(
 model.state_dict(),
 os.path.join(
 experiment.checkpoints,
 "fold{}".format(fold),
 "final_model.pth"
)
)

 # predictions
 model.load_best_model(fold)

 # validation

```

```

val_preds = model.predict(valid_loader)
val_predictions_df = pd.DataFrame(
 val_preds, columns=get_class_names_from_classmap(class_map))
val_predictions_df["fname"] = train_df.fname[valid].values
val_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "val_preds_fold_{}.csv".format(fold)
),
 index=False
)
del val_predictions_df

test
test_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.test_data_dir, fname)
 for fname in test_df.fname.values],
 transform=Compose([
 LoadAudio(),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
])
),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
 **loader_kwargs
)

test_preds = model.predict(test_loader)
test_predictions_df = pd.DataFrame(
 test_preds, columns=get_class_names_from_classmap(class_map))
test_predictions_df["fname"] = test_df.fname
test_predictions_df.to_csv(
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
),
 index=False
)
del test_predictions_df

holdout
if args.holdout_size:
 holdout_loader = torch.utils.data.DataLoader(
 SoundDataset(
 audio_files=[
 os.path.join(args.train_data_dir, fname)
 for fname in holdout_df.fname.values],
 labels=[item.split(",") for item in holdout_

```

```

df.labels.values],
 transform=Compose([
 LoadAudio(),
 MapLabels(class_map),
 audio_transform,
 DropFields(("audio", "filename", "sr")),
]),
 shuffle=False,
 batch_size=config.train.batch_size,
 collate_fn=make_collate_fn({"signal": audio_transform.padding_value}),
 **loader_kwargs
)

 holdout_metric = model.evaluate(holdout_loader)
 experiment.register_result(
 "fold{}.holdout_metric".format(fold), holdout_metric)

 print("\nHoldout metric: {:.4f}".format(holdout_metric))

 if args.device == "cuda":
 torch.cuda.empty_cache()

global metric
if all(
 "fold{}".format(k) in experiment.results.to_dict()
 for k in range(config.data._n_folds)):

 val_df_files = [
 os.path.join(
 experiment.predictions,
 "val_preds_fold{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

 val_predictions_df = pd.concat([
 pd.read_csv(file) for file in val_df_files]).reset_index(drop=True)

 labels = np.asarray([
 item["labels"] for item in SoundDataset(
 audio_files=train_df.fname.tolist(),
 labels=[item.split(",") for item in train_df.labels.values],
 transform=MapLabels(class_map)
)
])

 val_labels_df = pd.DataFrame(
 labels, columns=get_class_names_from_classmap(class_map))

```

```

val_labels_df["fname"] = train_df.fname

assert set(val_predictions_df.fname) == set(val_labels_d
f.fname)

val_predictions_df.sort_values(by="fname", inplace=True)
val_labels_df.sort_values(by="fname", inplace=True)

metric = lwlrp(
 val_labels_df.drop("fname", axis=1).values,
 val_predictions_df.drop("fname", axis=1).values
)

experiment.register_result("metric", metric)

submission

test_df_files = [
 os.path.join(
 experiment.predictions,
 "test_preds_fold_{}.csv".format(fold)
)
 for fold in range(config.data._n_folds)
]

if all(os.path.isfile for file in test_df_files):
 test_dfs = [pd.read_csv(file) for file in test_df_files]
 submission_df = pd.DataFrame({"fname": test_dfs[0].fname
.values})
 for c in get_class_names_from_classmap(class_map):
 submission_df[c] = np.mean([d[c].values for d in tes
t_dfs], axis=0)
 submission_df.to_csv(
 os.path.join(experiment.predictions, "submission.csv
"), index=False)=====

```

```
import os
import math
import itertools
from collections import defaultdict, OrderedDict, deque

from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
import torchvision.utils
from tensorboardX import SummaryWriter
from torch.nn.functional import binary_cross_entropy_with_logits

from ops.training import OPTIMIZERS, make_scheduler, make_step
from networks.losses import binary_cross_entropy, focal_loss, lsep_loss
from ops.utils import plot_projection

class APCModel(nn.Module):

 def __init__(self, experiment, device="cuda"):
 super().__init__()

 self.device = device

 self.experiment = experiment
 self.config = experiment.config

 self.input_norm = nn.LayerNorm(
 (self.config.data._input_dim,), elementwise_affine=False)

 self.rnn = nn.LSTM(
 self.config.data._input_dim, self.config.network.rnn_size,
 num_layers=self.config.network.rnn_layers,
 batch_first=True
)

 self.output_norm = nn.LayerNorm((self.config.network.rnn_size,))

 self.prediction_transforms = torch.nn.ModuleList([
 torch.nn.Sequential(
 torch.nn.Linear(
 self.config.network.rnn_size,
 self.config.data._input_dim)
)
])
 for steps in range(self.config.network.prediction)
```

```
n_steps)
])

 self.to(self.device)

 def forward(self, signal):

 # signal = signal.permute(0, 2, 1)
 signal = self.input_norm(signal)
 # signal = signal.permute(0, 2, 1)

 output, state = self.rnn(signal)
 output = self.output_norm(output)

 losses = []
 predictions = []

 for step, affine in enumerate(self.prediction_transforms, start=1):

 shifted_output = output[:, :-step, :]
 shifted_signal = signal.detach()[:, step:, :]

 prediction = affine(shifted_output)
 predictions.append(prediction)

 loss = torch.abs(shifted_signal - prediction)
 loss = loss.sum(-1)
 loss = loss.mean()

 losses.append(loss)

 r = dict(
 losses=losses,
 output=output,
 predictions=predictions
)

 return r

 def add_scalar_summaries(
 self, losses, writer, global_step):

 # scalars
 for k, loss in enumerate(losses, start=1):
 writer.add_scalar("loss_{k}".format(k=k), loss,
global_step)

 def add_image_summaries(
 self, signal, output, predictions, global_step, writer,
to_plot=8):
```

```

 if len(signal) > to_plot:
 signal = signal[:to_plot]
 output = output[:to_plot]
 predictions = [p[:to_plot] for p in predictions]

 # signal
 image_grid = torchvision.utils.make_grid(
 signal.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
 writer.add_image("signal", image_grid, global_step)
 # output
 image_grid = torchvision.utils.make_grid(
 output.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
 writer.add_image("output", image_grid, global_step)

 for k, p in enumerate(predictions, start=1):
 image_grid = torchvision.utils.make_grid(
 p.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
 writer.add_image(
 "prediction_{k}".format(k=k), image_grid, gl
obal_step)

 def add_projection_summary(self, image, global_step, wri
ter, name="projection"):
 writer.add_image(name, image.transpose(2, 0, 1), glo
bal_step)

 def train_epoch(self, train_loader,
 epoch, log_interval, write_summary=True)
:

 self.train()

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n
"

 .format(epoch=epoch)
)

 history = deque(maxlen=30)

 self.optimizer.zero_grad()
 accumulated_loss = 0

 with tqdm(total=len(train_loader), ncols=80) as pb:

 for batch_idx, sample in enumerate(train_loader)

```

```
:

 self.global_step += 1

 make_step(self.scheduler, step=self.global_s
tep)

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device)
)

 outputs = self(signal)

 losses = outputs["losses"]

 loss = (
 sum(losses)
) / self.config.train.accumulation_steps

 loss.backward()
 accumulated_loss += loss

 if batch_idx % self.config.train.accumulatio
n_steps == 0:
 self.optimizer.step()
 accumulated_loss = 0
 self.optimizer.zero_grad()

 history.append(loss.item())

 pb.update()
 pb.set_description(
 "Loss: {:.4f}".format(
 np.mean(history))
)

 if batch_idx % log_interval == 0:
 self.add_scalar_summaries(
 [loss.item() for loss in losses],
 self.train_writer, self.global_step)

 if batch_idx == 0:
 self.add_image_summaries(
 signal,
 outputs["output"],
 outputs["predictions"],
 self.global_step, self.train_writer)

 def evaluate(self, loader, verbose=False, write_summary=
False, epoch=None):

 self.eval()
```



```

 valid_losses = [0 for _ in range(self.config.network
.prediction_steps)]

 all_outputs = []
 all_labels = []

 with torch.no_grad():
 for batch_idx, sample in enumerate(loader):

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device)
)

 outputs = self(signal)

 losses = outputs["losses"]

 multiplier = len(signal) / len(loader.datase
t)

 for k, loss in enumerate(losses):
 valid_losses[k] += loss.item() * multipl
ier

 all_outputs.extend(
 outputs["output"].data.cpu().numpy())
 all_labels.extend(labels.data.cpu().numpy())

 valid_loss = sum(valid_losses)

 all_labels = np.array(all_labels)

 if write_summary:
 self.add_scalar_summaries(
 valid_losses,
 writer=self.valid_writer, global_step=self.g
lobal_step
)
 if epoch % self.config.train._proj_interval == 0
:
 self.add_projection_summary(
 plot_projection(
 all_outputs, all_labels, frames_per_
example=5, newline=True),
 writer=self.valid_writer, global_step=se
lf.global_step,
 name="projection_output")

 if verbose:
 print("\nValidation loss: {:.4f}".format(valid_l

```

```

oss))

 return -valid_loss

def validation(self, valid_loader, epoch):
 return self.evaluate(
 valid_loader,
 verbose=True, write_summary=True, epoch=epoch)

def predict(self, loader):
 self.eval()

 all_class_probs = []

 with torch.no_grad():
 for sample in loader:

 signal = sample["signal"].to(self.device)

 outputs = self(signal)

 class_logits = outputs["class_logits"].squeeze()

 class_probs = torch.sigmoid(class_logits).data.cpu().numpy()
 all_class_probs.extend(class_probs)

 all_class_probs = np.asarray(all_class_probs)

 return all_class_probs

def fit_validate(self, train_loader, valid_loader, epochs, fold,
 log_interval=25):

 self.experiment.register_directory("summaries")
 self.train_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "train"
)
)
 self.valid_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "valid"
)
)

```

```

)

 os.makedirs(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold)),
 exist_ok=True
)

 self.global_step = 0
 self.make_optimizer(max_steps=len(train_loader) * ep
ochs)

 scores = []
 best_score = 0

 for epoch in range(epochs):

 make_step(self.scheduler, epoch=epoch)

 if epoch == self.config.train.switch_off_augment
ations_on:
 train_loader.dataset.transform.switch_off_au
gmentations()

 self.train_epoch(
 train_loader, epoch,
 log_interval, write_summary=True
)
 validation_score = self.validation(valid_loader,
epoch)
 scores.append(validation_score)

 if epoch % self.config.train._save_every == 0:
 print("\nSaving model on epoch", epoch)
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "model_on_epoch_{}.pth".format(epoch
)
)
)

 if validation_score > best_score:
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"

```

```

)
)
 best_score = validation_score

 return scores

def make_optimizer(self, max_steps):

 optimizer = OPTIMIZERS[self.config.train.optimizer]
 optimizer = optimizer(
 self.parameters(),
 self.config.train.learning_rate,
 weight_decay=self.config.train.weight_decay
)
 self.optimizer = optimizer
 self.scheduler = make_scheduler(
 self.config.train.scheduler, max_steps=max_steps
)(optimizer)

def load_best_model(self, fold):

 self.load_state_dict(
 torch.load(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
)

=====
import os
import math
import itertools
from collections import defaultdict, OrderedDict, deque

from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
import torch.utils.model_zoo as model_zoo
import torchvision.utils
from tensorboardX import SummaryWriter
from pretrainedmodels.models import resnet18, resnet34

from ops.training import OPTIMIZERS, make_scheduler, make_step
from networks.losses import binary_cross_entropy, focal_loss, lsep_loss
from ops.utils import lwlap, make_mel_filterbanks, is_mel,

```

```
is_stft, compute_torch_stft
```

```
class ConvLockedDropout(nn.Module):
 def __init__(self, dropout_rate=0.0):
 super().__init__()
 self.dropout_rate = dropout_rate

 def forward(self, x):
 if not self.training or not self.dropout_rate:
 return x

 n, s, t = x.size()

 m = torch.zeros(n, s, 1, device=x.device).bernoulli_(1 - self.dropout_rate)
 m = m.expand_as(x)
 return m * x
```

```
class ResnetBlock(nn.Module):

 def __init__(self, depth):
 super().__init__()

 self.conv1 = nn.Conv1d(depth, depth, kernel_size=1)
 self.bn1 = nn.BatchNorm1d(depth)
 self.conv2 = nn.Conv1d(depth, depth, kernel_size=3,
padding=1)
 self.bn2 = nn.BatchNorm1d(depth)
 self.conv3 = nn.Conv1d(depth, depth, kernel_size=1)
 self.bn3 = nn.BatchNorm1d(depth)
 self.prelu1 = nn.PReLU(depth)
 self.prelu2 = nn.PReLU(depth)
 self.prelu3 = nn.PReLU(depth)

 def forward(self, x):
 identity = x

 out = self.conv1(x)
 out = self.bn1(out)
 out = self.prelu1(out)

 out = self.conv2(out)
 out = self.bn2(out)
 out = self.prelu2(out)

 out = self.conv3(out)
 out = self.bn3(out)

 out += identity
 out = self.prelu3(out)
```

```
 return out
```

```
class ResnetBlock2d(nn.Module):
```

```
 def __init__(self, depth):
 super().__init__()
```

```
 self.conv1 = nn.Conv2d(depth, depth, kernel_size=1)
 self.bn1 = nn.BatchNorm2d(depth)
 self.conv2 = nn.Conv2d(depth, depth, kernel_size=3,
padding=1)
 self.bn2 = nn.BatchNorm2d(depth)
 self.conv3 = nn.Conv2d(depth, depth, kernel_size=1)
 self.bn3 = nn.BatchNorm2d(depth)
 self.prelu1 = nn.PReLU(depth)
 self.prelu2 = nn.PReLU(depth)
 self.prelu3 = nn.PReLU(depth)
```

```
 def forward(self, x):
 identity = x
```

```
 out = self.conv1(x)
 out = self.bn1(out)
 out = self.prelu1(out)
```

```
 out = self.conv2(out)
 out = self.bn2(out)
 out = self.prelu2(out)
```

```
 out = self.conv3(out)
 out = self.bn3(out)
```

```
 out += identity
 out = self.prelu3(out)
```

```
 return out
```

```
class HierarchicalCNNClassificationModel(nn.Module):
```

```
 def __init__(self, experiment, device="cuda"):
 super().__init__()
```

```
 self.device = device
```

```
 self.experiment = experiment
 self.config = experiment.config
```

```
 if is_mel(self.config.data.features):
 self.filterbanks = torch.from_numpy(
```

```

 make_mel_filterbanks(self.config.data.features)).to(self.device)

 self.conv_modules = torch.nn.ModuleList()
 self.rnn = torch.nn.ModuleList()

 total_depth = 0
 rnn_size = 128

 for k in range(self.config.network.num_conv_blocks):

 input_size = self.config.data._input_dim if not
k else depth
 depth = int(
 self.config.network.growth_rate ** k
 * self.config.network.conv_base_depth)

 if k >= self.config.network.start_deep_supervisi
on_on:
 if self.config.network.aggregation_type == "
max":
 total_depth += depth
 elif self.config.network.aggregation_type ==
"rnn":
 total_depth += rnn_size * 2
 self.rnn.append(
 nn.Sequential(
 nn.LayerNorm((depth,)),
 nn.GRU(
 depth, rnn_size, batch_first
=True, bidirectional=True)
)
)

 modules = [nn.BatchNorm1d(input_size)]
 modules.extend([
 nn.Conv1d(
 input_size,
 depth,
 kernel_size=3,
 padding=1
),
 nn.MaxPool1d(kernel_size=2, stride=2),
 nn.BatchNorm1d(depth),
 nn.PReLU(depth),
 ResnetBlock(depth)
])

 self.conv_modules.append(nn.Sequential(*modules))

 self.global_maxpool = nn.AdaptiveMaxPool1d(1)

```

```

 self.output_transform = nn.Sequential(
 nn.BatchNorm1d(total_depth),
 nn.Linear(total_depth, total_depth),
 nn.BatchNorm1d(total_depth),
 nn.PReLU(total_depth),
 nn.Dropout(p=self.config.network.output_dropout)
),
 nn.Linear(total_depth, self.config.data._n_classes)
)

 self.to(self.device)

 def forward(self, signal):

 if is_stft(self.config.data.features) or is_mel(self.config.data.features):
 signal = compute_torch_stft(
 signal.squeeze(-1),
 self.config.data.features
)

 if is_stft(self.config.data.features):
 signal = torch.log(signal + 1e-4)

 if is_mel(self.config.data.features):
 signal = nn.functional.conv1d(
 signal,
 self.filterbanks.unsqueeze(-1)
)
 signal = torch.log(signal + 1e-4)

 features = []

 h = signal
 for k, module in enumerate(self.conv_modules):
 h = module(h)
 if k >= self.config.network.start_deep_supervision_on:
 if self.config.network.aggregation_type == "max":
 features.append(self.global_maxpool(h).squeeze(-1))
 elif self.config.network.aggregation_type == "rnn":
 rnn_input = h.permute(0, 2, 1)
 outputs, state = self.rnns[k - self.config.network.start_deep_supervision_on](rnn_input)
 features.append(
 state.permute(1, 0, 2).contiguous()
)

```



```

view(rnn_input.size(0), -1))

 features = torch.cat(features, -1)

 class_logits = self.output_transform(features)

 r = dict(
 class_logits=class_logits
)

 return r

def add_scalar_summaries(
 self, loss, metric, writer, global_step):

 # scalars
 writer.add_scalar("loss", loss, global_step)
 writer.add_scalar("metric", metric, global_step)

 def add_image_summaries(self, signal, global_step, write
r, to_plot=8):

 if len(signal) > to_plot:
 signal = signal[:to_plot]

 # image
 image_grid = torchvision.utils.make_grid(
 signal.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
 writer.add_image("signal", image_grid, global_step)

def train_epoch(self, train_loader,
 epoch, log_interval, write_summary=True)
:

 self.train()

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n
"

 .format(epoch=epoch)
)

 history = deque(maxlen=30)

 self.optimizer.zero_grad()
 accumulated_loss = 0

 with tqdm(total=len(train_loader), ncols=80) as pb:

 for batch_idx, sample in enumerate(train_loader)

```

```

:
 self.global_step += 1

 make_step(self.scheduler, step=self.global_s
tep)

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
)

 outputs = self(signal)

 class_logits = outputs["class_logits"].squeeze()

 loss = (
 lsep_loss(
 class_logits,
 labels
)
) / self.config.train.accumulation_steps

 loss.backward()
 accumulated_loss += loss

 if batch_idx % self.config.train.accumulation
n_steps == 0:
 self.optimizer.step()
 accumulated_loss = 0
 self.optimizer.zero_grad()

 probs = torch.sigmoid(class_logits).data.cpu
().numpy()
 labels = labels.data.cpu().numpy()

 metric = lwlap(labels, probs)
 history.append(metric)

 pb.update()
 pb.set_description(
 "Loss: {:.4f}, Metric: {:.4f}".format(
 loss.item(), np.mean(history))
)

 if batch_idx % log_interval == 0:
 self.add_scalar_summaries(
 loss.item(), metric, self.train_writ
er, self.global_step)

 if batch_idx == 0:
 self.add_image_summaries(

```

```

 signal, self.global_step, self.train
_writer)

 def evaluate(self, loader, verbose=False, write_summary=
False, epoch=None):

 self.eval()

 valid_loss = 0

 all_class_probs = []
 all_labels = []

 with torch.no_grad():
 for batch_idx, sample in enumerate(loader):

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
)

 outputs = self(signal)

 class_logits = outputs["class_logits"].squeeze()

 loss = (
 lsep_loss(
 class_logits,
 labels,
)
).item()

 multiplier = len(labels) / len(loader.dataset)

 valid_loss += loss * multiplier

 class_probs = torch.sigmoid(class_logits).data.cpu().numpy()
 labels = labels.data.cpu().numpy()

 all_class_probs.extend(class_probs)
 all_labels.extend(labels)

 all_class_probs = np.asarray(all_class_probs)
 all_labels = np.asarray(all_labels)

 metric = lwlap(all_labels, all_class_probs)

 if write_summary:
 self.add_scalar_summaries(

```

```

 valid_loss,
 metric,
 writer=self.valid_writer, global_step=self.global_step
)

 if verbose:
 print("\nValidation loss: {:.4f}".format(validation_loss))
 print("Validation metric: {:.4f}".format(metric))

 return metric

def validation(self, valid_loader, epoch):
 return self.evaluate(
 valid_loader,
 verbose=True, write_summary=True, epoch=epoch)

def predict(self, loader):
 self.eval()

 all_class_probs = []

 with torch.no_grad():
 for sample in loader:

 signal = sample["signal"].to(self.device)

 outputs = self(signal)

 class_logits = outputs["class_logits"].squeeze()

 class_probs = torch.sigmoid(class_logits).data.cpu().numpy()
 all_class_probs.extend(class_probs)

 all_class_probs = np.asarray(all_class_probs)

 return all_class_probs

def fit_validate(self, train_loader, valid_loader, epochs, fold,
 log_interval=25):

 self.experiment.register_directory("summaries")
 self.train_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,

```

```

 "fold_{}".format(fold),
 "train"
)
)
self.valid_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "valid"
)
)

os.makedirs(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold)),
 exist_ok=True
)

self.global_step = 0
self.make_optimizer(max_steps=len(train_loader) * ep
ochs)

scores = []
best_score = 0

for epoch in range(epochs):

 make_step(self.scheduler, epoch=epoch)

 if epoch == self.config.train.switch_off_augment
ations_on:
 train_loader.dataset.transform.switch_off_au
gmentations()

 self.train_epoch(
 train_loader, epoch,
 log_interval, write_summary=True
)
 validation_score = self.validation(valid_loader,
epoch)
 scores.append(validation_score)

 if epoch % self.config.train._save_every == 0:
 print("\nSaving model on epoch", epoch)
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "model_on_epoch_{}.pth".format(epoch
)

```

```

)
)

 if validation_score > best_score:
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
 best_score = validation_score

 return scores

def make_optimizer(self, max_steps):

 optimizer = OPTIMIZERS[self.config.train.optimizer]
 optimizer = optimizer(
 self.parameters(),
 self.config.train.learning_rate,
 weight_decay=self.config.train.weight_decay
)
 self.optimizer = optimizer
 self.scheduler = make_scheduler(
 self.config.train.scheduler, max_steps=max_steps
)(optimizer)

def load_best_model(self, fold):

 self.load_state_dict(
 torch.load(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
)

class TwoDimensionalCNNClassificationModel(nn.Module):

 def __init__(self, experiment, device="cuda"):
 super().__init__()

 self.device = device

 self.experiment = experiment
 self.config = experiment.config

```

```

 if is_mel(self.config.data.features):
 self.filterbanks = torch.from_numpy(
 make_mel_filterbanks(self.config.data.features)).to(self.device)

 self.conv_modules = torch.nn.ModuleList()
 self.rnns = torch.nn.ModuleList()

 total_depth = 0

 for k in range(self.config.network.num_conv_blocks):

 input_size = 2 if not k else depth
 depth = int(
 self.config.network.growth_rate ** k
 * self.config.network.conv_base_depth)

 rnn_size = 128

 if k >= self.config.network.start_deep_supervision:
 if self.config.network.aggregation_type == "max":
 total_depth += depth
 elif self.config.network.aggregation_type == "rnn":
 total_depth += rnn_size * 2
 self.rnns.append(
 nn.Sequential(
 nn.LayerNorm((depth,)),
 nn.GRU(
 depth, rnn_size, batch_first=True,
 bidirectional=True)
)
)

 modules = [nn.BatchNorm2d(input_size)]
 modules.extend([
 nn.Conv2d(
 input_size,
 depth,
 kernel_size=3,
 padding=1
),
 nn.MaxPool2d(kernel_size=2, stride=2),
 nn.BatchNorm2d(depth),
 nn.PReLU(depth),
 ResnetBlock2d(depth)
])

 self.conv_modules.append(nn.Sequential(*modules))

```

```

self.global_maxpool = nn.AdaptiveMaxPool2d(1)

self.output_transform = nn.Sequential(
 nn.BatchNorm1d(total_depth),
 nn.Linear(total_depth, total_depth),
 nn.BatchNorm1d(total_depth),
 nn.PReLU(total_depth),
 nn.Dropout(p=self.config.network.output_dropout)
,
 nn.Linear(total_depth, self.config.data._n_classes)
)

self.to(self.device)

def _add_frequency_encoding(self, x):
 n, d, h, w = x.size()

 vertical = torch.linspace(-1, 1, h, device=x.device)
 .view(1, 1, -1, 1)
 vertical = vertical.repeat(n, 1, 1, w)

 x = torch.cat([x, vertical], dim=1)

 return x

def forward(self, signal):

 if is_stft(self.config.data.features) or is_mel(self.config.data.features):
 signal = compute_torch_stft(
 signal.squeeze(-1),
 self.config.data.features
)

 if is_stft(self.config.data.features):
 signal = torch.log(signal + 1e-4)

 if is_mel(self.config.data.features):
 signal = nn.functional.conv1d(
 signal,
 self.filterbanks.unsqueeze(-1)
)
 signal = torch.log(signal + 1e-4)

 signal = signal.unsqueeze(1)
 signal = self._add_frequency_encoding(signal)

 features = []

 h = signal

```



```

 for k, module in enumerate(self.conv_modules):
 h = module(h)
 if k >= self.config.network.start_deep_supervision_on:
 if self.config.network.aggregation_type == "max":
 features.append(self.global_maxpool(h).squeeze(-1).squeeze(-1))
 elif self.config.network.aggregation_type == "rnn":
 rnn_input = torch.mean(h, 2).permute(0, 2, 1)
 outputs, state = self.rnnns[k - self.config.network.start_deep_supervision_on](rnn_input)
 features.append(state.permute(1, 0, 2).contiguous().view(rnn_input.size(0), -1))

 features = torch.cat(features, -1)

 class_logits = self.output_transform(features)

 r = dict(
 class_logits=class_logits
)

 return r

 def add_scalar_summaries(
 self, loss, metric, writer, global_step):

 # scalars
 writer.add_scalar("loss", loss, global_step)
 writer.add_scalar("metric", metric, global_step)

 def add_histogram_summaries(
 self, losses, writer, global_step):

 writer.add_histogram("losses", np.array(losses), global_step=global_step)

 def add_image_summaries(self, signal, global_step, writer, to_plot=8):

 if len(signal) > to_plot:
 signal = signal[:to_plot]

 # image
 image_grid = torchvision.utils.make_grid(
 signal.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True

```

```

)
 writer.add_image("signal", image_grid, global_step)
def train_epoch(self, train_loader,
 epoch, log_interval, write_summary=True)
:
 self.train()

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n
"
 .format(epoch=epoch)
)

 training_losses = []

 history = deque(maxlen=30)

 self.optimizer.zero_grad()
 accumulated_loss = 0

 with tqdm(total=len(train_loader), ncols=80) as pb:
 for batch_idx, sample in enumerate(train_loader)
:
 self.global_step += 1

 make_step(self.scheduler, step=self.global_s
tep)

 signal, labels, is_noisy = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
,
 sample["is_noisy"].to(self.device).float
()
)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 loss = (
 lsep_loss(
 class_logits,
 labels,
 average=False
)
) / self.config.train.accumulation_steps

```

```

 training_losses.extend(loss.data.cpu().numpy
))

 loss = loss.mean()

 loss.backward()
 accumulated_loss += loss

 if batch_idx % self.config.train.accumulatio
n_steps == 0:
 self.optimizer.step()
 accumulated_loss = 0
 self.optimizer.zero_grad()

 probs = torch.sigmoid(class_logits).data.cpu
().numpy()
 labels = labels.data.cpu().numpy()

 metric = lwrap(labels, probs)
 history.append(metric)

 pb.update()
 pb.set_description(
 "Loss: {:.4f}, Metric: {:.4f}".format(
 loss.item(), np.mean(history)))

 if batch_idx % log_interval == 0:
 self.add_scalar_summaries(
 loss.item(), metric, self.train_writ
er, self.global_step)

 if batch_idx == 0:
 self.add_image_summaries(
 signal, self.global_step, self.train
_writer)

 self.add_histogram_summaries(
 training_losses, self.train_writer, self.global_
step)

 def evaluate(self, loader, verbose=False, write_summary=
False, epoch=None):

 self.eval()

 valid_loss = 0

 all_class_probs = []
 all_labels = []

 with torch.no_grad():
 for batch_idx, sample in enumerate(loader):

```

```

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 loss = (
 lsep_loss(
 class_logits,
 labels,
)
).item()

 multiplier = len(labels) / len(loader.datase
t)

 valid_loss += loss * multiplier

 class_probs = torch.sigmoid(class_logits).da
ta.cpu().numpy()
 labels = labels.data.cpu().numpy()

 all_class_probs.extend(class_probs)
 all_labels.extend(labels)

 all_class_probs = np.asarray(all_class_probs)
 all_labels = np.asarray(all_labels)

 metric = lwlapr(all_labels, all_class_probs)

 if write_summary:
 self.add_scalar_summaries(
 valid_loss,
 metric,
 writer=self.valid_writer, global_step=se
lf.global_step
)

 if verbose:
 print("\nValidation loss: {:.4f}".format(val
id_loss))
 print("Validation metric: {:.4f}".format(met
ric))

 return metric

def validation(self, valid_loader, epoch):
 return self.evaluate(
 valid_loader,

```

```

 verbose=True, write_summary=True, epoch=epoch)

def predict(self, loader, n_tta=1):
 self.eval()

 all_class_probs = []

 for k in range(n_tta):
 tta_probs = []

 with torch.no_grad():
 for sample in loader:

 signal = sample["signal"].to(self.device)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 class_probs = torch.sigmoid(class_logits)
 tta_probs.extend(class_probs.data.cpu().numpy())

 tta_probs = np.array(tta_probs)
 all_class_probs.append(tta_probs)

 all_class_probs = np.mean(all_class_probs, 0)

 return all_class_probs

def fit_validate(self, train_loader, valid_loader, epochs, fold,
 log_interval=25):

 self.experiment.register_directory("summaries")
 self.train_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "train"
)
)
 self.valid_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "valid"
)
)

```

```

)

 os.makedirs(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold)),
 exist_ok=True
)

 self.global_step = 0
 self.make_optimizer(max_steps=len(train_loader) * epochs)

 scores = []
 best_score = 0

 for epoch in range(epochs):

 make_step(self.scheduler, epoch=epoch)

 if epoch == self.config.train.switch_off_augmentations_on:
 train_loader.dataset.transform.switch_off_augmentations()

 self.train_epoch(
 train_loader, epoch,
 log_interval, write_summary=True
)
 validation_score = self.validation(valid_loader, epoch)
 scores.append(validation_score)

 if epoch % self.config.train._save_every == 0:
 print("\nSaving model on epoch", epoch)
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "model_on_epoch_{}.pth".format(epoch)
)
)

 if validation_score > best_score:
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)

```

```

)
)
 best_score = validation_score

 return scores

def make_optimizer(self, max_steps):

 optimizer = OPTIMIZERS[self.config.train.optimizer]
 optimizer = optimizer(
 self.parameters(),
 self.config.train.learning_rate,
 weight_decay=self.config.train.weight_decay
)
 self.optimizer = optimizer
 self.scheduler = make_scheduler(
 self.config.train.scheduler, max_steps=max_steps
)(optimizer)

def load_best_model(self, fold):

 self.load_state_dict(
 torch.load(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
)

)

class CNNBackboneClassificationModel(nn.Module):

 def __init__(self, experiment, device="cuda"):
 super().__init__()

 self.device = device

 self.experiment = experiment
 self.config = experiment.config

 if is_mel(self.config.data.features):
 self.filterbanks = torch.from_numpy(
 make_mel_filterbanks(self.config.data.features)
).to(self.device)

 self.input_norm = nn.BatchNorm2d(3)

 if self.config.network.backbone == "resnet18":
 self.backbone = resnet18(pretrained=None)

```

```

elif self.config.network.backbone == "resnet34":
 self.backbone = resnet34(pretrained=None)

self.global_maxpool = nn.AdaptiveMaxPool2d(1)

total_depth = self.backbone.last_linear.in_features

self.output_transform = nn.Sequential(
 nn.BatchNorm1d(total_depth),
 nn.Linear(total_depth, total_depth),
 nn.BatchNorm1d(total_depth),
 nn.PReLU(total_depth),
 nn.Dropout(p=self.config.network.output_dropout)
,
 nn.Linear(total_depth, self.config.data._n_classes)
)

self.to(self.device)

def forward(self, signal):

 if is_stft(self.config.data.features) or is_mel(self
.config.data.features):
 signal = compute_torch_stft(
 signal.squeeze(-1),
 self.config.data.features
)

 if is_stft(self.config.data.features):
 signal = torch.log(signal + 1e-4)

 if is_mel(self.config.data.features):
 signal = nn.functional.conv1d(
 signal,
 self.filterbanks.unsqueeze(-1)
)
 signal = torch.log(signal + 1e-4)

 signal = signal.unsqueeze(1)
 signal = signal.repeat(1, 3, 1, 1)
 signal = self.input_norm(signal)

 h = self.backbone.features(signal)

 features = self.global_maxpool(h).squeeze(-1).squeeze
e(-1)

 class_logits = self.output_transform(features)

 r = dict(
 class_logits=class_logits

```



```

)

 return r

def add_scalar_summaries(
 self, loss, metric, writer, global_step):

 # scalars
 writer.add_scalar("loss", loss, global_step)
 writer.add_scalar("metric", metric, global_step)

def add_histogram_summaries(
 self, losses, writer, global_step):

 writer.add_histogram("losses", np.array(losses), glo
bal_step=global_step)

def add_image_summaries(self, signal, global_step, write
r, to_plot=8):

 if len(signal) > to_plot:
 signal = signal[:to_plot]

 # image
 image_grid = torchvision.utils.make_grid(
 signal.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
 writer.add_image("signal", image_grid, global_step)

def train_epoch(self, train_loader,
 epoch, log_interval, write_summary=True)
:

 self.train()

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n
"
 .format(epoch=epoch)
)

 training_losses = []

 history = deque(maxlen=30)

 self.optimizer.zero_grad()
 accumulated_loss = 0

 with tqdm(total=len(train_loader), ncols=80) as pb:

 for batch_idx, sample in enumerate(train_loader)

```

```

:
 self.global_step += 1
 make_step(self.scheduler, step=self.global_s
tep)

 signal, labels, is_noisy = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
 ,
 sample["is_noisy"].to(self.device).float
)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 loss = (
 lsep_loss(
 class_logits,
 labels,
 average=False
)
) / self.config.train.accumulation_steps

 training_losses.extend(loss.data.cpu().numpy
 ())

 loss = loss.mean()

 loss.backward()
 accumulated_loss += loss

 if batch_idx % self.config.train.accumulatio
n_steps == 0:
 self.optimizer.step()
 accumulated_loss = 0
 self.optimizer.zero_grad()

 probs = torch.sigmoid(class_logits).data.cpu
 ().numpy()

 labels = labels.data.cpu().numpy()

 metric = lwlap(labels, probs)
 history.append(metric)

 pb.update()
 pb.set_description(
 "Loss: {:.4f}, Metric: {:.4f}".format(
 loss.item(), np.mean(history))
)

```

```

 if batch_idx % log_interval == 0:
 self.add_scalar_summaries(
 loss.item(), metric, self.train_writer, self.global_step)

 if batch_idx == 0:
 self.add_image_summaries(
 signal, self.global_step, self.train_writer)

 self.add_histogram_summaries(
 training_losses, self.train_writer, self.global_step)

 def evaluate(self, loader, verbose=False, write_summary=False, epoch=None):

 self.eval()

 valid_loss = 0

 all_class_probs = []
 all_labels = []

 with torch.no_grad():
 for batch_idx, sample in enumerate(loader):

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device).float()
)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 loss = (
 lsep_loss(
 class_logits,
 labels,
)
).item()

 multiplier = len(labels) / len(loader.dataset)

 valid_loss += loss * multiplier

 class_probs = torch.sigmoid(class_logits).data.cpu().numpy()
 labels = labels.data.cpu().numpy()

```

```

 all_class_probs.extend(class_probs)
 all_labels.extend(labels)

 all_class_probs = np.asarray(all_class_probs)
 all_labels = np.asarray(all_labels)

 metric = lwlap(all_labels, all_class_probs)

 if write_summary:
 self.add_scalar_summaries(
 valid_loss,
 metric,
 writer=self.valid_writer, global_step=self.global_step
)

 if verbose:
 print("\nValidation loss: {:.4f}".format(valid_loss))
 print("Validation metric: {:.4f}".format(metric))

 return metric

def validation(self, valid_loader, epoch):
 return self.evaluate(
 valid_loader,
 verbose=True, write_summary=True, epoch=epoch)

def predict(self, loader, n_tta=1):
 self.eval()

 all_class_probs = []

 for k in range(n_tta):
 tta_probs = []

 with torch.no_grad():
 for sample in loader:

 signal = sample["signal"].to(self.device)

 outputs = self(signal)

 class_logits = outputs["class_logits"]

 class_probs = torch.sigmoid(class_logits)
 tta_probs.extend(class_probs)

```

```

 tta_probs = np.array(tta_probs)
 all_class_probs.append(tta_probs)

 all_class_probs = np.mean(all_class_probs, 0)

 return all_class_probs

def fit_validate(self, train_loader, valid_loader, epochs, fold,
 log_interval=25):

 self.experiment.register_directory("summaries")
 self.train_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "train"
)
)
 self.valid_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "valid"
)
)

 os.makedirs(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold)),
 exist_ok=True
)

 self.global_step = 0
 self.make_optimizer(max_steps=len(train_loader) * epochs)

 scores = []
 best_score = 0

 for epoch in range(epochs):

 make_step(self.scheduler, epoch=epoch)

 if epoch == self.config.train.switch_off_augmentations_on:
 train_loader.dataset.transform.switch_off_augmentations()

```

```

 self.train_epoch(
 train_loader, epoch,
 log_interval, write_summary=True
)
 validation_score = self.validation(valid_loader,
epoch)
 scores.append(validation_score)

 if epoch % self.config.train._save_every == 0:
 print("\nSaving model on epoch", epoch)
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "model_on_epoch_{}.pth".format(epoch
)
)
)

 if validation_score > best_score:
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
 best_score = validation_score

 return scores

def make_optimizer(self, max_steps):

 optimizer = OPTIMIZERS[self.config.train.optimizer]
 optimizer = optimizer(
 self.parameters(),
 self.config.train.learning_rate,
 weight_decay=self.config.train.weight_decay
)
 self.optimizer = optimizer
 self.scheduler = make_scheduler(
 self.config.train.scheduler, max_steps=max_steps
)(optimizer)

def load_best_model(self, fold):

 self.load_state_dict(
 torch.load(
 os.path.join(
 self.experiment.checkpoints,

```

```

 "fold_{}".format(fold),
 "best_model.pth"
)
)
)

=====
import os
import math
import itertools
from collections import defaultdict, OrderedDict, deque

from tqdm import tqdm
import numpy as np
import torch
import torch.nn as nn
import torchvision.utils
from tensorboardX import SummaryWriter
from torch.nn.functional import binary_cross_entropy_with_logits

from ops.training import OPTIMIZERS, make_scheduler, make_step
from networks.losses import binary_cross_entropy, focal_loss, lsep_loss
from ops.utils import plot_projection

class CausalConv1d(nn.Module):

 def __init__(self, in_channels, out_channels, kernel_size,
 stride=1):
 super().__init__()
 self.kernel_size = kernel_size
 self.conv = nn.Conv1d(
 in_channels, out_channels, kernel_size,
 stride=stride, padding=kernel_size)

 def forward(self, x):
 x = self.conv(x)
 return x[:, :, :-self.kernel_size]

class CPCModel(nn.Module):

 def __init__(self, experiment, device="cuda"):
 super().__init__()

 self.device = device

 self.experiment = experiment
 self.config = experiment.config

```

```

 encoder_layers = []

 for k in range(self.config.network.n_encoder_layers)
:
 input_size = self.config.data._input_dim if not
k else depth
 depth = int(
 self.config.network.growth_rate ** k
 * self.config.network.conv_base_depth)
 modules = [nn.BatchNorm1d(input_size)] if not k
else []
 modules.extend([
 CausalConv1d(
 input_size,
 depth,
 kernel_size=3,
 stride=2
),
 nn.PReLU(depth)
])
 encoder_layers.extend(modules)

 encoder_layers.append(nn.BatchNorm1d(depth))

 self.encoder = nn.Sequential(*encoder_layers)

 self.context_network = nn.GRU(
 depth, self.config.network.context_size,
 num_layers=1,
 batch_first=True
)

 self.coupling_transforms = torch.nn.ModuleList([
 torch.nn.Sequential(
 torch.nn.Conv1d(
 self.config.network.context_size, depth,
kernel_size=1)
)
 for steps in range(self.config.network.predictio
n_steps)
])

 self.to(self.device)

 def forward(self, signal):

 signal = signal.permute(0, 2, 1)

 # z is (n, depth, steps)
 z = self.encoder(signal)
 # c is (n, context_size, steps)

```



```
c, state = self.context_network(z.permute(0, 2, 1))
c = c.permute(0, 2, 1)

losses = []

 for step, affine in enumerate(self.coupling_transforms, start=1):

 a = affine(c)

 # logits is (n, steps, steps)
 logits = torch.bmm(z.permute(0, 2, 1), a)

 labels = torch.eye(logits.size(2) - step, device
=z.device)
 labels = torch.nn.functional.pad(labels, (0, step, step, 0))
 labels = labels.unsqueeze(0).expand_as(logits)

 loss = binary_cross_entropy_with_logits(logits,
labels)
 losses.append(loss)

 r = dict(
 losses=losses,
 z=z,
 c=c
)

 return r

def add_scalar_summaries(
 self, losses, writer, global_step):

 # scalars
 for k, loss in enumerate(losses, start=1):
 writer.add_scalar("loss_{k}".format(k=k), loss,
global_step)

 def add_image_summaries(self, signal, c, z, global_step,
writer, to_plot=8):

 if len(c) > to_plot:
 signal = signal[:to_plot]
 c = c[:to_plot]
 z = z[:to_plot]

 # signal
 image_grid = torchvision.utils.make_grid(
 signal.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
```

```

writer.add_image("signal", image_grid, global_step)
z
image_grid = torchvision.utils.make_grid(
 z.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
writer.add_image("z", image_grid, global_step)
c
image_grid = torchvision.utils.make_grid(
 c.data.cpu().unsqueeze(1),
 normalize=True, scale_each=True
)
writer.add_image("c", image_grid, global_step)

def add_projection_summary(self, image, global_step, writer, name="projection"):
 writer.add_image(name, image.transpose(2, 0, 1), global_step)

def train_epoch(self, train_loader, epoch, log_interval, write_summary=True):
:
 self.train()

 print(
 "\n" + " " * 10 + "***** Epoch {epoch} *****\n"
 .format(epoch=epoch)
)

 history = deque(maxlen=30)

 self.optimizer.zero_grad()
 accumulated_loss = 0

 with tqdm(total=len(train_loader), ncols=80) as pb:
 for batch_idx, sample in enumerate(train_loader):
:
 self.global_step += 1

 make_step(self.scheduler, step=self.global_step)

 signal, labels = (
 sample["signal"].to(self.device),
 sample["labels"].to(self.device)
)

 outputs = self(signal)

```

```
 losses = outputs["losses"]

 loss = (
 sum(losses)
) / self.config.train.accumulation_steps

 loss.backward()
 accumulated_loss += loss

 if batch_idx % self.config.train.accumulation_steps == 0:
 self.optimizer.step()
 accumulated_loss = 0
 self.optimizer.zero_grad()

 history.append(loss.item())

 pb.update()
 pb.set_description(
 "Loss: {:.4f}".format(
 np.mean(history))
)

 if batch_idx % log_interval == 0:
 self.add_scalar_summaries(
 [loss.item() for loss in losses],
 self.train_writer, self.global_step)

 if batch_idx == 0:
 self.add_image_summaries(
 signal,
 outputs["c"].permute(0, 2, 1),
 outputs["z"].permute(0, 2, 1),
 self.global_step, self.train_writer)

 def evaluate(self, loader, verbose=False, write_summary=False, epoch=None):
 self.eval()

 valid_losses = [0 for _ in range(self.config.network.prediction_steps)]

 all_c = []
 all_z = []
 all_labels = []

 with torch.no_grad():
 for batch_idx, sample in enumerate(loader):

 signal, labels = (
 sample["signal"].to(self.device),
```

```

 sample["labels"].to(self.device)
)
 outputs = self(signal)
 losses = outputs["losses"]
 multiplier = len(signal) / len(loader.datase
t)
 for k, loss in enumerate(losses):
 valid_losses[k] += loss.item() * multipl
ier
 all_c.extend(
 outputs["c"].permute(0, 2, 1).data.cpu()
.numpy())
 all_z.extend(
 outputs["z"].permute(0, 2, 1).data.cpu()
.numpy())
 all_labels.extend(labels.data.cpu().numpy())
 valid_loss = sum(valid_losses)
 all_labels = np.array(all_labels)
 if write_summary:
 self.add_scalar_summaries(
 valid_losses,
 writer=self.valid_writer, global_step=self.g
lobal_step
)
 if epoch % self.config.train._proj_interval == 0
:
 self.add_projection_summary(
 plot_projection(
 all_c, all_labels, frames_per_exempl
e=5, newline=True),
 writer=self.valid_writer, global_step=se
lf.global_step,
 name="projection_c")
 self.add_projection_summary(
 plot_projection(all_z, all_labels, frame
s_per_example=5),
 writer=self.valid_writer, global_step=se
lf.global_step,
 name="projection_z")
 if verbose:
 print("\nValidation loss: {:.4f}".format(valid_l
oss))

```

```

 return -valid_loss

 def validation(self, valid_loader, epoch):
 return self.evaluate(
 valid_loader,
 verbose=True, write_summary=True, epoch=epoch)

 def predict(self, loader):
 self.eval()

 all_class_probs = []

 with torch.no_grad():
 for sample in loader:

 signal = sample["signal"].to(self.device)

 outputs = self(signal)

 class_logits = outputs["class_logits"].squeeze()

 class_probs = torch.sigmoid(class_logits).data.cpu().numpy()
 all_class_probs.extend(class_probs)

 all_class_probs = np.asarray(all_class_probs)

 return all_class_probs

 def fit_validate(self, train_loader, valid_loader, epochs, fold,
 log_interval=25):

 self.experiment.register_directory("summaries")
 self.train_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "train"
)
)
 self.valid_writer = SummaryWriter(
 log_dir=os.path.join(
 self.experiment.summaries,
 "fold_{}".format(fold),
 "valid"
)
)

```

```

 os.makedirs(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold)),
 exist_ok=True
)

 self.global_step = 0
 self.make_optimizer(max_steps=len(train_loader) * epochs)

 scores = []
 best_score = 0

 for epoch in range(epochs):

 make_step(self.scheduler, epoch=epoch)

 if epoch == self.config.train.switch_off_augmentations_on:
 train_loader.dataset.transform.switch_off_augmentations()

 self.train_epoch(
 train_loader, epoch,
 log_interval, write_summary=True
)
 validation_score = self.validation(valid_loader, epoch)
 scores.append(validation_score)

 if epoch % self.config.train._save_every == 0:
 print("\nSaving model on epoch", epoch)
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "model_on_epoch_{}.pth".format(epoch)
)
)

 if validation_score > best_score:
 torch.save(
 self.state_dict(),
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)

```

```

 best_score = validation_score

 return scores

def make_optimizer(self, max_steps):

 optimizer = OPTIMIZERS[self.config.train.optimizer]
 optimizer = optimizer(
 self.parameters(),
 self.config.train.learning_rate,
 weight_decay=self.config.train.weight_decay
)
 self.optimizer = optimizer
 self.scheduler = make_scheduler(
 self.config.train.scheduler, max_steps=max_steps
)(optimizer)

def load_best_model(self, fold):

 self.load_state_dict(
 torch.load(
 os.path.join(
 self.experiment.checkpoints,
 "fold_{}".format(fold),
 "best_model.pth"
)
)
)

=====
import torch
import torch.nn.functional as F

def focal_loss(input, target, focus=2.0, raw=True):

 if raw:
 input = torch.sigmoid(input)

 eps = 1e-7

 prob_true = input * target + (1 - input) * (1 - target)
 prob_true = torch.clamp(prob_true, eps, 1-eps)
 modulating_factor = (1.0 - prob_true).pow(focus)

 return (-modulating_factor * prob_true.log()).mean()

def binary_cross_entropy(input, target, raw=True):
 if raw:
 input = torch.sigmoid(input)
 return torch.nn.functional.binary_cross_entropy(input, t

```

```
arget)
```

```
def lsep_loss_stable(input, target, average=True):
 n = input.size(0)

 differences = input.unsqueeze(1) - input.unsqueeze(2)
 where_lower = (target.unsqueeze(1) < target.unsqueeze(2)
).float()

 differences = differences.view(n, -1)
 where_lower = where_lower.view(n, -1)

 max_difference, index = torch.max(differences, dim=1, ke
epdim=True)
 differences = differences - max_difference
 exps = differences.exp() * where_lower

 lsep = max_difference + torch.log(torch.exp(-max_differe
nce) + exps.sum(-1))

 if average:
 return lsep.mean()
 else:
 return lsep

def lsep_loss(input, target, average=True):
 differences = input.unsqueeze(1) - input.unsqueeze(2)
 where_different = (target.unsqueeze(1) < target.unsqueez
e(2)).float()

 exps = differences.exp() * where_different
 lsep = torch.log(1 + exps.sum(2).sum(1))

 if average:
 return lsep.mean()
 else:
 return lsep=====
import random

import numpy as np
import librosa
import scipy.signal

from sklearn.utils import gen_even_slices

def compute_stft(audio, window_size, hop_size, log=True, eps
=1e-4):
```



```
f, t, s = scipy.signal.stft(
 audio, nperseg=window_size, noverlap=hop_size)

s = np.abs(s)

if log:
 s = np.log(s + eps)

return s

def trim_audio(audio):
 audio, interval = librosa.effects.trim(audio, top_db=60)
 return audio

def read_audio(file):
 audio, sr = librosa.load(file, sr=None)
 return audio, sr

def mix_audio_and_labels(first_audio, second_audio, first_labels, second_labels):
 new_labels = np.clip(first_labels + second_labels, 0, 1)
 a = np.random.uniform(0.4, 0.6)
 shorter, longer = first_audio, second_audio
 if shorter.size == longer.size:
 return (shorter + longer) / 2, new_labels
 if first_audio.size > second_audio.size:
 shorter, longer = longer, shorter
 start = random.randint(0, longer.size - 1 - shorter.size)
 end = start + shorter.size
 longer *= a
 longer[start:end] += shorter * (1 - a)
 return longer, new_labels

def shuffle_audio(audio, chunk_length=0.5, sr=None):
 n_chunks = int((audio.size / sr) / chunk_length)
 if n_chunks in (0, 1):
 return audio
```

```

 slices = list(gen_even_slices(audio.size, n_chunks))
 random.shuffle(slices)

 shuffled = np.concatenate([audio[s] for s in slices])

 return shuffled

def cutout(audio, area=0.25):

 area = int(audio.size * area)

 start = random.randrange(audio.size)
 end = start + area

 audio[start:end] = 0

 return audio
=====
from sklearn.model_selection import KFold
from iterstrat.ml_stratifiers import MultilabelStratifiedKFold
import numpy as np

def train_validation_data(ids, labels, n_folds, seed):

 for train, valid in KFold(
 n_folds, shuffle=True, random_state=seed).split(ids,
 labels):
 yield train, valid

def train_validation_data_stratified(
 ids, labels, classmap, n_folds, seed):

 binary_labels = np.zeros(
 (len(labels), len(classmap)), dtype=np.float32)
 for k, item in enumerate(labels.values):
 for label in item.split(","):
 binary_labels[k, classmap[label]] = 1

 for train, valid in MultilabelStratifiedKFold(
 n_folds, shuffle=True, random_state=seed).split(ids,
 binary_labels):
 yield train, valid=====
import random
from copy import deepcopy

import numpy as np
from torch.utils.data.dataloader import default_collate

```

```

def make_collate_fn(padding_values):
 def _collate_fn(batch):
 for name, padding_value in padding_values.items():
 lengths = [len(sample[name]) for sample in batch]
 max_length = max(lengths)

 for n, size in enumerate(lengths):
 p = max_length - size
 if p:
 pad_width = [(0, p)] + [(0, 0)] * (batch
[n][name].ndim - 1)
 if padding_value == "edge":
 batch[n][name] = np.pad(
 batch[n][name], pad_width,
 mode="edge")
 else:
 batch[n][name] = np.pad(
 batch[n][name], pad_width,
 mode="constant", constant_values
=padding_value)

 return default_collate(batch)

 return _collate_fn

class BucketingSampler:
 def __init__(self, dataset, max_batch_elems, buckets):
 self.buckets = buckets
 self.dataset = dataset
 self.max_batch_elems = max_batch_elems

 self._create_batches()

 def _create_batches(self):
 self.n_bins = len(self.buckets)
 binned_sizes = np.digitize(self.dataset.lengths, self
.buckets)

 batches = []

 for bin_idx in range(1, self.n_bins):

```

```

 ids = np.nonzero(binned_sizes == bin_idx)[0]
 random.shuffle(ids)

 current_len = 0
 batch = []

 for id in ids:
 if current_len < self.max_batch_elems:
 batch.append(id)
 current_len += self.dataset.lengths[id]
 else:
 batches.append(batch)
 current_len = self.dataset.lengths[id]
 batch = [id]

 if batch:
 batches.append(batch)

 random.shuffle(batches)

 self.n_batches = len(batches)
 self.batches = batches

 def __iter__(self):
 return iter(self.batches)

 def __len__(self):
 return self.n_batches=====
from functools import partial

import numpy as np
import torch
from torch.optim import Optimizer
from torch.optim.lr_scheduler import StepLR, CosineAnnealing
LR, _LRScheduler

OPTIMIZERS = {
 "adam": partial(torch.optim.Adam, amsgrad=True),
 "momentum": partial(torch.optim.SGD, momentum=0.9, nestero
rov=True)
}

def make_scheduler(params, max_steps):

 name, *args = params.split("_")

 if name == "steplr":

 step_size, gamma = args
 step_size = int(step_size)

```

```

 gamma = float(gamma)

 return partial(StepLR, step_size=step_size, gamma=gamma)

elif name == "1cycle":

 min_lr, max_lr = args
 min_lr = float(min_lr)
 max_lr = float(max_lr)

 return partial(
 OneCycleScheduler, min_lr=min_lr, max_lr=max_lr,
 max_steps=max_steps)

def make_step(scheduler, epoch=None, step=None, val_score=None):

 if isinstance(scheduler, StepLR) and epoch is not None:
 scheduler.step(epoch)

 elif isinstance(scheduler, OneCycleScheduler) and step is not None:
 scheduler.step()

class CyclicLR:
 """Sets the learning rate of each parameter group according to
 cyclical learning rate policy (CLR). The policy cycles the learning
 rate between two boundaries with a constant frequency, as detailed in
 the paper `Cyclical Learning Rates for Training Neural Networks`.
 The distance between the two boundaries can be scaled on a per-iteration
 or per-cycle basis.

 Cyclical learning rate policy changes the learning rate after every batch.
 `batch_step` should be called after a batch has been used for training.
 To resume training, save `last_batch_iteration` and use it to instantiate `CycleLR`.

 This class has three built-in policies, as put forth in the paper:
 "triangular":
 A basic triangular cycle w/ no amplitude scaling.
 "triangular2":

```

A basic triangular cycle that scales initial amplitude by half each cycle.

"exp\_range":

A cycle that scales initial amplitude by  $\gamma^{(cycle\ iterations)}$  at each cycle iteration.

This implementation was adapted from the github repo: `bckenstler/CLR`\_

Args:

optimizer (Optimizer): Wrapped optimizer.

base\_lr (float or list): Initial learning rate which is the lower boundary in the cycle for eachparam groups

Default: 0.001

max\_lr (float or list): Upper boundaries in the cycle for each parameter group. Functionally, it defines the cycle amplitude ( $\max\_lr - \text{base\_lr}$ ).

The lr at any cycle is the sum of base\_lr and some scaling of the amplitude; therefore max\_lr may not actually be reached depending on scaling function. Default: 0.006

step\_size (int): Number of training iterations per half cycle. Authors suggest setting step\_size 2-8 x training iterations in epoch. Default: 200

mode (str): One of {triangular, triangular2, exp\_range}.

Values correspond to policies detailed above. If scale\_fn is not None, this argument is ignored.

Default: 'triangular'

gamma (float): Constant in 'exp\_range' scaling function:

$\gamma^{(cycle\ iterations)}$   
Default: 1.0

scale\_fn (function): Custom scaling policy defined by a single

argument lambda function, where  $0 \leq \text{scale\_fn}(x) \leq 1$  for all  $x \geq 0$ . mode parameter is ignored  
Default: None

scale\_mode (str): {'cycle', 'iterations'}.

Defines whether scale\_fn is evaluated on cycle number or cycle iterations (training iterations since start of cycle).  
Default: 'cycle'

last\_batch\_iteration (int): The index of the last batch

tch. Default: -1

Example:

```
>>> optimizer = torch.optim.SGD(model.parameters(),
lr=0.1, momentum=0.9)
>>> scheduler = torch.optim.CyclicLR(optimizer)
>>> data_loader = torch.utils.data.DataLoader(...)
>>> for epoch in range(10):
>>> for batch in data_loader:
>>> scheduler.batch_step()
>>> train_batch(...)
```

.. \_Cyclical Learning Rates for Training Neural Networks  
: <https://arxiv.org/abs/1506.01186>  
.. \_bckenstler/CLR: <https://github.com/bckenstler/CLR>  
"""

```
def __init__(self, optimizer, base_lr=1e-3, max_lr=6e-3,
 step_size=2000, mode='triangular', gamma=1.
,
 scale_fn=None, scale_mode='cycle', last_batch_
iteration=-1):

 if not isinstance(optimizer, Optimizer):
 raise TypeError('{} is not an Optimizer'.format(
 type(optimizer).__name__))
 self.optimizer = optimizer

 if isinstance(base_lr, list) or isinstance(base_lr,
tuple):
 if len(base_lr) != len(optimizer.param_groups):
 raise ValueError("expected {} base_lr, got {}".format(
 len(optimizer.param_groups), len(base_lr)
))
 self.base_lrs = list(base_lr)
 else:
 self.base_lrs = [base_lr] * len(optimizer.param_
groups)

 if isinstance(max_lr, list) or isinstance(max_lr, tu
ple):
 if len(max_lr) != len(optimizer.param_groups):
 raise ValueError("expected {} max_lr, got {}".format(
 len(optimizer.param_groups), len(max_lr)
))
 self.max_lrs = list(max_lr)
 else:
 self.max_lrs = [max_lr] * len(optimizer.param_gr
oups)
```

```

 self.step_size = step_size

 if mode not in ['triangular', 'triangular2', 'exp_range'] \
 and scale_fn is None:
 raise ValueError('mode is invalid and scale_fn is None')

 self.mode = mode
 self.gamma = gamma

 if scale_fn is None:
 if self.mode == 'triangular':
 self.scale_fn = self._triangular_scale_fn
 self.scale_mode = 'cycle'
 elif self.mode == 'triangular2':
 self.scale_fn = self._triangular2_scale_fn
 self.scale_mode = 'cycle'
 elif self.mode == 'exp_range':
 self.scale_fn = self._exp_range_scale_fn
 self.scale_mode = 'iterations'
 else:
 self.scale_fn = scale_fn
 self.scale_mode = scale_mode

 self.batch_step(last_batch_iteration + 1)
 self.last_batch_iteration = last_batch_iteration

 def batch_step(self, batch_iteration=None):
 if batch_iteration is None:
 batch_iteration = self.last_batch_iteration + 1
 self.last_batch_iteration = batch_iteration
 for param_group, lr in zip(self.optimizer.param_groups, self.get_lr()):
 param_group['lr'] = lr

 def _triangular_scale_fn(self, x):
 return 1.

 def _triangular2_scale_fn(self, x):
 return 1 / (2. ** (x - 1))

 def _exp_range_scale_fn(self, x):
 return self.gamma**(x)

 def get_lr(self):
 step_size = float(self.step_size)
 cycle = np.floor(1 + self.last_batch_iteration / (2
* step_size))
 x = np.abs(self.last_batch_iteration / step_size - 2
* cycle + 1)

```



```

 lrs = []
 param_lrs = zip(self.optimizer.param_groups, self.base_lrs, self.max_lrs)
 for param_group, base_lr, max_lr in param_lrs:
 base_height = (max_lr - base_lr) * np.maximum(0,
(1 - x))
 if self.scale_mode == 'cycle':
 lr = base_lr + base_height * self.scale_fn(cycle)
 else:
 lr = base_lr + base_height * self.scale_fn(self.last_batch_iteration)
 lrs.append(lr)
 return lrs

def annealing_linear(start, end, r):
 return start + r * (end - start)

def annealing_cos(start, end, r):
 cos_out = np.cos(np.pi * r) + 1
 return end + (start - end) / 2 * cos_out

class OneCycleScheduler:
 def __init__(
 self, optimizer,
 min_lr, max_lr,
 max_steps, annealing=annealing_linear):

 self.optimizer = optimizer
 self.min_lr = min_lr
 self.max_lr = max_lr
 self.max_steps = max_steps
 self.annealing = annealing
 self.epoch = -1

 def step(self):
 self.epoch += 1

 mid = int(round(self.max_steps * 0.3))

 if self.epoch < mid:
 r = self.epoch / mid
 lr = self.annealing(self.min_lr, self.max_lr, r)
 else:
 r = (self.epoch - mid) / (self.max_steps - mid)
 lr = self.annealing(self.max_lr, self.min_lr / 1e3, r)

 for param_group in self.optimizer.param_groups:
 param_group['lr'] = lr

```

```
=====
import random
import math
from functools import partial
import json

import pysndfx
import librosa
import numpy as np
import torch

from ops.audio import (
 read_audio, compute_stft, trim_audio, mix_audio_and_labels,
 shuffle_audio, cutout
)

SAMPLE_RATE = 44100

class Augmentation:
 """A base class for data augmentation transforms"""
 pass

class MapLabels:

 def __init__(self, class_map, drop_raw=True):
 self.class_map = class_map

 def __call__(self, dataset, **inputs):
 labels = np.zeros(len(self.class_map), dtype=np.float32)
 for c in inputs["raw_labels"]:
 labels[self.class_map[c]] = 1.0

 transformed = dict(inputs)
 transformed["labels"] = labels
 transformed.pop("raw_labels")

 return transformed

class MixUp(Augmentation):

 def __init__(self, p):
 self.p = p
```

```
def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if np.random.uniform() < self.p:
 first_audio, first_labels = inputs["audio"], inputs["labels"]
 random_sample = dataset.random_clean_sample()
 new_audio, new_labels = mix_audio_and_labels(
 first_audio, random_sample["audio"],
 first_labels, random_sample["labels"]
)

 transformed["audio"] = new_audio
 transformed["labels"] = new_labels

 return transformed

class FlipAudio(Augmentation):
 def __init__(self, p):
 self.p = p

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if np.random.uniform() < self.p:
 transformed["audio"] = np.flipud(inputs["audio"])

 return transformed

class AudioAugmentation(Augmentation):
 def __init__(self, p):
 self.p = p

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if np.random.uniform() < self.p:
 effects_chain = (
 pysndfx.AudioEffectsChain()
 .reverb(
 reverberance=random.randrange(50),
```

```
 room_scale=random.randrange(50),
 stereo_depth=random.randrange(50)
)
 .pitch(shift=random.randrange(-300, 300))
 .overdrive(gain=random.randrange(2, 10))
 .speed(random.uniform(0.9, 1.1))
)
transformed["audio"] = effects_chain(inputs["audio"])

return transformed
```

```
class LoadAudio:
```

```
 def __init__(self):
 pass

 def __call__(self, dataset, **inputs):
 audio, sr = read_audio(inputs["filename"])

 transformed = dict(inputs)
 transformed["audio"] = audio
 transformed["sr"] = sr

 return transformed
```

```
class STFT:
```

```
 eps = 1e-4

 def __init__(self, n_fft, hop_size):
 self.n_fft = n_fft
 self.hop_size = hop_size

 def __call__(self, dataset, **inputs):
 stft = compute_stft(
 inputs["audio"],
 window_size=self.n_fft, hop_size=self.hop_size,
 eps=self.eps)

 transformed = dict(inputs)
 transformed["stft"] = np.transpose(stft)

 return transformed
```

```
class AudioFeatures:

 eps = 1e-4

 def __init__(self, descriptor, verbose=True):

 name, *args = descriptor.split("_")

 self.feature_type = name

 if name == "stft":

 n_fft, hop_size = args
 self.n_fft = int(n_fft)
 self.hop_size = int(hop_size)

 self.n_features = self.n_fft // 2 + 1
 self.padding_value = 0.0

 if verbose:
 print(
 "\nUsing STFT features with params:\n",
 "n_fft: {}, hop_size: {}".format(
 n_fft, hop_size
)
)

 elif name == "mel":

 n_fft, hop_size, n_mel = args
 self.n_fft = int(n_fft)
 self.hop_size = int(hop_size)
 self.n_mel = int(n_mel)

 self.n_features = self.n_mel
 self.padding_value = 0.0

 if verbose:
 print(
 "\nUsing mel features with params:\n",
 "n_fft: {}, hop_size: {}, n_mel: {}".for
mat(
 n_fft, hop_size, n_mel
)
)

 elif name == "raw":

 self.n_features = 1
 self.padding_value = 0.0

 if verbose:
```

```

 print(
 "\nUsing raw waveform features."
)

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if self.feature_type == "stft":
 # stft = compute_stft(
 # inputs["audio"],
 # window_size=self.n_fft, hop_size=self.hop_
size,
 # eps=self.eps, log=True
 #)

 transformed["signal"] = np.expand_dims(inputs["a
udio"], -1)

 elif self.feature_type == "mel":
 stft = compute_stft(
 inputs["audio"],
 window_size=self.n_fft, hop_size=self.hop_si
ze,
 eps=self.eps, log=False
)

 transformed["signal"] = np.expand_dims(inputs["a
udio"], -1)

 elif self.feature_type == "raw":
 transformed["signal"] = np.expand_dims(inputs["a
udio"], -1)

 return transformed

class SampleSegment(Augmentation):
 def __init__(self, ratio=(0.3, 0.9), p=1.0):
 self.min, self.max = ratio
 self.p = p

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if np.random.uniform() < self.p:

```

```
 original_size = inputs["audio"].size
 target_size = int(np.random.uniform(self.min, self.max) * original_size)
 start = np.random.randint(original_size - target_size - 1)
 transformed["audio"] = inputs["audio"][start:start+target_size]

 return transformed
```

```
class ShuffleAudio(Augmentation):
```

```
 def __init__(self, chunk_length=0.5, p=0.5):

 self.chunk_length = chunk_length
 self.p = p

 def __call__(self, dataset, **inputs):

 transformed = dict(inputs)

 if np.random.uniform() < self.p:
 transformed["audio"] = shuffle_audio(
 transformed["audio"], self.chunk_length, sr=
transformed["sr"])

 return transformed
```

```
class CutOut(Augmentation):
```

```
 def __init__(self, area=0.25, p=0.5):

 self.area = area
 self.p = p

 def __call__(self, dataset, **inputs):

 transformed = dict(inputs)

 if np.random.uniform() < self.p:
 transformed["audio"] = cutout(
 transformed["audio"], self.area)

 return transformed
```

```
class SampleLongAudio:
```

```
 def __init__(self, max_length):
```

```
 self.max_length = max_length

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 if (inputs["audio"].size / inputs["sr"]) > self.max_
length:

 max_length = self.max_length * inputs["sr"]

 start = np.random.randint(0, inputs["audio"].siz
e - max_length)
 transformed["audio"] = inputs["audio"][start:sta
rt+max_length]

 return transformed
```

```
class OneOf:
```

```
 def __init__(self, transforms):
 self.transforms = transforms

 def __call__(self, dataset, **inputs):
 transform = random.choice(self.transforms)
 return transform(**inputs)
```

```
class DropFields:
```

```
 def __init__(self, fields):
 self.to_drop = fields

 def __call__(self, dataset, **inputs):
 transformed = dict()

 for name, input in inputs.items():
 if not name in self.to_drop:
 transformed[name] = input

 return transformed
```

```
class RenameFields:
```

```
 def __init__(self, mapping):
```



```

 self.mapping = mapping

 def __call__(self, dataset, **inputs):
 transformed = dict(inputs)

 for old, new in self.mapping.items():
 transformed[new] = transformed.pop(old)

 return transformed

class Compose:

 def __init__(self, transforms):
 self.transforms = transforms

 def switch_off_augmentations(self):
 for t in self.transforms:
 if isinstance(t, Augmentation):
 t.p = 0.0

 def __call__(self, dataset=None, **inputs):
 for t in self.transforms:
 inputs = t(dataset=dataset, **inputs)

 return inputs

class Identity:

 def __call__(self, dataset=None, **inputs):

 return inputs=====
import json

import torch
import umap
import numpy as np
from sklearn.manifold import TSNE
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import label_ranking_average_precision_score, accuracy_score
from matplotlib import pyplot as plt
import librosa

Calculate the overall lwlap using sklearn.metrics function.

```

```

def lwlap(truth, scores):
 """Calculate the overall lwlap using sklearn.metrics.lrap
 ."""
 # sklearn doesn't correctly apply weighting to samples with
 # no labels, so just skip them.
 sample_weight = np.sum(truth > 0, axis=1)
 nonzero_weight_sample_indices = np.flatnonzero(sample_weight
 ht > 0)
 overall_lwlap = label_ranking_average_precision_score(
 truth[nonzero_weight_sample_indices, :] > 0,
 scores[nonzero_weight_sample_indices, :],
 sample_weight=sample_weight[nonzero_weight_sample_indices])
 return overall_lwlap

def load_json(file):
 with open(file, "r") as f:
 return json.load(f)

def get_class_names_from_classmap(classmap):
 r = dict((v, k) for k, v in classmap.items())
 return [r[label] for label in sorted(classmap.values())]

def plot_projection(vectors, labels, frames_per_example=3, newline=False):

 representations = []
 classes = []
 for sample, label in zip(vectors, labels):
 if sum(label) > 1:
 continue
 choices = np.random.choice(
 np.arange(len(sample)), replace=False,
 size=min(frames_per_example, len(sample)))
 representations.extend(sample[choices])
 classes.extend([label.tolist().index(1)] * len(choices))

 representations = np.array(representations)

 # fit a simple model to estimate the quality of the learned representations
 X_train, X_valid, y_train, y_valid = train_test_split(
 representations, classes, shuffle=False, test_size=0.2)

 scaler = StandardScaler().fit(X_train)
 X_train = scaler.transform(X_train)
 X_valid = scaler.transform(X_valid)

```

```
model = KNeighborsClassifier(n_neighbors=5)
model.fit(X_train, y_train)

score = accuracy_score(y_valid, model.predict(X_valid))
if newline:
 print()
print("Classification accuracy: {:.4f}".format(score))

plot projection
embeddings = TSNE().fit_transform(representations)

fig = plt.figure(figsize=(10, 10))
ax = fig.add_subplot(111)
ax.scatter(embeddings[:, 0], embeddings[:, 1], c=classes
, s=10)

fig.canvas.draw()

image = np.array(fig.canvas.renderer._renderer)

plt.close()

return image

def make_mel_filterbanks(descriptor, sr=44100):

 name, *args = descriptor.split("_")

 n_fft, hop_size, n_mel = args
 n_fft = int(n_fft)
 hop_size = int(hop_size)
 n_mel = int(n_mel)

 filterbank = librosa.filters.mel(
 sr=sr, n_fft=n_fft, n_mels=n_mel,
 fmin=5, fmax=None
).astype(np.float32)

 return filterbank

def is_mel(descriptor):
 return descriptor.startswith("mel")

def is_stft(descriptor):
 return descriptor.startswith("stft")

def compute_torch_stft(audio, descriptor):
```

```

 name, *args = descriptor.split("_")

 n_fft, hop_size, *rest = args
 n_fft = int(n_fft)
 hop_size = int(hop_size)

 stft = torch.stft(
 audio,
 n_fft=n_fft,
 hop_length=hop_size,
 window=torch.hann_window(n_fft, device=audio.device)
)

 stft = torch.sqrt((stft ** 2).sum(-1))

 return stft

=====
import os
import glob
import pickle
import random
import json

import torch
from tqdm import tqdm
import torch.utils.data as data
import numpy as np
import pandas as pd

class SoundDataset(data.Dataset):

 def __init__(
 self, audio_files, labels=None,
 transform=None, is_noisy=None, clean_transform=None)
 :

 self.transform = transform
 self.clean_transform = clean_transform
 self.audio_files = audio_files
 self.labels = labels
 self.is_noisy = is_noisy or np.zeros(len(self.audio_
files))

 def __getitem__(self, index):

 sample = dict(
 filename=self.audio_files[index],
 is_noisy=self.is_noisy[index]
)

```

```
 if self.labels is not None:
 sample["raw_labels"] = self.labels[index]

 if self.transform is not None:
 sample = self.transform(dataset=self, **sample)

 return sample

 def random_clean_sample(self):

 index = random.randint(0, len(self) - 1)

 sample = dict(
 filename=self.audio_files[index],
 is_noisy=self.is_noisy[index]
)

 if self.labels is not None:
 sample["raw_labels"] = self.labels[index]

 if self.clean_transform is not None:
 sample = self.clean_transform(dataset=self, **sa
mple)

 return sample

 def __len__(self):
 return len(self.audio_files)
```