

Introduction

see V56 for the best result of LB632 -- Finally I beat the current best public kernel using Keras :) -- This probably be my last update on this kernel -- If you find this kernel helpful, please upvote

Version upto V60 have a silly bug of 'if <-- elif' so that model selection is wrong

This is my effort to do a Keras replication with comparable baseline to the great kernel of @mhiro2 <https://www.kaggle.com/mhiro2/simple-2d-cnn-classifier-with-pytorch> (<https://www.kaggle.com/mhiro2/simple-2d-cnn-classifier-with-pytorch>) (and further improved by @peining), which in turns use the excellent pre-processed data of @daisukelab <https://www.kaggle.com/daisukelab/creating-fat2019-preprocessed-data> (<https://www.kaggle.com/daisukelab/creating-fat2019-preprocessed-data>) -- Note that to inference to the private data in stage-2, you have to preprocess data yourself.

One change I made in a Keras version, in addition to a simple conv net, we can also use a pre-defined architectures [trained from scratch] MobileNetV2 , InceptionV3 and Xception where you can choose in the kernel. Also, many ideas borrow from a nice kernel of @voglinio <https://www.kaggle.com/voglinio/keras-2d-model-5-fold-log-spectrogram-curved-only> (<https://www.kaggle.com/voglinio/keras-2d-model-5-fold-log-spectrogram-curved-only>) , I also borrow the SoftMax+BCE loss & TTA ideas from Giba's kernel (BTW, we all know Giba without having to mention his user :).

I apologize that my code is not at all clean; some of the pytorch code is still here albeit not used.

Major Updates

- V1 [CV680, LB574]
- V4 [CV66x, LB576]
- V5 [] Add image augmentation module
- V9 [CV679] Add lwrap TF metric (credit @rio114 : <https://www.kaggle.com/rio114/keras-cnn-with-lwrap-evaluation> (<https://www.kaggle.com/rio114/keras-cnn-with-lwrap-evaluation>))
- V11 [] Employ list of augmentations mentioned in https://github.com/sainathadapa/kaggle-freesound-audio-tagging/blob/master/approaches_all.md (https://github.com/sainathadapa/kaggle-freesound-audio-tagging/blob/master/approaches_all.md)
- V16 [] Add BCEwithLogits (use only with ACTIVATION = 'linear')
- V17 add SimpleCNN similar to the pytorch baseline
- V22 add Curated-Only, Train-augment options
- V23 add CRNN model
- **V30 LB598 with shallow CNN in 400s, set iteration to 150**
- **V39 LB608 with CoarseDropout Augmentation**
- V40 Simple Snapshot (Checkpoint) Ensemble
- **V52 [CV811, LB616] MixUp+CoarseDropout : credit <https://www.kaggle.com/mathormad/resnet50-v2-keras-focal-loss-mix-up> (<https://www.kaggle.com/mathormad/resnet50-v2-keras-focal-loss-mix-up>)**
- **V56 [CV830, LB632] Change Architecture to get the best result**
- V61 fix silly bugs on model selection

```

In [1]: import gc
import os
import pickle
import random
import time
from collections import Counter, defaultdict
from functools import partial
from pathlib import Path
from psutil import cpu_count
import matplotlib.pyplot as plt

import librosa
import numpy as np
import pandas as pd
from PIL import Image
from sklearn.model_selection import train_test_split
from imgaug import augmenters as iaa
#from skmultilearn.model_selection import iterative_train_test_split

import torch
import torch.nn as nn
import torch.nn.functional as F
from fastprogress import master_bar, progress_bar
from torch.optim import Adam
from torch.optim.lr_scheduler import CosineAnnealingLR
from torch.utils.data import Dataset, DataLoader
from torchvision.transforms import transforms

```

utils

```

In [2]: NUM_CLASSES = 80
SIZE=128
checkpoint_file = ['model_best1.h5', 'model_best2.h5', 'model_best3.h5']
# See Version40 for 3 snapshots (or you can use only 1 which is normal run)
EPOCHS = [432, 0, 0] #150 for inception, 100 for xception
TTA = [19, 0, 0] #Number of test-time augmentation
BATCH_SIZE = 32

LR = 4e-4
PATIENCE = 10 #ReduceOnPlateau option
LR_FACTOR = 0.8 #ReduceOnPlateau option
CURATED_ONLY = True # use only curated data for training
TRAIN_AUGMENT = True # use augmentation for training data?
VALID_AUGMENT = False
MODEL = 'mobile' # 'cnn8th' # choose among 'xception', 'inception', 'mobile',
'crnn', 'simple'
SEED = 520

USE_MIXUP = True
MIXUP_PROB = 0.275

# No K-Fold implementation yet
# NUM_K_FOLDS = 5 # how many folds (K) you gonna splits
# NUM_MODEL_RUN = 5 # how many models (<= K) you gonna train [e.g. set to 1
for a simple train/test split]

# if use BCEwithLogits loss, use Activation = 'linear' only
ACTIVATION = 'linear'
# ACTIVATION = 'softmax'
# ACTIVATION = 'sigmoid'

# LOSS = 'categorical_crossentropy'
# LOSS = 'binary_crossentropy'
LOSS = 'BCEwithLogits'

```

```
In [3]: def seed_everything(seed):  
        random.seed(seed)  
        os.environ['PYTHONHASHSEED'] = str(seed)  
        np.random.seed(seed)  
        torch.manual_seed(seed)  
        torch.cuda.manual_seed(seed)  
        torch.backends.cudnn.deterministic = True  
  
        seed_everything(SEED)
```

```

In [4]: # from official code https://colab.research.google.com/drive/1AgPdhSp7ttY180
3fEoHQKlt_3HJDLi8#scrollTo=cRCaCIb9oguU
def _one_sample_positive_class_precisions(scores, truth):
    """Calculate precisions for each true class for a single sample.

    Args:
        scores: np.array of (num_classes,) giving the individual classifier scores.
        truth: np.array of (num_classes,) bools indicating which classes are true.

    Returns:
        pos_class_indices: np.array of indices of the true classes for this sample.
        pos_class_precisions: np.array of precisions corresponding to each of those classes.
    """
    num_classes = scores.shape[0]
    pos_class_indices = np.flatnonzero(truth > 0)
    # Only calculate precisions if there are some true classes.
    if not len(pos_class_indices):
        return pos_class_indices, np.zeros(0)
    # Retrieval list of classes for this sample.
    retrieved_classes = np.argsort(scores)[::-1]
    # class_rankings[top_scoring_class_index] == 0 etc.
    class_rankings = np.zeros(num_classes, dtype=np.int)
    class_rankings[retrieved_classes] = range(num_classes)
    # Which of these is a true label?
    retrieved_class_true = np.zeros(num_classes, dtype=np.bool)
    retrieved_class_true[class_rankings[pos_class_indices]] = True
    # Num hits for every truncated retrieval list.
    retrieved_cumulative_hits = np.cumsum(retrieved_class_true)
    # Precision of retrieval list truncated at each hit, in order of pos_labels.
    precision_at_hits = (
        retrieved_cumulative_hits[class_rankings[pos_class_indices]] /
        (1 + class_rankings[pos_class_indices].astype(np.float)))
    return pos_class_indices, precision_at_hits

def calculate_per_class_lwlrp(truth, scores):
    """Calculate label-weighted label-ranking average precision.

    Arguments:
        truth: np.array of (num_samples, num_classes) giving boolean ground-truth
        of presence of that class in that sample.
        scores: np.array of (num_samples, num_classes) giving the classifier-
        under-
        test's real-valued score for each class for each sample.

    Returns:
        per_class_lwlrp: np.array of (num_classes,) giving the lwlrp for each
        class.
        weight_per_class: np.array of (num_classes,) giving the prior of each
        class within the truth labels. Then the overall unbalanced lwlrp is
        simply np.sum(per_class_lwlrp * weight_per_class)
    """
    assert truth.shape == scores.shape
    num_samples, num_classes = scores.shape
    # Space to store a distinct precision value for each class on each sample.
    # Only the classes that are true for each sample will be filled in.
    precisions_for_samples_by_classes = np.zeros((num_samples, num_classes))
    for sample num in range(num samples):

```

In [5]: `import tensorflow as tf`

```

# from https://www.kaggle.com/riol14/keras-cnn-with-lwlrwrap-evaluation/
def tf_one_sample_positive_class_precisions(y_true, y_pred) :
    num_samples, num_classes = y_pred.shape

    # find true labels
    pos_class_indices = tf.where(y_true > 0)

    # put rank on each element
    retrieved_classes = tf.nn.top_k(y_pred, k=num_classes).indices
    sample_range = tf.zeros(shape=tf.shape(tf.transpose(y_pred)), dtype=tf.int32)
    sample_range = tf.add(sample_range, tf.range(tf.shape(y_pred)[0], delta=1))
    sample_range = tf.transpose(sample_range)
    sample_range = tf.reshape(sample_range, (-1, num_classes*tf.shape(y_pred)[0]))
    retrieved_classes = tf.reshape(retrieved_classes, (-1, num_classes*tf.shape(y_pred)[0]))
    retrieved_class_map = tf.concat((sample_range, retrieved_classes), axis=0)
    retrieved_class_map = tf.transpose(retrieved_class_map)
    retrieved_class_map = tf.reshape(retrieved_class_map, (tf.shape(y_pred)[0], num_classes, 2))

    class_range = tf.zeros(shape=tf.shape(y_pred), dtype=tf.int32)
    class_range = tf.add(class_range, tf.range(num_classes, delta=1))

    class_rankings = tf.scatter_nd(retrieved_class_map,
                                   class_range,
                                   tf.shape(y_pred))

    #pick up ranks
    num_correct_until_correct = tf.gather_nd(class_rankings, pos_class_indices)

    # add one for division for "presicion_at_hits"
    num_correct_until_correct_one = tf.add(num_correct_until_correct, 1)
    num_correct_until_correct_one = tf.cast(num_correct_until_correct_one, tf.float32)

    # generate tensor [num_sample, predict_rank],
    # top-N predicted elements have flag, N is the number of positive for each sample.
    sample_label = pos_class_indices[:, 0]
    sample_label = tf.reshape(sample_label, (-1, 1))
    sample_label = tf.cast(sample_label, tf.int32)

    num_correct_until_correct = tf.reshape(num_correct_until_correct, (-1, 1))
    retrieved_class_true_position = tf.concat((sample_label,
                                                num_correct_until_correct), axis=1)
    retrieved_pos = tf.ones(shape=tf.shape(retrieved_class_true_position)[0], dtype=tf.int32)
    retrieved_class_true = tf.scatter_nd(retrieved_class_true_position,
                                         retrieved_pos,
                                         tf.shape(y_pred))

    # cumulate predict_rank
    retrieved_cumulative_hits = tf.cumsum(retrieved_class_true, axis=1)

    # find positive position
    pos_ret_indices = tf.where(retrieved_class_true > 0)

    # find cumulative hits

```

```
In [6]: from keras import backend as k
def BCEwithLogits(y_true, y_pred):
    return K.mean(K.binary_crossentropy(y_true, y_pred, from_logits=True), axis=-1)
```

Using TensorFlow backend.

dataset

```
In [7]: dataset_dir = Path('../input/freesound-audio-tagging-2019')
preprocessed_dir = Path('../input/fat2019_prep_mels1')
```

```
In [8]: csvs = {
    'train_curated': dataset_dir / 'train_curated.csv',
    # 'train_noisy': dataset_dir / 'train_noisy.csv',
    'train_noisy': preprocessed_dir / 'trn_noisy_best50s.csv',
    'sample_submission': dataset_dir / 'sample_submission.csv',
}

dataset = {
    'train_curated': dataset_dir / 'train_curated',
    'train_noisy': dataset_dir / 'train_noisy',
    'test': dataset_dir / 'test',
}

mels = {
    'train_curated': preprocessed_dir / 'mels_train_curated.pkl',
    'train_noisy': preprocessed_dir / 'mels_trn_noisy_best50s.pkl',
    'test': preprocessed_dir / 'mels_test.pkl', # NOTE: this data doesn't work at 2nd stage
}
```

```
In [9]: train_curated = pd.read_csv(csvs['train_curated'])
train_noisy = pd.read_csv(csvs['train_noisy'])
if CURATED_ONLY:
    train_df = train_curated
else:
    train_df = pd.concat([train_curated, train_noisy], sort=True, ignore_index=True)
train_df.head()
```

Out[9]:

	fname	labels
0	0006ae4e.wav	Bark
1	0019ef41.wav	Raindrop
2	001ec0ad.wav	Finger_snapping
3	0026c7cb.wav	Run
4	0026f116.wav	Finger_snapping

```
In [10]: test_df = pd.read_csv(csvs['sample_submission'])
test_df.head()
```

```
Out[10]:
```

	fname	Accelerating_and_revving_and_vroom	Accordion	Acoustic_guitar	Applause	Bark	Bass
0	000ccb97.wav	0	0	0	0	0	
1	0012633b.wav	0	0	0	0	0	
2	001ed5f1.wav	0	0	0	0	0	
3	00294be0.wav	0	0	0	0	0	
4	003fde7a.wav	0	0	0	0	0	

```
In [11]: labels = test_df.columns[1:].tolist()
labels[:10]
```

```
Out[11]: ['Accelerating_and_revving_and_vroom',
'Accordion',
'Acoustic_guitar',
'Applause',
'Bark',
'Bass_drum',
'Bass_guitar',
'Bathtub_(filling_or_washing)',
'Bicycle_bell',
'Burping_and_eructation']
```

```
In [12]: num_classes = len(labels)
num_classes
```

```
Out[12]: 80
```

```
In [13]: y_train = np.zeros((len(train_df), num_classes)).astype(int)
for i, row in enumerate(train_df['labels'].str.split(',')):
    for label in row:
        idx = labels.index(label)
        y_train[i, idx] = 1

y_train.shape
```

```
Out[13]: (4970, 80)
```

```
In [14]: with open(mels['train_curated'], 'rb') as curated, open(mels['train_noisy'],
'rb') as noisy:
    x_train = pickle.load(curated)
    if CURATED_ONLY == False:
        x_train.extend(pickle.load(noisy))

with open(mels['test'], 'rb') as test:
    x_test = pickle.load(test)

len(x_train), len(x_test)
```

```
Out[14]: (4970, 1120)
```

```
In [15]: for ii in range(5):  
          print(x_train[ii].shape) #x_train is of shape (TRAIN_NUM,128,LEN,3) [4D Tensor]  
          print(x_test[ii].shape, '\n') #x_test of shape (TEST_NUM,128,LEN,3) [4D Tensor]
```

```
(128, 448, 3)  
(128, 128, 3)
```

```
(128, 131, 3)  
(128, 1021, 3)
```

```
(128, 128, 3)  
(128, 300, 3)
```

```
(128, 1623, 3)  
(128, 1146, 3)
```

```
(128, 128, 3)  
(128, 1442, 3)
```

model


```
In [16]: from keras.layers import *
from keras.models import Sequential, load_model, Model
from keras import metrics
from keras.optimizers import Adam
from keras import backend as K
import keras
from keras.models import Model
from keras.applications.inception_v3 import InceptionV3
from keras.applications.inception_v3 import preprocess_input as preprocess_inception
from keras.applications.mobilenet_v2 import MobileNetV2
from keras.applications.mobilenet_v2 import preprocess_input as preprocess_mobile
from keras.applications.xception import Xception
from keras.applications.xception import preprocess_input as preprocess_xception

from keras.utils import Sequence
from sklearn.utils import shuffle
def create_model_inception(n_out=NUM_CLASSES):

    base_model = InceptionV3(weights=None, include_top=False)

    x0 = base_model.output
    x1 = GlobalAveragePooling2D()(x0)
    x2 = GlobalMaxPooling2D()(x0)
    x = Concatenate()([x1,x2])

    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    x = Dense(256, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    predictions = Dense(n_out, activation=ACTIVATION)(x)

    # this is the model we will train
    model = Model(inputs=base_model.input, outputs=predictions)
    return model
```

```
In [17]: def create_model_xception(n_out=NUM_CLASSES):

    base_model = Xception(weights=None, include_top=False)

    x0 = base_model.output
    x1 = GlobalAveragePooling2D()(x0)
    x2 = GlobalMaxPooling2D()(x0)
    x = Concatenate()([x1,x2])

    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    x = Dense(256, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    # x = Dense(128, activation='relu')(x)
    # x = BatchNormalization()(x)
    # x = Dropout(0.3)(x)

    predictions = Dense(n_out, activation=ACTIVATION)(x)

    # this is the model we will train
    model = Model(inputs=base_model.input, outputs=predictions)
    return model
```

```
In [18]: def create_model_mobile(n_out=NUM_CLASSES):

    base_model = MobileNetV2(weights=None, include_top=False)

    x0 = base_model.output
    x1 = GlobalAveragePooling2D()(x0)
    x2 = GlobalMaxPooling2D()(x0)
    x = Concatenate()([x1,x2])

    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    x = Dense(256, activation='relu')(x)
    x = BatchNormalization()(x)
    x = Dropout(0.5)(x)

    # x = Dense(128, activation='relu')(x)
    # x = BatchNormalization()(x)
    # x = Dropout(0.25)(x)

    predictions = Dense(n_out, activation=ACTIVATION)(x)

    # this is the model we will train
    model = Model(inputs=base_model.input, outputs=predictions)
    return model
```

```
In [19]: def conv_simple_block(x, n_filters):

    x = Convolution2D(n_filters, (3,1), padding="same")(x)
    x = BatchNormalization()(x)
    x = Activation("relu")(x)

    x = Convolution2D(n_filters, (3,1), padding="same")(x)
    x = BatchNormalization()(x)
    x = Activation("relu")(x)
    x = AveragePooling2D()(x)

    return x

def create_model_simplecnn(n_out=NUM_CLASSES):

    inp = Input(shape=(128,128,3))
    # inp = Input(shape=(None,None,3))
    x = conv_simple_block(inp,64)
    x = conv_simple_block(x,128)
    x = conv_simple_block(x,256)
    x = conv_simple_block(x,128)

    # x1 = GlobalAveragePooling2D()(x)
    # x2 = GlobalMaxPooling2D()(x)
    # x = Add()([x1,x2])

    x = Flatten()(x)
    x = Dropout(0.2)(x)

    x = Dense(128, activation='linear')(x)
    x = PReLU()(x)
    x = BatchNormalization()(x)
    x = Dropout(0.2)(x)
    predictions = Dense(n_out, activation=ACTIVATION)(x)

    model = Model(inputs=inp, outputs=predictions)
    return model
```

```

In [20]: def output_of_lambda(input_shape):
          return (input_shape[0], input_shape[2], input_shape[3])

def my_max(x):
    return K.max(x, axis=1, keepdims=False)

def crnn_simple_block(x, n_filters):

    x = Convolution2D(n_filters, (3,1), padding="same")(x)
    x = Activation("relu")(x)

    x = Convolution2D(n_filters, (3,1), padding="same")(x)
    x = Activation("relu")(x)
    x = MaxPooling2D()(x)
    x = Dropout(0.2)(x)

    return x

def create_model_crnn(n_out=NUM_CLASSES):

    # inp = Input(shape=(128,128,3))
    inp = Input(shape=(128,None,3))
    x = crnn_simple_block(inp,64)
    x = crnn_simple_block(x,128)
    x = crnn_simple_block(x,256)

    # eliminate the frequency dimension, x = (batch, time, channels)
    x = Lambda(my_max, output_shape=output_of_lambda)(x)

    x = Bidirectional(CuDNNGRU(128, return_sequences=True))(x)
    # x = Bidirectional(CuDNNLSTM(64, return_sequences=True))(x)
    x = GlobalMaxPooling1D()(x)
    x = Dense(128, activation='linear')(x)
    x = PReLU()(x)
    x = BatchNormalization()(x)
    x = Dropout(0.2)(x)
    predictions = Dense(n_out, activation=ACTIVATION)(x)

    model = Model(inputs=inp, outputs=predictions)
    return model

```

```

In [21]: # from the 8th solution in 2018 competition
# https://github.com/sainathadapa/kaggle-freesound-audio-tagging
def create_model_cnn8th(n_out=NUM_CLASSES):
    regu=0
    inp = Input(shape=(128,128,3))

    x = Conv2D(48, 11, strides=(1,1),kernel_initializer='he_uniform', activation='relu', padding='same',kernel_regularizer=regularizers.l2(regu))(inp)
    x = BatchNormalization()(x)
    x = Conv2D(48, 11, strides=(2,3),kernel_initializer='he_uniform', activation='relu', padding='same',kernel_regularizer=regularizers.l2(regu))(x)
    x = MaxPooling2D(3, strides=(1,2))(x)
    x = BatchNormalization()(x)

    x = Conv2D(128, 5, strides=(1,1),kernel_initializer='he_uniform', activation='relu', padding='same',kernel_regularizer=regularizers.l2(regu))(x)
    x = BatchNormalization()(x)
    x = Conv2D(128, 5, strides=(2,3),kernel_initializer='he_uniform', activation='relu', padding='same',kernel_regularizer=regularizers.l2(regu))(x)
    x = MaxPooling2D(3, strides=2)(x)
    x = BatchNormalization()(x)

    x = Conv2D(192, 3, strides=1,kernel_initializer='he_uniform', activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
    x = Conv2D(192, 3, strides=1,kernel_initializer='he_uniform', activation='relu', padding='same')(x)
    x = BatchNormalization()(x)
    x = Conv2D(128, 3, strides=1,kernel_initializer='he_uniform', activation='relu', padding='same',kernel_regularizer=regularizers.l2(regu))(x)
    x = MaxPooling2D(3, strides=(1,2))(x)
    x = BatchNormalization()(x)

    x = Flatten()(x)
    x = Dense(256, activation='relu')(x)
    x = Dropout(0.5)(x)
    x = Dense(256, activation='relu')(x)
    x = Dropout(0.5)(x)

    predictions = Dense(n_out, activation=ACTIVATION)(x)

    model = Model(inputs=inp, outputs=predictions)
    return model

```

```
In [22]: K.clear_session()
        '''Choose your model here'''
        if MODEL == 'xception':
            preprocess_input = preprocess_xception
            model = create_model_xception(n_out=NUM_CLASSES)
        elif MODEL == 'inception':
            preprocess_input = preprocess_inception
            model = create_model_inception(n_out=NUM_CLASSES)
        elif MODEL == 'mobile':
            preprocess_input = preprocess_mobile
            model = create_model_mobile(n_out=NUM_CLASSES)
        elif MODEL == 'crnn':
            preprocess_input = preprocess_mobile
            model = create_model_crnn(n_out=NUM_CLASSES)
        elif MODEL == 'cnn8th':
            preprocess_input = preprocess_mobile
            model = create_model_cnn8th(n_out=NUM_CLASSES)
        else:
            preprocess_input = preprocess_mobile
            model = create_model_simplecnn(n_out=NUM_CLASSES)

        print(MODEL)
        model.summary()
```

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.
Instructions for updating:
Colocations handled automatically by placer.
WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
mobile

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, None, None, 3 0		
Conv1_pad (ZeroPadding2D) [0][0]	(None, None, None, 3 0		input_1
Conv1 (Conv2D) [0][0]	(None, None, None, 3 864		Conv1_pad
bn_Conv1 (BatchNormalization)	(None, None, None, 3 128		Conv1[0][0]
Conv1_relu (ReLU) [0][0]	(None, None, None, 3 0		bn_Conv1
expanded_conv_depthwise (Depthw [0][0]	(None, None, None, 3 288		Conv1_relu
expanded_conv_depthwise_BN (Bat v_depthwise[0][0]	(None, None, None, 3 128		expanded_con
expanded_conv_depthwise_relu (R v_depthwise_BN[0][0]	(None, None, None, 3 0		expanded_con
expanded_conv_project (Conv2D) v_depthwise_relu[0][0]	(None, None, None, 1 512		expanded_con
expanded_conv_project_BN (Batch v_project[0][0]	(None, None, None, 1 64		expanded_con
block_1_expand (Conv2D) v_project_BN[0][0]	(None, None, None, 9 1536		expanded_con
block_1_expand_BN (BatchNormali nd[0][0]	(None, None, None, 9 384		block_1_exp
block_1_expand_relu (ReLU) nd_BN[0][0]	(None, None, None, 9 0		block_1_exp
block_1_pad (ZeroPadding2D)	(None, None, None, 9 0		block_1_exp

train

```
In [23]: import numpy as np
xx = np.random.rand(1)
print(xx.shape,xx)

xx = np.random.rand(1,1)
print(xx.shape)

(1,) [0.70582626]
(1, 1)
```

```
In [24]:
```



```

In [24]: # If you want, you can try more advanced augmentation like this
augment_img = iaa.Sequential([
#     iaa.ContrastNormalization((0.9, 1.1)),
#     iaa.Multiply((0.9, 1.1), per_channel=0.2),
    iaa.Fliplr(0.5),
#     iaa.GaussianBlur(sigma=(0, 0.1)),
#     iaa.Affine( # x-shift
#         translate_percent={"x": (-0.1, 0.1), "y": (-0.0, 0.0)},
#     ),
    iaa.CoarseDropout(0.12, size_percent=0.05) # see examples : https://g
ithub.com/aleju/imgaug
    ], random_order=True)

# Or you can choose this simplest augmentation (like pytorch version)
# augment_img = iaa.Fliplr(0.5)

# This is my ugly modification; sorry about that
class FATTrainDataset(Sequence):

    def mix_up(x, y):
        x = np.array(x, np.float32)
        lam = np.random.beta(1.0, 1.0)
        ori_index = np.arange(int(len(x)))
        index_array = np.arange(int(len(x)))
        np.random.shuffle(index_array)

        mixed_x = lam * x[ori_index] + (1 - lam) * x[index_array]
        mixed_y = lam * y[ori_index] + (1 - lam) * y[index_array]

        return mixed_x, mixed_y

    def getitem(image):
        # crop 2sec

        base_dim, time_dim, _ = image.shape
        crop = random.randint(0, time_dim - base_dim)
        image = image[:, crop:crop+base_dim, :]

        image = preprocess_input(image)

#         label = self.labels[idx]
        return image
    def create_generator(train_X, train_y, batch_size, shape, augment=False
e, shuffling=False, test_data=False, mixup=False, mixup_prob=0.3):
        assert shape[2] == 3
        while True:
            if shuffling:
                train_X, train_y = shuffle(train_X, train_y)

            for start in range(0, len(train_y), batch_size):
                end = min(start + batch_size, len(train_y))
                batch_images = []
                X_train_batch = train_X[start:end]
                if test_data == False:
                    batch_labels = train_y[start:end]

                for i in range(len(X_train_batch)):
                    image = FATTrainDataset.getitem(X_train_batch[i])
                    if augment:
                        image = FATTrainDataset.augment(image)
                    batch_images.append(image)

                if (mixup and test_data == False):
                    dice = np.random.rand(1)
                    if dice > mixup_prob:
                        batch_images, batch_labels = FATTrainDataset.mix up

```

```
In [25]: from keras.callbacks import (ModelCheckpoint, LearningRateScheduler,
                                         EarlyStopping, ReduceLR0nPlateau, CSVLogger)

from sklearn.model_selection import train_test_split, KFold

reduceLR0nPlat = ReduceLR0nPlateau(monitor='val_tf_lwlap', factor=LR_FACTOR,
                                   patience=PATIENCE,
                                   verbose=1, mode='max', min_delta=0.0001,
                                   cooldown=2, min_lr=1e-5 )

csv_logger = CSVLogger(filename='../working/training_log.csv',
                       separator=',',
                       append=True)

checkpoint = ModelCheckpoint(checkpoint_file[0], monitor='val_tf_lwlap', verbose=1,
                             save_best_only=True, mode='max', save_weights_only = False)
callbacks_list = [checkpoint, csv_logger, reduceLR0nPlat]
```

```
In [26]: # split data into train, valid
x_trn, x_val, y_trn, y_val = train_test_split(x_train, y_train, test_size=0.2,
                                              random_state=SEED)

# create train and valid datagens
train_generator = FATTrainDataset.create_generator(
    x_trn, y_trn, BATCH_SIZE, (SIZE, SIZE, 3), augment=TRAIN_AUGMENT, shuffling=True,
    mixup = USE_MIXUP, mixup_prob = MIXUP_PROB)
validation_generator = FATTrainDataset.create_generator(
    x_val, y_val, BATCH_SIZE, (SIZE, SIZE, 3), augment=VALID_AUGMENT, shuffling=False)
```

```
In [27]: train_steps = np.ceil(float(len(x_trn)) / float(BATCH_SIZE))
val_steps = np.ceil(float(len(x_val)) / float(BATCH_SIZE))
train_steps = train_steps.astype(int)
val_steps = val_steps.astype(int)
print(train_steps, val_steps)
print(len(x_trn))
```

125 32
3976

```
In [28]: print(LOSS)
if LOSS=='BCEwithLogits':
    model.compile(loss=BCEwithLogits,
                  optimizer=Adam(lr=LR),
                  metrics=[tf_lwlap, 'categorical_accuracy'])
else:
    model.compile(loss=LOSS,
                  optimizer=Adam(lr=LR),
                  metrics=[tf_lwlap, 'categorical_accuracy'])
```

BCEwithLogits

```
In [29]: print(LR, PATIENCE, LR_FACTOR, BATCH_SIZE, TRAIN_AUGMENT, USE_MIXUP, MIXUP_PROB)
```

0.0004 10 0.8 32 True True 0.275

```
In [30]: hist = model.fit_generator(  
        train_generator,  
        steps_per_epoch=train_steps,  
        validation_data=validation_generator,  
        validation_steps=val_steps,  
        epochs=EPOCHS[0],  
        verbose=1,  
        callbacks=callbacks_list)
```

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

WARNING:tensorflow:From /opt/conda/lib/python3.6/site-packages/tensorflow/python/ops/math_grad.py:102: div (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Deprecated in favor of operator or tf.math.divide.

Epoch 1/432

125/125 [=====] - 30s 238ms/step - loss: 0.8264 - tf_lwlap: 0.0999 - categorical_accuracy: 0.0133 - val_loss: 0.6638 - val_tf_lwlap: 0.0758 - val_categorical_accuracy: 0.0141

Epoch 00001: val_tf_lwlap improved from -inf to 0.07578, saving model to model_best1.h5

Epoch 2/432

125/125 [=====] - 16s 129ms/step - loss: 0.6238 - tf_lwlap: 0.0978 - categorical_accuracy: 0.0090 - val_loss: 0.4021 - val_tf_lwlap: 0.0717 - val_categorical_accuracy: 0.0221

Epoch 00002: val_tf_lwlap did not improve from 0.07578

Epoch 3/432

125/125 [=====] - 16s 129ms/step - loss: 0.3651 - tf_lwlap: 0.0926 - categorical_accuracy: 0.0103 - val_loss: 0.1388 - val_tf_lwlap: 0.0615 - val_categorical_accuracy: 0.0080

Epoch 00003: val_tf_lwlap did not improve from 0.07578

Epoch 4/432

125/125 [=====] - 16s 130ms/step - loss: 0.1845 - tf_lwlap: 0.0982 - categorical_accuracy: 0.0148 - val_loss: 0.1091 - val_tf_lwlap: 0.0715 - val_categorical_accuracy: 0.0231

Epoch 00004: val_tf_lwlap did not improve from 0.07578

Epoch 5/432

125/125 [=====] - 16s 129ms/step - loss: 0.1222 - tf_lwlap: 0.0999 - categorical_accuracy: 0.0108 - val_loss: 0.1082 - val_tf_lwlap: 0.0747 - val_categorical_accuracy: 0.0191

Epoch 00005: val_tf_lwlap did not improve from 0.07578

Epoch 6/432

125/125 [=====] - 16s 129ms/step - loss: 0.0965 - tf_lwlap: 0.1110 - categorical_accuracy: 0.0168 - val_loss: 0.0800 - val_tf_lwlap: 0.0819 - val_categorical_accuracy: 0.0161

Epoch 00006: val_tf_lwlap improved from 0.07578 to 0.08190, saving model to model_best1.h5

Epoch 7/432

125/125 [=====] - 16s 130ms/step - loss: 0.0881 - tf_lwlap: 0.1187 - categorical_accuracy: 0.0205 - val_loss: 0.0843 - val_tf_lwlap: 0.0866 - val_categorical_accuracy: 0.0191

Epoch 00007: val_tf_lwlap improved from 0.08190 to 0.08659, saving model to model_best1.h5

Epoch 8/432

125/125 [=====] - 17s 134ms/step - loss: 0.0851 - tf_lwlap: 0.1251 - categorical_accuracy: 0.0213 - val_loss: 0.1017 - val_tf_lwlap: 0.0875 - val_categorical_accuracy: 0.0211

Epoch 00008: val_tf_lwlap improved from 0.08659 to 0.08750, saving model to model_best1.h5

Epoch 9/432

125/125 [=====] - 16s 129ms/step - loss: 0.0835 - tf_lwlap: 0.1282 - categorical_accuracy: 0.0238 - val_loss: 0.0800 - val_tf_lwlap: 0.0975 - val_categorical_accuracy: 0.0211

Epoch 00009: val_tf_lwlap improved from 0.08750 to 0.09754, saving model to

```
In [31]: print(K.eval(model.optimizer.lr))
```

```
1e-05
```

```
In [32]: # if LOSS=='BCEwithLogits':
#         model.compile(loss=BCEwithLogits,
#                       optimizer=Adam(lr=3e-4),
#                       metrics=[tf_lwlr, 'categorical_accuracy'])
#     else:
#         model.compile(loss=LOSS,
#                       optimizer=Adam(lr=3e-4),
#                       metrics=[tf_lwlr, 'categorical_accuracy'])

# train_generator = FATTrainDataset.create_generator(
#     x_trn, y_trn, BATCH_SIZE, (SIZE,SIZE,3), augment=TRAIN_AUGMENT,
#     shuffling=True, mixup = False, mixup_prob=0.1)

# EPOCHS = [100, 66, 0]

# print(K.eval(model.optimizer.lr))
```

```
In [33]: if EPOCHS[1] > 0:
        checkpoint = ModelCheckpoint(checkpoint_file[1], monitor='val_tf_lwlr',
        verbose=1,
                                save_best_only=True, mode='max', save_weights_o
nly = False)
        callbacks_list = [checkpoint, csv_logger, reduceLROnPlat]

        hist = model.fit_generator(
            train_generator,
            steps_per_epoch=train_steps,
            validation_data=validation_generator,
            validation_steps=val_steps,
            epochs=EPOCHS[1],
            verbose=1,
            callbacks=callbacks_list)
```

```
In [34]: print(K.eval(model.optimizer.lr))
```

```
1e-05
```

```
In [35]: if EPOCHS[2] > 0:
        checkpoint = ModelCheckpoint(checkpoint_file[2], monitor='val_tf_lwlr',
        verbose=1,
                                save_best_only=True, mode='max', save_weights_o
nly = False)
        callbacks_list = [checkpoint, csv_logger, reduceLROnPlat]

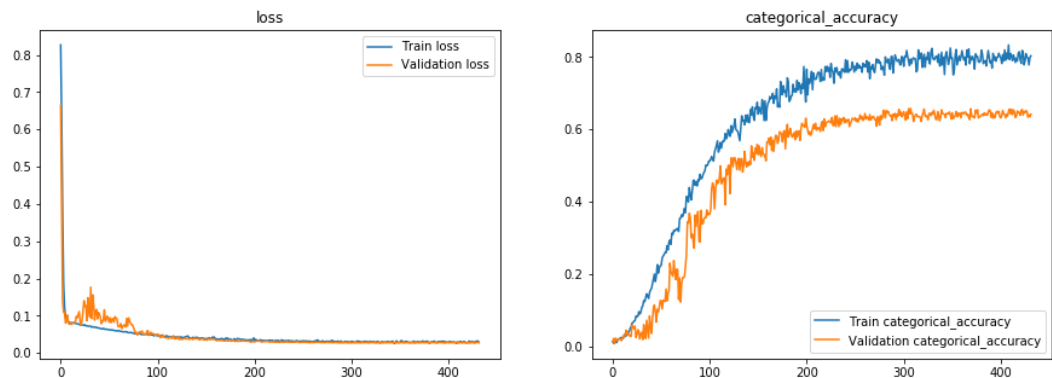
        hist = model.fit_generator(
            train_generator,
            steps_per_epoch=train_steps,
            validation_data=validation_generator,
            validation_steps=val_steps,
            epochs=EPOCHS[2],
            verbose=1,
            callbacks=callbacks_list)
```

```
In [36]: print(K.eval(model.optimizer.lr))
```

```
1e-05
```


```
In [37]: fig, ax = plt.subplots(1, 2, figsize=(15,5))
ax[0].set_title('loss')
ax[0].plot(hist.epoch, hist.history["loss"], label="Train loss")
ax[0].plot(hist.epoch, hist.history["val_loss"], label="Validation loss")
ax[1].set_title('categorical_accuracy')
ax[1].plot(hist.epoch, hist.history["categorical_accuracy"], label="Train ca
tegorical_accuracy")
ax[1].plot(hist.epoch, hist.history["val_categorical_accuracy"], label="Vali
dation categorical_accuracy")
ax[0].legend()
ax[1].legend()
```

Out[37]: <matplotlib.legend.Legend at 0x7fd0e9d46898>



```
In [38]: fig, ax = plt.subplots(1, 2, figsize=(15,5))
ax[0].set_title('tf_lwlrp')
ax[0].plot(hist.epoch, hist.history["tf_lwlrp"], label="Train lwlrp")
ax[0].plot(hist.epoch, hist.history["val_tf_lwlrp"], label="Validation lwlr
ap")
ax[1].set_title('categorical_accuracy')
ax[1].plot(hist.epoch, hist.history["categorical_accuracy"], label="Train ca
tegorical_accuracy")
ax[1].plot(hist.epoch, hist.history["val_categorical_accuracy"], label="Vali
dation categorical_accuracy")
ax[0].legend()
ax[1].legend()
```

Out[38]: <matplotlib.legend.Legend at 0x7fd4b26e4e48>



In [39]:

Calculate Validation Score using TTA

Note that we have to initiate validation_generation everytime before doing a new prediction as `model.fit_generator` will mis-index examples at the end of epoch (and you will get random score)

```
In [39]: model.load_weights(checkpoint_file[0])

validation_generator = FATTrainDataset.create_generator(
    x_val, y_val, BATCH_SIZE, (SIZE,SIZE,3), augment=False, shuffling=False)
pred_val_y = model.predict_generator(validation_generator, steps=val_steps, verbose=1)

for kk in range(len(TTA)):
    for ii in range(TTA[kk]):
        validation_generator = FATTrainDataset.create_generator(
            x_val, y_val, BATCH_SIZE, (SIZE,SIZE,3), augment=False, shuffling=False)

        pred_val_y += model.predict_generator(validation_generator, steps=val_steps, verbose=1)

    if kk+1 < len(TTA) and TTA[kk+1] > 0:
        model.load_weights(checkpoint_file[kk+1])

'''Since the score is based on ranking, we do not need to normalize the prediction'''
# pred_val_y = pred_val_y/10
```

```
32/32 [=====] - 2s 52ms/step
32/32 [=====] - 0s 15ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 0s 15ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 0s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 0s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
32/32 [=====] - 1s 16ms/step
```

```
Out[39]: 'Since the score is based on ranking, we do not need to normalize the prediction'
```

```
In [40]: train_generator = FATTrainDataset.create_generator(
    x_trn, y_trn, BATCH_SIZE, (SIZE,SIZE,3), augment=False, shuffling=False)
pred_train_y = model.predict_generator(train_generator, steps=train_steps, verbose=1)

125/125 [=====] - 2s 15ms/step
```

```
In [41]: import sklearn.metrics
def calculate_overall_lwlap_sklearn(truth, scores):
    """Calculate the overall lwlap using sklearn.metrics.lrap."""
    # sklearn doesn't correctly apply weighting to samples with no labels, so just skip them.
    sample_weight = np.sum(truth > 0, axis=1)
    nonzero_weight_sample_indices = np.flatnonzero(sample_weight > 0)
    overall_lwlap = sklearn.metrics.label_ranking_average_precision_score(
        truth[nonzero_weight_sample_indices, :] > 0,
        scores[nonzero_weight_sample_indices, :],
        sample_weight=sample_weight[nonzero_weight_sample_indices])
    return overall_lwlap
```

```
In [42]: print(pred_val_y.shape, y_val.shape)
print(np.sum(pred_val_y), np.sum(y_val))
# for ii in range(len(y_val)):
#     print(np.sum(pred_val_y[ii]), np.sum(y_val[ii]))
```

```
(994, 80) (994, 80)
-11160808.0 1128
```

```
In [43]: print("lwlap from sklearn.metrics for training data =", calculate_overall_lwlap_sklearn(y_trn, pred_train_y))
print("val lwlap from sklearn.metrics =", calculate_overall_lwlap_sklearn(y_val, pred_val_y/10))

score, weight = calculate_per_class_lwlap(y_val, pred_val_y)
lwlap = (score * weight).sum()
print('direct calculation of val lwlap : %.4f' % (lwlap))
```

```
lwlap from sklearn.metrics for training data = 0.9937980679367769
val lwlap from sklearn.metrics = 0.7947953814827754
direct calculation of val lwlap : 0.7948
```

Simple Error Analysis

```
In [44]: idx = np.sum(y_val,axis=1) > 1
print(y_val[idx, :].shape, y_val[idx==False, :].shape)

print("val lwlap for multi-labels =", calculate_overall_lwlap_sklearn(y_val[idx, :], pred_val_y[idx, :]))
print("val lwlap for single-label =", calculate_overall_lwlap_sklearn(y_val[idx==False, :], pred_val_y[idx==False, :]))
```

```
(125, 80) (869, 80)
val lwlap for multi-labels = 0.7788500531195031
val lwlap for single-label = 0.7995477865991014
```

Predict Test Data with TTA

```
In [45]:
```

```
In [45]: test_steps = np.ceil(float(len(x_test)) / float(BATCH_SIZE)).astype(int)
```



```

In [46]: model.load_weights(checkpoint_file[0])

test_generator = FATTrainDataset.create_generator(
    x_test, x_test, BATCH_SIZE, (SIZE,SIZE,3), augment=False, shuffling=False, test_data=True)
pred_test_y = model.predict_generator(test_generator, steps=test_steps, verbose=1)

for kk in range(len(TTA)):
    for ii in range(TTA[kk]):
        test_generator = FATTrainDataset.create_generator(
            x_test, x_test, BATCH_SIZE, (SIZE,SIZE,3), augment=False, shuffling=False, test_data=True)

        pred_test_y += model.predict_generator(test_generator, steps=test_steps, verbose=1)

    if kk+1 < len(TTA) and TTA[kk+1] > 0:
        model.load_weights(checkpoint_file[kk+1])

35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 15ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 15ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 15ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 15ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 16ms/step
35/35 [=====] - 1s 15ms/step

```

```

In [47]: sort_idx = np.argsort(labels).astype(int)

```

```

In [48]: print(sort_idx)

```

```

[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47
 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71
 72 73 74 75 76 77 78 79]

```

```
In [49]: sample_sub = pd.read_csv('../input/freesound-audio-tagging-2019/sample_submission.csv')
test_Y_sort = pred_test_y[:, sort_idx]
sample_sub.iloc[:, 1:] = test_Y_sort
sample_sub.to_csv('submission.csv', index=False)

sample_sub.head()
```

Out[49]:

	fname	Accelerating_and_revving_and_vroom	Accordion	Acoustic_guitar	Applause	B
0	000ccb97.wav	-154.853333	-178.622345	-137.645279	-154.654419	-124.198
1	0012633b.wav	-62.905537	-155.926025	-138.813126	-119.286285	-128.623
2	001ed5f1.wav	-101.077438	-147.583435	-141.133286	-120.893265	-146.542
3	00294be0.wav	-173.398453	-147.612381	-148.860428	-205.831070	-120.307
4	003fde7a.wav	-167.963501	-154.409439	-134.879196	-167.252792	-130.279